

Awareness Growth in Bayesian Networks:

The Role of Structural Assumptions

July 28, 2022

Wordcount (including footnotes): 5,392

We examine different counterexamples to Reverse Bayesianism, a popular theory that addresses the problem of awareness growth. We agree with the general skepticism toward Reverse Bayesianism, but submit that the problem of awareness growth cannot be tackled in an algorithmic manner, because subject-matter, structural assumptions need to be made explicit. Thanks to their ability to express probabilistic dependencies, we illustrate how Bayesian networks can help to model awareness growth in the Bayesian framework.

1 Introduction

Learning is modeled in the Bayesian framework by the rule of conditionalization. This rule posits that the agent's new degree of belief in a proposition H after a learning experience E should be the same as the agent's old degree of belief in H conditional on E . That is,

$$P^E(H) = P(H|E),$$

where $P()$ represents the agent's old degree of belief (before the learning experience E) and $P^E()$ represents the agent's new degree of belief (after the learning experience E).

Both E and H belong to the agent's algebra of propositions. This algebra models the agent's awareness state, the propositions taken to be live possibilities. Conditionalization never modifies the algebra and thus makes it impossible for an agent to learn something they have never thought about. Even before learning about E , the agent must already have assigned a degree of belief to any proposition conditional on E . This picture commits the agent to the specification of their 'total possible future experience' (Howson, 1976), as though learning was

1 confined to an ‘initial prison’ (Lakatos, 1968).

2 But, arguably, the learning process is more complex than what conditionalization allows.
3 Not only do we learn that some propositions that we were entertaining are true or false, but
4 we may also learn new propositions that we did not entertain before. Or we may entertain
5 new propositions—without necessarily learning that they are true or false—and this change
6 in awareness may in turn change what we already believe. How should this more complex
7 learning process be modeled by Bayesianism? Call this the problem of awareness growth.

8 The algebra of propositions need not be so narrowly construed that it only contains proposi-
9 tions that are presently under consideration. The algebra may also contain propositions which,
10 though outside the agent’s present consideration, are still the object, perhaps implicitly, of
11 certain dispositions to believe.¹ But even this expanded algebra will have to be revised sooner
12 or later. The algebra of propositions could in principle contain anything that could possibly be
13 conceived, expressed, thought of. Such a rich algebra would not need to change at any point,
14 but this is an implausible model of ordinary agents with bounded resources such as ourselves.

15 Critics of Bayesianism and sympathizers alike have been discussing the problem of awareness
16 growth under different names for quite some time, at least since the eighties. This problem arises
17 in a number of different contexts, for example, new scientific theories (Chihara, 1987; Earman,
18 1992; Glymour, 1980), language changes and paradigm shifts (Williamson, 2003), and theories
19 of induction (Zabell, 1992). A proposal that has attracted considerable scholarly attention
20 in recent years is Reverse Bayesianism (Bradley, 2017; Karni & Vierø, 2015; Wenmackers
21 & Romeijn, 2016). The idea is to model awareness growth as a change in the algebra while
22 ensuring that the proportions of probabilities of the propositions shared between the old and
23 new algebra remain the same in a sense to be specified.

24 Let \mathcal{F} be the initial algebra of propositions and let \mathcal{F}^+ the algebra after the agent’s awareness
25 state has grown. Both algebras contain the contradictory and tautologous propositions \perp and
26 \top , and they are closed under connectives such as disjunction \vee , conjunction \wedge and negation \neg .
27 Denote by X and X^+ the subsets of these algebras that contain only basic propositions, namely
28 those without connectives. **Reverse Bayesianism** posits that the ratio of probabilities for any
29 basic propositions A and B in both X and X^+ —the basic propositions shared by the old and

¹Roussos (2021) notes that, for the sake of clarity, the problem of awareness growth should only address propositions which agents are *truly* unaware of (say new scientific theories), not propositions that were temporarily forgotten or set aside. This is a helpful clarification to keep in mind, although the recent literature on the topic does not make a sharp distinction between true unawareness and temporary unawareness.

1 new algebra—remain constant through the process of awareness growth:

$$\frac{P(A)}{P(B)} = \frac{P^+(A)}{P^+(B)},$$

2 where $P()$ represents the agent's degree of belief before awareness growth and $P^+()$ represents
3 the agent's degree of belief after awareness growth.

4 Reverse Bayesianism is an elegant theory that manages to cope with a seemingly intractable
5 problem. As the awareness state of an agent grows, the agent would prefer not to throw
6 away completely the epistemic work they have done previously. The agent may desire to retain
7 as much of their old degrees of beliefs as possible. Reverse Bayesianism provides a simple
8 recipe to do that. It also coheres with the conservative spirit of Bayesian conditionalization
9 which preserves the old probability distribution conditional on what is learned.

10 Unfortunately, Reverse Bayesianism does not deliver the intuitive results in all cases. There
11 is no shortage of counterexamples against it in the recent philosophical literature (Mathani,
12 2020; Steele & Stefánsson, 2021). In addition, attempts to extend traditional arguments in
13 defense of Bayesian conditionalization to the case of awareness growth seem to hold little
14 promise (Pettigrew, forthcoming). If the consensus in the literature is that Reverse Bayesianism
15 is not the right theory of awareness growth, what theory (if any) should replace it?

16 Here we offer a diagnosis of what is wrong with Reverse Bayesianism and outline an
17 alternative proposal. The problem of awareness growth—we hold—cannot be tackled in an
18 algorithmic manner because subject-matter assumptions, both probabilistic and structural, need
19 to be made explicit. So any theory of awareness growth cannot be a purely formal theory. This
20 does not mean, however, we should give up on probability theory altogether in this context.
21 Thanks to its ability to express probabilistic dependencies, we think that the theory of Bayesian
22 networks can help to model awareness growth in the Bayesian framework. We illustrate this
23 claim as we examine different counterexamples to Reverse Bayesianism.

24 **2 Expansion**

25 A common set of cases of awareness growth are usually referred in the literature with the label
26 *awareness expansion*. A precise definition of expansion can be tricky to provide, but a rough
27 characterization will suffice here. Suppose, as is customary, propositions are interpreted as sets

1 of possible worlds, where the set of all possible worlds is the possibility space. Awareness
2 expansion occurs when a new proposition is added to the algebra and its interpretation includes
3 possible worlds not in the original possibility space. So the addition of the new proposition
4 causes the original space of possibilities to expand.

5 Perhaps the most straightforward example of awareness expansion occurs when you become
6 aware of a new explanation for the evidence at your disposal which you had not considered
7 before. This can happen in many fields of inquiry: medicine, law, science, everyday affairs.
8 Here is a scenario by Steele & Stefánsson (2021):

9 FRIENDS: Suppose you happen to see your partner enter your best friend's house
10 on an evening when your partner had told you she would have to work late. At
11 that point, you become convinced that your partner and best friend are having
12 an affair, as opposed to their being warm friends or mere acquaintances. You
13 discuss your suspicion with another friend of yours, who points out that perhaps
14 they were meeting to plan a surprise party to celebrate your upcoming birthday—a
15 possibility that you had not even entertained. Becoming aware of this possible
16 explanation for your partner's behaviour makes you doubt that she is having an
17 affair with your friend, relative, for instance, to their being warm friends. (Steele
18 & Stefánsson, 2021, sec. 5, Example 2)

19 Initially, the algebra only contained the hypotheses 'my partner and my best friend met to
20 have an affair' (*Affair*) and 'my partner and my best friend met as friends or acquaintances'
21 (*Friends/acquaintances*). These were the only explanations you considered for your evidence,
22 namely that your partner and your best friend met one night without telling you (*Secretive*).
23 Given this evidence, *Affair* is more probable than *Friends/acquaintances*:²

$$P(\textit{Affair}|\textit{Secretive}) > P(\textit{Friends/acquaintances}|\textit{Secretive}). \quad (>)$$

24 When the algebra changes, a new hypothesis is added which you had not considered before:
25 your partner and your best friends met to plan a surprise party for your upcoming birthday
26 (*Surprise*). This hypothesis seems to be a much better explanation. So, given the same evidence,

²This assumes that the prior probabilities of the two hypotheses were not strongly skewed in one direction. If you were initially nearly certain your partner could not possibly have an affair, even the fact they behaved very secretly or lied to you might not affect the probability of the two hypotheses.

1 the hypothesis *Surprise* is more likely than the hypothesis *Affair*:

$$P^+(Surprise|Secretive) > P^+(Affair|Secretive).$$

2 And since *Surprise* implies *Friends/acquaintances*—after all, in order to prepare a surprise
3 party, your partner and best friend have to be at least acquaintances— *Friends/acquaintances* is
4 now more likely than *Affair*:

$$P^+(Affair|Secretive) < P^+(Friends/acquaintances|Secretive). \quad (<)$$

5 As Steele & Stefánsson note, the conjunction of ($>$) and ($<$) violates Reverse Bayesianism.
6 But, they also note, a quick fix is available. For consider the following condition, called
7 **Awareness Rigidity**:

$$P^+(A|T^*) = P(A),$$


8 where T^* corresponds to a proposition that picks out, from the vantage point of the new
9 awareness state, the entire possibility space before the episode of awareness growth. Aware-
10 ness rigidity establishes that, once a suitable proposition T^* is identified, the old probability
11 assignments can be kept conditional on T^* . In our running example, $\neg Surprise$ is the suitable
12 proposition T^* : *that* there were was no surprise party in the making picks out the original
13 possibility space in its entirety. After all, before awareness growth, the eventuality that your
14 partner and best friend could be organizing a surprise for you had been tacitly ruled out. So,
15 conditional on $\neg Surprise$, no probability assignment should change, including the probability
16 of *Affair*. Thus,

$$P^+(Affair|Secretive \& \neg Surprise) > P^+(Friends/acquaintances|Secretive \& \neg Surprise).$$

17 Steele and Stefansson, however, go on to show that Awareness Rigidity does not hold in
18 other cases, what they call awareness refinement. But before considering them, it is important
19 to understand why Awareness Rigidity is an adequate principle to govern awareness growth in
20 expansion cases. As we shall soon see, Bayesian networks helps to model these cases more
21 precisely. A Bayesian network is a compact formalism to represent probabilistic dependencies.
22 A Bayesian network consists of a direct acyclic graph (DAG) accompanied by a probability

1 distribution. The nodes in the graph represent random variables that can take different values.
2 We will use ‘nodes’ and ‘variables’ interchangeably. The nodes are connected by arrows, but
3 no loops are allowed, hence the name direct acyclic graph. Bayesian networks are relied upon
4 in many fields, but have been rarely deployed to model awareness growth (the exception is
5 Williamson (2003)). We think instead they are a good framework for this purpose. Awareness
6 growth can be modeled as a change in the graphical network—nodes and arrows are changed,
7 added or erased—as well as a change in the probability distribution from the old to the new
8 network.

9 To model FRIENDS with Bayesian networks, we start with this graph, which is the usual
10 hypothesis-evidence idiom (Fenton, Neil, & Lagnado, 2013):



12 where H is the hypothesis node and E the evidence node. If an arrow goes from H to E , the
13 full probability distribution associated with the Bayesian network is defined by two probability
14 tables.³ One table defines the prior probabilities for all the states of H , and another table defines
15 the conditional probabilities of the form $P(E = e | H = h)$, where uppercase letters represent
16 the variables (nodes) and lower case letters represent the values of these variables. These two
17 probability tables are sufficient to specify the full probability distribution.

18 Initially, before awareness growth, the hypothesis H takes only two values, *Friends/acquaintances*
19 and *Affair*. These two values are meant to be exhaustive. So, initially, *Affair* functions as
20 the negation of *Friends/acquaintances*, or vice versa. After awareness growth—specifically,
21 awareness expansion—the two values of the node H are no longer considered exhaustive. The
22 node now has a third value: *Surprise*. The rest of the structure of the network remains intact.

23 So, expansion simply consists in the addition of an extra value in one of the nodes of the
24 network. Specifically, in FRIENDS, the extra value was added to the upstream node H . Since
25 the extra value was added in the upstream node, this requires a change in the probability table
26 for the conditional probabilities $P(E = e | H = h)$. However, for these conditional probabilities,
27 a restricted variant of Reverse Bayesianism still holds. For all values e and h of upstream node

³A major point of contention in the interpretation of Bayesian networks is the meaning of the directed arrows. They could be interpreted causally—as though the direction of causality proceeds from the events described by the hypothesis to event described by the evidence—but they need not be; see footnote 11.

1 H and downstream node E in the old network, the following constraint holds:

$$\frac{P(E = e|H = h)}{P(E = e|H = h)} = \frac{P^+(E = e|H = h)}{P^+(E = e|H = h)}. \quad (C)$$

2 What would happen is the extra value had been added to the downstream node E ? For suppose
3 that the evidence node E could initially take only two values, say *Secretive* and *Public*. You then
4 realize the evidence node could also take a third value, say *Ambiguous*. Would that realization
5 mandate a change in the old conditional probabilities $P(E = e|H = h)$? Since the change is
6 downstream the old conditional probabilities $P(E = e|H = h)$ would remain the same, and thus
7 *a fortiori*, constraint (C) would be satisfied.

8 The upshot of this discussion is this. Awareness expansion can be modeled by changes
9 in the Bayesian network used to represent the epistemic state of the agent. The structure
10 of the network does not change, but values are added to the upstream or downstream node.
11 This addition of values to the nodes can be carried while satisfying a restricted version of
12 Reverse Bayesian, what we called constraint (C). This constraint achieves the same result as
13 Awareness Rigidity, but more easily applies to the conditional probabilities that are needed in
14 the probability tables of Bayesian Networks. The challenge now is to develop a systematic
15 method to determine when constraint (C) is satisfied and when it fails. The structure of the
16 Bayesian network will be our guide. This will afford us a firmer foundation to develop a general
17 theory of awareness growth.

18 **3 Refinement**


19 We turn now from cases of awareness expansion to cases of awareness refinement. In the
20 framework of Bayesian networks, expansion consisted in added values to nodes in the network.
21 Refinement, instead, can be modeled by adding nodes to the network without necessarily add
22 any new values to the existing nodes. Intuitively, refinement takes place when an epistemic
23 agent acquires a more-fined grained picture of the situation, say instead of thinking that the
24 political spectrum is divided into liberal and conservatives, the political spectrum can be further
25 divided into traditional-liberal, new-liberal, traditional-conservative and new-conservative. The
26 political spectrum is still divided into liberal and conservative—not expansion occurred—but
27 these two categories have been further refined.

1 Although there is no shortage of counterexamples to Reverse Bayesianism when it comes
2 to awareness refinement, we begin with our own. This will allow us to underscore the role
3 of subject-matter assumptions in theorizing about awareness growth. Consider the following
4 scenario:

5 LIGHTING: You have evidence that favors a certain hypothesis, say a witness saw
6 the defendant around the crime scene. You give some weight to this evidence.
7 In your assessment, that the defendant was seen around the crime scene (your
8 evidence) raises the probability that the defendant was actually there (your hypoth-
9 esis). But now you ask, what if it was dark when the witness saw the defendant?
10 In light of your realization that it could have been dark, you wonder whether (and
11 if so how) you should change the probability that you assigned to the hypothesis
12 that the defendant was around the crime scene.

13 As your awareness grows, you do not learn anything specific about the lighting conditions,
14 neither that they were bad nor that they were good. You simply wonder what they were, a
15 variable you had previously not considered. So no Bayesian updating takes place in the strict
16 sense, although broadly speaking some new information has been introduced.⁴ Something has
17 changed in your epistemic state—you have a more fine-grained assessment of what could have
18 happened—but it is not clear what you should do in this scenario. Since the lighting conditions
19 could have been bad but could also have been good, perhaps you should just stay put until you
20 learn something more specific.

21 In what follows, we illustrate how Bayesian networks helps to model what is going on in
22 LIGHTING and conclude that you should probably revise downward your confidence in the
23 hypothesis that the defendant was around the crime scene. The starting point of our analysis is
24 the usual hypothesis-evidence idiom, repeated below for convenience:



26 Since you trust the evidence, you think that the evidence is more likely under the hypothesis


⁴HERE EXPLAIN DIFFERENCE WITH STEELE AND STEFANSSON. The process of awareness growth in LIGHTING adds only one extra variable, lighting conditions, while MOVIES adds two extra variables, language difficulty and whether the owner is simple-minded or not. Further, MOVIES contains a clear-cut case of learning, that the owner *is* simple-minded. This is not so in LIGHTING. Strictly speaking, you are learning that it is *possible* that the lighting conditions were bad. However, you are not conditioning on the proposition ‘the lighting conditions were bad’ or ‘the lighting conditions were good’. So you are not learning about the lighting conditions in the sense in which learning is understood in this paper.

1 that the defendant was present at the crime scene than under the alternative hypothesis:

$$P(E=seen|H=present) > P(E=seen|H=absent)$$

2 The inequality is a qualitative ordering of how plausible the evidence is in light of competing
3 hypotheses. No matter the numbers, by the probability calculus, it follows that the evidence
4 raises the probability of the hypothesis $H=present$.

5 Now, as you wonder about the lighting conditions, the graph should be amended:



6
7 where the node L can have two values, $L=good$ and $L=bad$. Commonsense as well as psycho-
8 logical findings suggest that when the visibility deteriorates, people's ability to identify faces
9 worsen. So a plausible way to modify your assessment of the evidence is as follows:

$$P^+(E=seen|H=present \wedge L=good) > P^+(E=seen|H=absent \wedge L=good)$$

10

$$P^+(E=seen|H=present \wedge L=bad) = P^+(E=seen|H=absent \wedge L=bad)$$

11 In words, if the lighting conditions were good, you still trust the evidence like you did before
12 (first line), but if the lighting conditions were bad, you regard the evidence as no better than
13 chance (second line). These probabilistic constraints are plausible, but should ultimately be
14 grounded on verifiable empirical regularities.

15 Despite the change in awareness, you have not learned anything in the strict sense. Your
16 new stock of evidence does not contain neither the information that the lighting conditions
17 were bad nor that they were good. But the Bayesian network structure that represents your
18 epistemic state is now more fine-grained. The network contains the new variable L which it
19 did not contain prior to the episode of awareness growth. In addition—and this is the crucial

1 point—the new variable bears certain *structural relationships* with the variables H and E . The
 2 graphical network represents a direct probabilistic dependency between the lighting conditions
 3 L and the witness sensory experience E , but does not allow for any direct dependency between
 4 the lighting conditions and the fact that the defendant was (or was not) at the crime scene.
 5 There is no direct arrow between the nodes L and H . This structure of dependencies captures
 6 our causal intuitions about the scenario: the lighting conditions do affect what the witness
 7 could see, but do not directly affect what the defendant might have done.

8 Without Bayesian networks, episodes of awareness growth could only be modeled by the
 9 addition of new propositions that were not previously in the algebra. But this approach would
 10 fail to capture crucial information. When awareness growth takes place against the background
 11 of an intuitive causal structure of the world—as in the case of LIGHTING—this structure should
 12 also be modeled. Bayesian networks offer a formal framework that can do precisely that.

13 This model of causal structure can now guide us to decide whether the restricted version of
 14 Reverse Bayesianism, what we called constraint (C), holds in this scenario. Specifically, we
 15 need to assess whether the following holds:

$$\frac{P(E = \text{seen} | H = \text{present})}{P(E = \text{seen} | H = \text{absent})} = \frac{P^+(E = \text{seen} | H = \text{present})}{P^+(E = \text{seen} | H = \text{absent})}.$$

16 The question here is whether you should assess the evidence at your disposal—that the witness
 17 saw the defendant at the crime scene—any differently than before. As noted earlier, without a
 18 clear model of the scenario, it might seem that you should simply stay put. After all, besides
 19 the sensory experience of the witness, you have gained no novel information about the lighting
 20 conditions. Should you thus conclude that the evidence has the same value before and after the
 21 realization that lighting could have been bad?

22 The evidence would have the same value if the likelihood ratios associated with it relative to
 23 the competing hypotheses were the same before and after awareness growth. But, in changing
 24 the probability function from $P()$ to $P^+()$, it would be quite a coincidence if this were true. In
 25 our example, many possible probability assignments violate this equality. If before awareness
 26 growth you thought the evidence favored the hypothesis $H = \text{present}$ to some extent, after the

1 growth in awareness, the evidence is likely to appear less strong.⁵ If this is correct, this outcome
 2 violates constraint (C). Reverse Bayesianism is also violated since the ratio of the probabilities
 3 of $H=present$ to $E=seen$, before and after awareness growth, has changed:

$$\frac{P^{E=seen}(H=present)}{P^{E=seen}(E=seen)} \neq \frac{P^{+,E=seen}(H=present)}{P^{+,E=seen}(E=seen)},$$

4 where $P^{E=seen}()$ and $P^{+,E=seen}()$ represent the agent's degrees of belief, before and after aware-
 5 ness growth, updated by the evidence $E=seen$.⁶

6 The general lesson to be learned here has to do with the importance of formalizing structural
 7 assumptions and the role of Bayesian networks in modeling awareness growth. Modeling those

⁵By the law of total probability, the right hand side of the equality in (C) should be expanded, as follows:

$$\frac{P^+(E=e|H=h)}{P^+(E=e|H=h')} = \frac{P^+(E=seen \wedge L=good|H=present) + P^+(E=seen \wedge L=bad|H=present)}{P^+(E=seen \wedge L=good|H=absent) + P^+(E=seen \wedge L=bad|H=absent)}.$$

For concreteness, let's use some numbers:

$$P(E=seen|H=present) = P^+(E=seen|H=present \wedge L=good) = .8$$

$$P(E=seen|H=absent) = P^+(E=seen|H=absent \wedge L=good) = .4$$

$$P^+(E=seen|H=present \wedge L=bad) = P^+(E=seen|H=absent \wedge L=bad) = .5.$$

$$P^+(L=bad) = P^+(L=good) = .5.$$

So the ratio $\frac{P(E=seen|H=present)}{P(E=seen|H=absent)}$ equals 2. After the growth in awareness, the ratio $\frac{P^+(E=seen|H=present)}{P^+(E=seen|H=absent)}$ will drop to $\frac{.65}{.45} \approx 1.44$. The calculations here rely on the dependency structure encoded in the Bayesian network (see starred step below).

$$\begin{aligned} P^+(E=seen|H=present) &= P^+(E=seen \wedge L=good|H=present) + P^+(E=seen \wedge L=bad|H=present) \\ &= P^+(E=seen|H=present \wedge L=good) \times P^+(L=good|H=present) \\ &\quad + P^+(E=seen|H=present \wedge L=bad) \times P^+(L=bad|H=present) \\ &= * P^+(E=seen|H=present \wedge L=good) \times P^+(L=good) \\ &\quad + P^+(E=seen|H=present \wedge L=bad) \times P^+(L=bad) \\ &= .8 \times .5 + .5 \times .5 = .65 \end{aligned}$$

$$\begin{aligned} P^+(E=seen|H=absent) &= P^+(E=seen \wedge L=good|H=absent) + P^+(E=seen \wedge L=bad|H=absent) \\ &= P^+(E=seen|H=absent \wedge L=good) \times P^+(L=good|H=absent) \\ &\quad + P^+(E=seen|H=absent \wedge L=bad) \times P^+(L=bad|H=absent) \\ &= * P^+(E=seen|H=absent \wedge L=good) \times P^+(L=good) \\ &\quad + P^+(E=seen|H=absent \wedge L=bad) \times P^+(L=bad) \\ &= .4 \times .5 + .5 \times .5 = .45 \end{aligned}$$

This argument can be repeated with many other numerical assignments.

⁶The scenario also violates Awareness Rigidity which requires that $P^+(A|T^*) = P(A)$, where T^* corresponds to a proposition that picks out, from the vantage point of the new awareness state, the entire possibility space before the episode of awareness growth. In LIGHTING, however, T^* does not change, so Awareness Rigidity would require that $P^+(A) = P(A)$, and instead in the scenario, we have

$$P^+(H=present|E=seen) \neq P(H=present|E=seen).$$


1 structural assumptions allows us to see that constraint (C)—as well as Reverse Bayesianism
 2 more generally—fails here. To strengthen this point, consider this variation of the LIGHTING
 3 scenario:

4 VERACITY: A witness saw that the defendant was around the crime scene and
 5 you initially took this to be evidence that the defendant was actually there. But
 6 then you worry that the witness might be lying or misremembering what happened.
 7 Perhaps, the witness was never there, made things up or mixed things up. Should
 8 you reassess the evidence at your disposal? If so, how?

9 It might seem that this scenario is no different from LIGHTING. The realization that lighting
 10 could be bad should make you less confident in the truthfulness of the sensory evidence. And
 11 the same conclusion should presumably follow from the realization that the witness could be
 12 lying. So both scenarios would be counterexamples to Reverse Bayesianism. But, upon closer
 13 scrutiny, things are not that simple. To run the two scenarios together would be a mistake.

14 The evidence at your disposal in LIGHTING is the sensory evidence—the experience of
 15 seeing—and the possibility of bad lighting does affect the quality of your visual experience.
 16 So, if lighting was indeed bad, this would warrant lowering your confidence in the truthfulness
 17 of the visual experience. But the possibility of lying in VERACITY does not affect the quality
 18 of the visual experience in and of itself, although it affects the quality of the *reporting* of
 19 that experience. So, if the witness did lie, this would not warrant lowering your confidence
 20 in the truthfulness of the visual experience, only in the truthfulness of the reporting of that
 21 experience. The distinction between the visual experience and its reporting is crucial here.
 22 Bayesian networks help to model this distinction precisely, and then see why LIGHTING and
 23 VERACITY are structurally different.

24 The graphical network should initially look like the initial DAG for LIGHTING, consisting
 25 of the hypothesis node H upstream and the evidence node E downstream. As your awareness
 26 grows, the graphical network should be updated by adding another node R further downstream:



28 As before, the hypothesis node H bears on the whereabouts of the defendant and has two values,
 29 $H=present$ and $H=absent$. Note the difference between E and R . The evidence node E bears
 30 on the visual experience had by the witness. The reporting node R , instead, bears on what the

1 witness reports to have seen. The chain of transmission from ‘visual experience’ to ‘reporting’
2 may fail for various reasons, such as lying or misremembering.

3 In VERACITY, the conditional probabilities, $P(E = e|H = h)$ should be the same as $P^+(E =$
4 $e|H = h)$ for any values e and h of the variables H and E that are shared before and after
5 awareness growth. In comparing the old and new Bayesian network, this equality falls out
6 from their structure, as the connection between H and E remains unchanged. Thus, constraint
7 (C)—along with Reverse Bayesianism—is perfectly fine in scenarios such as VERACITY.

8 This does not mean that the assessment of the probability of the hypothesis $H=present$ should
9 undergo no change. If you worry that the witness could have lied, this should presumably
10 make you less confident about $H=present$. To accommodate this intuition, VERACITY can
11 be interpreted as a scenario in which an episode of awareness refinement takes place together
12 with a form of retraction. At first, after the learning episode, you update your belief based on
13 the *visual experience* of the witness. But after the growth in awareness, you realize that your
14 learning is in fact limited to what the witness *reported* to have seen. The previous learning
15 episode is retracted and replaced by a more careful statement of what you learned: instead
16 of conditioning on $E=seen$, you should condition on what the witness reported to have seen,
17 $R=seen-reported$. This retraction will affect the probability of the hypothesis $H=present$.

18 Where does this leave us? Refinement cases that might at first appear similar can be
19 structurally different in important ways, and this difference can be appreciated by looking
20 at the Bayesian networks used to model them. In modeling VERACITY, the new node is
21 added downstream, while in modeling LIGHTING, it is added upstream. This difference affects
22 how probability assignments should be revised. Since the conditional probabilities associated
23 with the upstream nodes are unaffected, Reverse Bayesianism is satisfied in VERACITY.⁷ By
24 contrast, since the conditional probabilities associated with the downstream node will often
25 have to change, Reverse Bayesianism fails in LIGHTING.

26 This further corroborates our working hypothesis: structural features about how we con-
27 ceptualize a specific scenario are the guiding principles about how we update the probability
28 function through awareness growth, not a formal principle like Reverse Bayesianism. We
29 further elaborate on this conjecture by drawing on some examples from Anna Mathani.

⁷Note that $P(H=present|E=seen) \neq P(H=present|R=seen-reported)$, but since you are conditioning on different propositions, this does not conflict with Reverse Bayesianism.

1 3.1 Sure no-gain bets

2 Suppose the witness reports to have seeing the defendant around the crime scene. You are not
3 aware that the witness could be lying. Thus, you are 100% confident that the witness saw is
4 what they report to have seeing. In fact, you make no distinction between reporting to have
5 seeing and seeing itself. So you would be willing to buy for 1\$ the following bet: if the witness
6 saw the defendant, you get 1\$, and 0\$ otherwise. If the witness did see the defendant, you
7 get you 1\$ back, and otherwise you loose \1\$. You are 100% sure the witness did see the
8 defendant, so—by your lights—you stand to loose no money whatsoever from this bet. But
9 suppose that, as a matter of fact, there is a difference between reporting and seeing. So, the
10 witness might report to have seeing something without actually having seeing it. So, contrary
11 to your conviction, that the witness saw the defendant is not 100% probable. This means that
12 you would be willing to engage in a bet in which you are guaranteed not to win any money and
13 could potentially lose money. If the witness did see the defendant you would get your 1\$ back,
14 but if not, you would lose it.

15 4 Mathani's counterexamples

16 Mathani (2020) offers two counterexamples to Reverse Bayesianism. The first goes like this:

17 TENANT: Suppose that you are staying at Bob's flat which he shares with his
18 landlord. You know that Bob is a tenant, and that there is only one landlord, and
19 that this landlord also lives in the flat. In the morning you hear singing coming
20 from the shower room, and you try to work out from the sounds who the singer
21 could be. At this point you have two relevant propositions that you consider
22 possible ... *Landlord* standing for the possibility that the landlord is the singer,
23 and *Bob* standing for the possibility that Bob is the singer ... Because you know
24 that Bob is a tenant in the flat, you also have a credence in the proposition *Tenant*
25 that the singer is a tenant. Your credence in *Tenant* is the same as your credence
26 in *Bob*, for given your state of awareness these two propositions are equivalent
27 ... Now let's suppose the possibility suddenly occurs to you that there might be
28 another tenant living in the same flat (*Other*).

29 Initially, you thought the singer could either be the landlord or Bob, the tenant. Then you come

1 to the realization that a third person could be the singer, another tenant. Before awareness
2 growth, that Bob is in the shower and that a tenant is in the shower are equivalent descriptions.
3 After awareness growth, this equivalence breaks down.

4 Why is this scenario problematic? Suppose, after you hear singing in the shower, you become
5 sure someone is in there, but you cannot tell who. So $P(Landlord) = P(Bob) = 1/2$, and since
6 *Bob* and *Tenant* are equivalent, also $P(Tenant) = 1/2$. Now, *Landlord*, *Bob* and *Tenant* are all
7 propositions that you were originally aware of, and thus Reverse Bayesianism requires that their
8 assigned probabilities should remain in the same proportion after your awareness grows. But
9 note that *Other* entails *Tenant* and *Bob* and *Other* are disjoint, so it follows that $P^+(Other)$
10 must have zero probability.⁸ This is an undesired outcome that rules out the possibility that
11 there could be a third person in the shower.⁹

12 Consider now Mathani's second counterexample:

13 COIN: You know that I am holding a fair ten pence UK coin which I am about to
14 toss. You have a credence of 0.5 that it will land *Heads*, and a credence of 0.5 that
15 it will land *Tails*. You think that the tails side always shows an engraving of a lion.
16 So you also have a credence of 0.5 that (*Lion*) it will land with the lion engraving
17 face-up: relative to your state of awareness *Tails* and *Lion* are equivalent.... Now
18 let's suppose that you somehow become aware that occasionally ten pence coins
19 have an engraving of Stonehenge on the tails side.

20 *Tails* and *Lion* are equivalent propositions prior to awareness growth. Suppose you initially
21 gave *Tails* and *Lion* the same credence. Reverse Bayesianism requires that their relative
22 proportions should stay the same after awareness grow. The same applies to *Heads* and *Tails*.
23 But since *Lion* and *Stonehenge* are incompatible and the latter entails *Tails*, you should have
24 $P^+(Stonehenge) = 0$, again an undesirable conclusion.


25 Mathani notes that COIN has the same structure as TENANT. This is true to some extent, but
26 there is also an interesting asymmetry between the two scenarios. In TENANT, it is natural to
27 assign 1/3 to *Landlord*, *Bob* and *Other* after awareness growth. That someone is singing in
28 the shower is evidence that someone must be in there, but without any more discriminating
29 evidence, each person should be assigned the same probability. Consequently, a probability

⁸If $P^+(Other) > 0$, the proportion of *Tenant* to *Landlord* or the proportion of *Bob* to *Landlord* should change.

⁹Awareness Rigidity is no of help either because it would require that $P^+(Landlord|Landlord \vee Tenant) = P^+(Bob|Landlord \vee Tenant)$ both equal 1/2, thus forcing $P^+(Other|Landlord \vee Tenant)$ to zero.

1 of 2/3 should be assigned to *Tenant*. On this picture, the proportion of *Landlord* to *Tenant*
2 changes from 1:1 (before awareness growth) to 1:2 (after awareness growth). But, in COIN,
3 the relative proportion of *Heads* to *Tails* should remain constant throughout, unless evidence
4 emerges that the coin is not fair. One might have expected that *Landlord* and *Tenant* would
5 behave just like *Heads* and *Tails*, but actually they do not.

6 Bayesian networks can help to model the asymmetry between these two scenarios. Consider
7 COIN first. The structure of the scenario is represented by the following graph:



9 The upstream node *Outcome* has two states, *tails* and *heads*. These two states remain the same
10 throughout. What changes are the states associated with the *Image* node downstream. Before
11 awareness growth, the node *Image* has two states: *lions* and *heads-image*.¹⁰ You assume that
12 *Image = lions* is true if and only if *Outcome = tails* is true. Then, you come to the realization
13 that the imagines for tails include a lion or a stonehenge engraving. So, after awareness growth,
14 the node *Image* contains three states: *lion*, *stonehenge* and *heads-image*. Consider now the
15 other scenario, TENANT. We start with the following graph:



17 Initially, the upstream node *Person* has two possible states, representing who is in the bathroom
18 singing: *landlord-person* and *bob*. To simplify things, the assumption here is that the evidence
19 of singing has already ruled out the possibility that no one would be in the shower. The
20 downstream node *Role* has also two values, *landlord* and *tenant*. After your awareness grows,
21 the upstream node *Person* should now have one more possible state, *other*.

22 The difference in modeling the two scenarios is this. In COIN, the states of the upstream
23 node remain fixed, whereas in TENANT, they change. After awareness growth, no new state
24 is added to *Outcome*, but an additional state, *other*, is added to *Person*. Plausible probability
25 distributions for the Bayesian networks associated with the two scenarios are displayed in
26 Table 1. How the networks should be built and which probabilities should shift is based on
27 our background knowledge. This knowledge tells us that the equiprobability of *heads* and

¹⁰The heads side must have some image, not specified in the scenario.

1 *tails* should not be affected by realizing that *stonhenge* is another possible engraving for the
2 tails side. It also tells us that the probabilities of *landlord* and *tenant* should be affected by
3 realizing that a third person could be in the shower.

4 We conclude with some programmatic remarks. We think that the awareness of agents grows
5 while holding fixed certain material structural assumptions, based on commonsense, semantic
6 stipulations or causal dependency.¹¹ To model awareness growth, we need a formalism that
7 can express these material structural assumptions. This can be done using Bayesian networks,
8 and we offered some illustrations of this strategy. These material assumptions also guide
9 us in formulating the adequate conservative constraints, and these will inevitably vary on
10 a case-by-case basis. The literature on awareness growth from a Bayesian perspective is
11 primarily concerned with a formal, almost algorithmic solution to the problem. Insofar as
12 Reverse Bayesianism is an expression of this formalistic aspiration, we agree with Steele and
13 Stefánsson that we are better off looking elsewhere.

14 **5 Towards a general theory**

15 Awareness growth can occur in different ways. The key question is to what extent probability
16 assignments that were made prior to the episode of awareness growth can be retained. There
17 seems to be no clear rule that can decide that. We propose the following procedure. Construct a
18 Bayesian network prior to awareness growth and compare it with the new Bayesian network
19 after awareness growth. If the new arrows and nodes are all downstream, the old probabilities
20 table should not be changed. The paradigmatic cases of this are scenarios VERACITY and
21 COIN. If, instead, the new arrows and nodes are upstream, the old probabilities tables
22 should be changed. The paradigmatic examples are LIGHTING and TENANT.

23 **6 Counterexamples**

24 In this section, we rehearse two of the counterexamples to Reverse Bayesianism by Steele and
25 Stefánsson. One example targets awareness expansion and the other awareness refinement
26 (more on this distinction soon). We show why they make a limited case against Reverse

¹¹ Arrows in Bayesian networks are often taken to represent causal relationships, but other interpretations exist. Schaffer (2016) discusses an interpretation in which arrows represent grounding relations rather than causality.

P(Image Outcome)		Outcome	
Image	lion	heads	tails
	heads-image	0	1
		1	0
P+(Image Outcome)		Outcome	
Image	lion	heads	tails
	stonehenge	0	1/2
	heads-image	0	1/2
		1	0
P(Outcome) = P+(Outcome)		Outcome	
	heads	tails	
	1/2	1/2	

P(Role Person)		Person		
Role	tenant	landlord-person	bob	
	landlord	0	1	
		1	0	
P+(Role Person)		Person		
Role	tenant	landlord-person	bob	other
	landlord	0	1/2	1/2
		1	0	0
P(Person)		Person		
	landlord-person	bob		
	1/2	1/2		
P+(Person)		Person		
	landlord-person	bob	other	
	1/3	1/3	1/3	

Table 1: Top table displays a plausible probability distribution for COIN and bottom table does the same for TENANT.

1 Bayesianism and then provide a better counterexample with the aid of Bayesian networks.

2 6.1 Friends and Movies

3 The difference between expansion and refinement is intuitively plausible, but can be tricky
4 to pin down formally. A rough characterization will suffice here. Suppose, as is customary,
5 propositions are interpreted as sets of possible worlds, where the set of all possible worlds is
6 the possibility space. An algebra of propositions thus interpreted induces a partition of the
7 possibility space. Refinement occurs when the new proposition added to the algebra induces a
8 more fine-grained partition of the possibility space. Expansion occurs when the new proposition
9 is inconsistent with the existing ones, thus making the old partition no longer exhaustive.

10 This is not the end of the story, however. Steele and Stefánsson offer another counterexample
11 that also works against Awareness Rigidity, this time targeting a case of refinement:

12 MOVIES: Suppose you are deciding whether to see a movie at your local cinema.
13 You know that the movie's predominant language and genre will affect your
14 viewing experience. The possible languages you consider are French and German
15 and the genres you consider are thriller and comedy. But then you realise that,
16 due to your poor French and German skills, your enjoyment of the movie will
17 also depend on the level of difficulty of the language. Since it occurs to you that
18 the owner of the cinema is quite simple-minded, you are, after this realisation,
19 much more confident that the movie will have low-level language than high-level
20 language. Moreover, since you associate low-level language with thrillers, this
21 makes you more confident than you were before that the movie on offer is a thriller
22 as opposed to a comedy. (Steele & Stefánsson, 2021, sec. 5, Example 3)

23 This is a case of refinement. For you initially categorized movies by just language and genre,
24 and then you refined your categorization by adding another variable, level of difficulty. Without
25 considering language difficulty, you assigned the same probability to the hypotheses *Thriller*
26 and *Comedy*. But learning that the owner was simple-minded made you think that the level
27 of linguistic difficulty must be low and the movie most likely a thriller rather than a comedy
28 (perhaps because thrillers are simpler—linguistically—than comedies). So, against Reverse
29 Bayesianism, MOVIES violates the condition $\frac{P(\textit{Thriller})}{P(\textit{Comedy})} = \frac{P^+(\textit{Thriller})}{P^+(\textit{Comedy})}$.

30 The counterexample also violates Awareness Rigidity. For consider a proposition that picks

1 out the entire possibility space, for example, *Thriller* \vee *Comedy*.¹² Awareness Rigidity would
2 require that $P(\textit{Thriller}) = P^+(\textit{Thriller}|\textit{Thriller} \vee \textit{Comedy})$. But MOVIES does not satisfy this
3 equality since the probability of *Thriller* has gone up.

4 How good of a counterexample is this? Steele and Stefánsson consider an objection:

5 It might be argued that our examples are not illustrative of ... a simple growth in
6 awareness; rather, our examples illustrate and should be expressed formally as
7 complex learning experiences, where first there is a growth in awareness, and then
8 there is a further learning event ... In this way, one could argue that the awareness-
9 growth aspect of the learning event always satisfies Reverse Bayesianism.

10 Admittedly, MOVIES can be split into two episodes. In the first, you entertain a new variable
11 besides language and genre, namely the language difficulty of the movie. In the second episode,
12 you learn something you did not consider before, namely that the owner is simple-minded.
13 Could Reverse Bayesianism still work for the first episode, but not the second? Steele and
14 Stefánsson do not address this question explicitly, but insist that no matter the answer both
15 episodes are instances of awareness growth. We agree with them on this point. Awareness
16 growth is both *entertaining* a new proposition not in the initial awareness state of the agent
17 and *learning* a new proposition. Nonetheless, many could still wonder. Is the second episode
18 (learning something new) necessary for the counterexample to work together with the first
19 episode (mere refinement without learning)?

20 Suppose the counterexample did work only in tandem with an episode of learning something
21 new. If that were so, defenders of Reverse Bayesianism or Awareness Rigidity could still
22 claim that their theory applies to a large class of cases. It applies to cases of awareness
23 refinement without learning and also to cases of awareness expansion. For recall that the
24 first putative counterexample featuring awareness expansion—FRIENDS—did not challenge
25 Reverse Bayesianism insofar as the latter is formulated in terms of its close cousin, Awareness
26 Rigidity. So the force of Steele and Stefánsson's counterexamples would be rather limited.

27 References

28 Bradley, R. (2017). *Decision theory with a human face*. Cambridge University Press.

¹²Since MOVIES is a case of refinement, *Thriller* \vee *Comedy* picks out the entire possibility space both before and after awareness growth.

- 1 Chihara, C. S. (1987). Some problems for bayesian confirmation theory. *British Journal for*
2 *the Philosophy of Science*, 38(4), 551–560.
- 3 Earman, J. (1992). *Bayes or bust? A critical examination of bayesian confirmation theory*.
4 MIT press.
- 5 Fenton, N., Neil, M., & Lagnado, D. A. (2013). A General Structure for Legal Arguments
6 About Evidence Using Bayesian Networks. *Cognitive Science*, 37(1), 61–102. <https://doi.org/10.1111/cogs.12004>
- 7 Glymour, C. (1980). *Theory and evidence*. Princeton University Press.
- 8 Howson, C. (1976). The development of logical probability. In *Essays in memory of imre*
9 *lakatos. Boston studies in the philosophy of science* (pp. 277–298). Springer.
- 10 Karni, E., & Vierø, M.-L. (2015). Probabilistic sophistication and reverse bayesianism. *Journal*
11 *of Risk and Uncertainty Volume*, 50, 189–208.
- 12 Lakatos, I. (1968). Changes in the problem of inductive logic. *Studies in Logic and the*
13 *Foundations of Mathematics*, 51, 315–417.
- 14 Mathani, A. (2020). Awareness growth and dispositional attitudes. *Synthese*, 198(9), 8981–
15 8997.
- 16 Pettigrew, R. (forthcoming). How should your beliefs change when your awareness grows?
17 *Episteme*, 1–25. <https://doi.org/10.1017/epi.2022.33>
- 18 Roussos, J. (2021). Awareness growth and belief revision. *Manuscript*.
- 19 Schaffer, J. (2016). Grounding in the image of causation. *Philosophical Studies*, 173, 49–100.
- 20 Steele, K., & Stefánsson, O. (2021). Belief revision for growing awareness. *Mind*, 130(520),
21 1207–1232.
- 22 Wenmackers, S., & Romeijn, J.-W. (2016). New theory about old evidence: A framework for
23 open-minded bayesianism. *Synthese Volume*, 193, 1225–1250.
- 24 Williamson, J. (2003). Bayesianism and language change. *Journal of Logic, Language, and*
25 *Information*, 12(1), 53–97.
- 26 Zabell, S. (1992). Predicting the unpredictable. *Synthese*, 90(1), 205–232.
- 27