

Likelihood ratio and evidence strength

Marcello Di Bello and Rafal Urbaniak

Contents

1. Likelihood ratio as a measure of evidence strength	1
2. The risk of false positive and its impact	6
3. Eyewitness identification and likelihood ratio	13
4. Hypothesis choice	17
5. Levels of Hypotheses and the two-stain problem	20
6. Relevance and the small-town murder scenario	23
Appendices	25
A. The cold-hit confusion	25
B. Likelihood ratio and cold-hit DNA matches	32
C. Confirmation measures	34
References	36

Contents

1. Likelihood ratio as a measure of evidence strength

The fallacies we considered earlier in the book — such as the base rate fallacy, the prosecutor’s fallacy, and the defense attorney’s fallacy—show how the posterior probability can be misjudged, upwards or downwards, even if the subject gets the likelihoods right. These examples illustrate that the assessment of the posterior probability of a hypothesis given the evidence depends also on the prior probability of the hypothesis. The correctness of such an assessment therefore requires that the priors are chosen sensibly (or that a range of sensible priors is considered) and appropriately put together with the likelihoods involved. Quite crucially, the posterior probability given a piece of evidence should not be confused with the probative value of a given piece of evidence itself with respect to the hypothesis in question.

Consider the following examples. Suppose the prior probability of a given hypothesis H is low, say $P(H) = .001$, but taking evidence E into account brings this probability up to $.35$, that is, $P(H|E) = .35$. This is a dramatic upward shift. Even though the posterior probability of H given E is not very high, E strongly favors H . Conversely, suppose the prior probability of H is extremely high, say $P(H) = .999$, but taking evidence E into account brings this probability down to $.75$, that is, $P(H|E) = .75$. This

Marcello: Comments directed at Marcello

Alicja: Comments directed at Alicja

Rafal: Comments directed at Rafal

General comments.

Read and discuss the decision Marcello sent

I’m not convinced. Where do you think this fits and how deeply do you want to go into this??? Allen objection to infinite complexity, plausibility and probability?, The nature of juridical proof: Probability as a tool in plausible reasoning, s intuitive, reduction of complexities in choosing hypotheses in LR

write the conclusions section

add crossref

is a dramatic downward shift. Even though the posterior probability of H given E is still quite high, E speaks against H . Or suppose an expert testifies that the blood found at the crime scene matches the defendant's and it is .5 probable that a person unrelated to the crime would match by coincidence. Absent any other compelling evidence to the contrary, it should initially be very likely that the defendant, as anyone else, had little to do with the crime. Say, for illustrative purposes, that the prior probability is .01, and let the probability of a match if the suspect is the source be approximately 1. In this setting, the posterior probability that the defendant is the source is roughly .17. This might seem unimpressive, but it does not mean that the incriminating evidence was weak. While the match was not not strong enough to make it very likely that the defendant was the source of the traces, the posterior probability is seventeen times larger than the prior. Similarly, in the Collins case, the posterior probability jumped from the $1/6 \times 10^6$ prior to .7 after taking the match into account. Still not enough for a conviction, but a remarkable increase nonetheless. These examples illustrate how measuring the strength of evidence in terms of the posterior it leads to seems inappropriate.

So how do we capture the strength of an item of evidence that reflects the impact the evidence on the posterior probability? One measure of the strength of evidence is the likelihood of the evidence compared to the prior of the evidence (this measure is sometimes called the *Bayes factor*):

$$\text{BF}(E, H) = \frac{P(E|H)}{P(E)}. \quad (\text{BF})$$

The Bayes factor is one probabilistic measure of the extent to which the evidence, regardless of the absolute posterior probability, supports or does not support the hypothesis. It seem to be an intuitively plausible measure of evidential strength. Note that by Bayes' theorem

$$P(H|E) = \text{BF}(H, E) \times P(H)$$

and so the Bayes factor is greater than one if and only if the posterior probability $P(H|E)$ is higher than the prior probability $P(H)$, $P(H) < P(H|E)$. So E *positively* supports H whenever the Bayes factor is greater than one. The greater the Bayes factor (for values above one), the greater the upward shift from prior to posterior probability, the more strongly E positively supports H . In line with the motivating examples, the posterior probability of H given E could still be low even if the Bayes factor is significantly above one. Conversely, again by Bayes' theorem, the probability of H given E is lower than the probability of H , $P(H) > P(H|E)$ just in case the Bayes factor is less than one. So E *negatively* supports H whenever the Bayes factor is less than one. In general, the smaller the Bayes factor (for values below one), the greater the downward shift from prior to posterior probability, the more strongly E negatively supports H . If $P(H) = P(H|E)$, the evidence has no impact on the probability of H .

One reason to think the Bayes Factor is a useful measure of evidential strength is that it appropriately deviates from 1, its point of neutrality. But let us pause a moment to think about the denominator in (BF). It can be calculated following the law of total probability:

$$P(E) = P(E|H)P(H) + P(E|\neg H)P(\neg H). \quad (1)$$

The catch-all alternative hypothesis $\neg H$ can be replaced by a more fine-grained set of alternatives, say H_1, H_2, \dots, H_k , provided H and these alternatives are exclusive and cover the entire space of possibilities (that is, they form a partition). The law of total probability would then read:

$$P(E) = P(E|H)P(H) + \sum_{i=1}^k P(E|H_i)P(H_i). \quad (2)$$

For simplicity, let's stick to (1) for now, and use it to rewrite (BF):

$$\text{BF}(E, H) = \frac{P(E|H)}{P(E|H)P(H) + P(E|\neg H)P(\neg H)}. \quad (3)$$

What should be clear from this formulation is that the Bayes factor depends on the prior probabilities. Indeed, suppose $P(E|H) = 1$ and $P(E|\neg H) = .1$. If $P(H) = .1$, $P(E)$, the denominator, is .19 and so the

Bayes Factor is approximately 5.26. If, however, $P(H) = .2$, the denominator is .28 and the Bayes Factor is approximately 3.57. In fact, a more general look (Figure 1) shows the prior probability can have larger impact on the Bayes factor than the likelihood $P(E|\sim H)$.

This seems to be a reason against using the Bayes factor as a measure of evidential strength in the context of legal fact-finding. After all, in such contexts, we would like the expert's assessment not to depend on the expert's prior convictions about the hypothesis, and for all the agents involved in legal fact-finding to understand the expert's statement in the same way, even if they assign different priors to the hypothesis.

Bayes factor as a function of prior and $P(E|\sim H)$.




Figure 1: Impact of the prior and likelihood of E given H for probabilities in $(0, 0.05)$ and Bayes Factor restricted to $(0, 250)$ for visibility.

A related reason to worry about the denominator of (3) is that assessing the strength of evidence using the Bayes factor seems to impose too great a cognitive burden on an agent, since it would require estimating $P(E)$. This rarely can be done directly, and estimation using the denominator of (3) or (2) (in a more complex case) not only requires that the agent sifts through the entire space of possibilities, but also that the agent uses as weights a sensible selection of priors for the hypotheses involved.

Here is another reason to hesitate about measuring evidential support with BF, stemming from (Gillies, 1986).¹ Consider the hypothesis H =: “the suspect is guilty” and suppose it is already known that E =: “the suspect killed the victim” (say here guilt requires both *actus reus*, the killing, and *mens rea*, the intention). Clearly, H entails E , and E provides positive support for H . Now consider a composite hypothesis $H' =$: “the suspect is guilty and we live in a simulation built by aliens.” We hope the reader shares the intuition that the support E provides for H' is somewhat weaker than it does for H (yes, the example of H' is really far-fetched, but this is the point: addition of far-fetched hypotheses should decrease the support). This feature, however, cannot be captured by BF. In general, suppose $H \models E$ (and so, also, $H \wedge X \models E$). Then both $P(E|H)$ and $P(E|H \wedge X)$ equal 1. But this means that $BF(H, E) = BF(H \wedge X, E) = 1/P(E)$, and so the Bayes factors for the two support relations are equal. This is the problem of irrelevant conjuncts.

Other probabilistic measures of confirmation have been proposed and do deserve a discussion of their applicability in the context of legal evidence evaluation. However, as this would detract us from the study of the properties and applications of the most promising candidate that in fact has been used by legal probabilists (likelihood ratio), we postpone painting a more complex picture of various confirmation measures, the conditions they satisfy, and the reasons why they are not preferable to likelihood ratio to Appendix C.

Suppose we are after a measure that does not depend on priors, puts no such cognitive requirements on an agent, and is sensitive to the addition of irrelevant conjuncts. Clearly, we should not simply

¹See also (Fitelson, 1999) for a discussion.

Think about this again later. Plausible, but I think this needs some more careful discussion. Diagnostic properties of medical tests might change depending on base rate?

Marcello: added this flow here, and introduced the notion of irrelevant conjuncts above, check

Marcello: Now clearly introduced this notion when criticizing BF, check.

use $P(E|H)$. For one thing, in many interesting cases this conditional probability will be very close to one and will not allow us to distinguish between the strengths of pieces of evidence that we should distinguish. For instance, what is the probability that the blood types match if the accused is the source? Well, one, pretty much. What is the probability that the DNA profiles match if the accused is the source? Again, one, pretty much. But obviously a DNA profile match is not on par with a blood type match insofar as strength of evidence is involved.

Even if this conditional probability is not close to one, there are reasons not to use it as a measure of evidential strength. Consider an example by Triggs & Buckleton (2004). In a child abuse case, the prosecutor offers evidence that a couple's child rocks and that only 3% of non-abused children rock, $P(\text{child rocks}|\text{no abuse}) = .3$. If it is unlikely that a non-abused child would rock, the fact that this child rocks might seem strong evidence of abuse. But this reading of the 3% figure is mistaken. It could well be that 3% of abused children rock, $P(\text{child rocks}|\text{abuse}) = .3$. Note that the two probabilities need not add up to 1. Similarly, learning only that $P(\text{child rocks}|\text{abuse}) = .3$ does not provide full information needed for evidence evaluation, and one also needs information about $P(\text{child rocks}|\text{no abuse}) = .3$. In our particular case, given that rocking is equally unlikely under either hypothesis, rocking cannot count as evidence of abuse, and any of the low conditional probabilities involved alone does not allow us to notice this. Thus, in order to avoid exaggerations of the evidence both conditional probabilities need to be involved in the evidence strength evaluation (ENFSI, 2015; Royall, 1997; Triggs & Buckleton, 2004).

One issue that these considerations illustrate is that what matters is also the probability of the evidence if the hypothesis is false. If the accused is not the source, the probability of a blood match if the accused is not the source, while small, is much higher than the probability of a DNA profile match, and this seems to explain why the latter piece of evidence is stronger. So, both the probability of the evidence given the hypothesis, and the probability of evidence given an alternative hypothesis should be somehow factored into a useful measure of evidential strength.

One straightforward way to implement this is to use the **likelihood ratio**, a comparative measure of whether evidence E supports a hypothesis H more than a competing hypothesis H' , in symbols:

$$LR(E, H, H') = \frac{P(E|H)}{P(E|H')}. \quad (LR)$$

If the evidence supports H more than H' , the ratio would be above one, and if the evidence supports H' more than H , the ratio would be below one. So, as with the Bayes factor, support levels correspond to deviations from one. The greater the likelihood ratio (for values above one), the stronger the evidence in favor of H as contrasted with H' . The smaller the likelihood ratio (for values below one), the stronger the evidence in favor of the competing hypothesis H' as contrasted with H .

One reason to prefer LR to BF is that LR is not susceptible to the addition of irrelevant hypotheses. Even if $P(E|H) = P(E|H \wedge X) = 1$, we have that $LR(E, H) = 1/P(E|\neg H)$, while $LR(E, H \wedge X) = 1/P(E|\neg H \vee \neg X)$. Crucially, the denominators might differ and so might the resulting likelihoods.

The likelihood ratio does not depend on the priors, and is a simpler and more workable measure than the Bayes factor, since it does not require one to think about the probability of the evidence in general, namely $P(E)$. Think of blood match evidence. To calculate the Bayes factor, we need two things: $P(E|H)$, which here—say—is close to one, and $P(E)$, which is the probability of a blood match *not* conditional on whether the suspect is the source. Yet, direct and reliable estimation of this probability seems to be quite hard. Of course, one could try to use the law of total probability, as in (1), but then they would need to use $P(E|\neg H)$ which, with $P(E|H)$, would already suffice for the calculation of LR, and additional information about other probabilities used in the right-hand side. In this sense, it seems that the calculation of BF is more committing and requires more information than the calculation of LR.

This apparent simplicity of LR, however, can often give rise to errors in the assessment of the evidence, especially if the two hypotheses are not chosen carefully. As it will transpire, the choice of the hypotheses that are conditioned upon is crucial. In the most straightforward case, H' is simply the negation of H . In many practical contexts such a simplistic set-up, however, is not viable. We will discuss these issues in detail in this chapter later on.

Moreover, LR allows for a clearer separation of the impact of the priors from the impact of the evidence on the posterior probability. The relationship between likelihood ratio $P(E|H)/P(E|H')$ and posterior odds $P(H|E)/P(H'|E)$ is apparent in the odds version of Bayes' theorem:

$$\frac{P(H|E)}{P(H'|E)} = \frac{P(E|H)}{P(E|H')} \times \frac{P(H)}{P(H')}. \quad (4)$$

If the likelihood ratio is greater (lower) than one, the posterior odds will be greater (lower) than the prior odds of H . The likelihood ratio, then, is a measure of the upward or downward impact of the evidence on the prior odds of two hypotheses H and H' .

Moreover, the division of labour in legal fact-finding also seems to recommend LR over BF. A prominent forensic scientist recommends that ‘in criminal adjudication, the values of the prior odds and the posterior odds are matters for the judge and jury, in accordance with the normal division of labour in forensic fact-finding.’ (Aitken & Taroni, 2008, p. 194) and that the experts should ‘not trespass on the province of the jury by commenting directly on the accused’s guilt or innocence, . . . and should generally confine their testimony to presenting the likelihood of their evidence under competing propositions’ (Aitken, Roberts, & Jackson, 2010, p. 42). If however, the expert were to report BF, this would mean they are competent to estimate $P(E)$ directly, which is unlikely, or that they implicitly estimate it relying on their estimation of $P(H)$ in the background (if they use the law of total probability). In contrast, LR is one way to describe the value of evidence while abiding by these constraints, and it in fact is used in practice. Experts sometimes do testify by offering the likelihood ratio as a measure of the strength of the evidence. An expert, for instance, may testify that the blood-staining on the jacket of the defendant is ten times more likely to be seen if the wearer of the jacket hit the victim (prosecutor’s hypothesis) rather than if he did not (defense’s hypothesis) (Aitken et al., 2010, p. 38).

The idea that both conditional probabilities involved in likelihood ratio should be used in evidence strength evaluation applies generally to all forms of evidence, inclusive of DNA evidence, although it might not always make a practical difference. For suppose an expert testifies that the crime traces genetically match the defendant and that the **random match probability** is extremely low, say 1 in 100 million. Is the match strong evidence that the defendant is the source of the traces? The random match probability—often interpreted as the probability that someone who is not the source would coincidentally match, $P(\text{match}|\neg\text{source})$ —is a common measure of the strength of a DNA match. The lower this probability, the more strongly incriminating the match. This is sensible because a low random match probability suggests it is unlikely two people could share the same DNA profile. This is, however, also in agreement with the use of likelihood ratio in evidence evaluation, because $P(\text{match}|\text{source})$ is practically equal to one, so neglecting it in evidence strength reporting does not make any real difference. That $P(\text{match}|\neg\text{source})$ is low is in such contexts enough to ensure that the likelihood ratio is significantly above one. For practical purposes, then, a suitably low random match probability does capture the idea that the evidence is strongly incriminating evidence. The conceptual point still stands, though. If $P(\text{match}|\text{source})$ was significantly different from one, reporting only $P(\text{match}|\neg\text{source})$ would be misleading.

Our goal in this chapter is to reach a balanced view of the utility of LR. On one hand, we want to motivate its use and appreciate its utility. On the other, we want to point out reasons why LR might be misleading or unhelpful. In the next subsection, we will focus on a part of the positive task. A good example of the positive role of texts{LR} in theorizing about fact-finding is its use in studying the impact of false positive risk. We will take a look at this example now, in Section 2. Later on (Section A), we will take a more detailed look at another positive example, the use of LR in studying the value of the so-called cold-hit matches.²

Our goal in the remainder of this chapter is to take a closer look at an important tool of evidence evaluation proposed by legal probabilists: the likelihood ratio. In Section 1 we explain what this measure of evidential strength is, how it compares to another one, sometimes called the Bayes factor, and why it is preferable to it. In Section 2 we give an example of its utility. We explain the reasons to take the risk of false positives in DNA identification seriously and use likelihood ratio to illustrate the impact of such a risk on the value of DNA evidence. Another example, developed in Section 3, is meant to remove the impression that likelihood ratios are useful only for clearly quantitative evidence such as DNA evidence. We talk about the insights that probabilistic thinking can bring into eyewitness evidence evaluation and incorporation.

²In principle, it would be possible to go over analogous considerations in terms of BF, however, as we already argued, there are reasons to prefer the use LR, and calculations in terms of LR are simpler and assume that less information is available to the agent than those in terms of BF.

Marcello: moved the chapter structure here, deleted the introduction section, is this better now?

Likelihood ratio, while useful, is sometimes hard to interpret in practice, as it depends on the choice of the hypotheses involved. We discuss this problem in Section 4. One reason why many hypotheses can be considered in such contexts is that they admit different levels: we discuss this phenomenon in Section 5, illustrating it with the example of the so-called two-stain problem. The general lesson from these sections is that likelihood ratio, while a very useful measure of evidential value, when presented alone without proper attention paid to hypothesis formulation and the impact of the choice of the hypotheses on the likelihood ratio itself, might be misleading.

Next, in Section 6 we discuss another problem put forward against the utility of likelihood ratios: the problem of relevance. On this objection, likelihood ratio sometimes fails to identify all the relevant items of evidence. We explain what the objection is, and why it is misled. However, the discussion leads us to another reason why likelihood ratio when presented alone might be unhelpful. Indeed, it might be the case that a piece of evidence is irrelevant with respect to a certain choice of hypotheses, but relevant to another, and all of these hypotheses might play an important role in the fact-finding process. As relevance is to be established in connection with various elements of a complex case, instead of a single likelihood ratio, one should consider various likelihood ratios arising for various selections of hypotheses at play, and so a readable representation of such complexities involved in a given case would be useful for putting various likelihood ratios to a proper use.

To illustrate the utility of thinking in terms of likelihood ratios in theorizing about the value of evidence, we spend some time in the appendix to go over a debate about the value of cold-hit DNA matches in which various agencies studying or performing DNA evaluation have still failed to reach agreement (Appendix A), and argue that a proper use of likelihood ratios leads to a fairly clear resolution (Appendix B). Finally, a more philosophically minded reader, who might recall that there are quite a few probabilistic confirmation measures in the vicinity and might wonder why almost none of them were discussed in the chapter, in Appendix C we explain why we think these other measures are not fit for the particular purpose at hand: evidence evaluation in legal-fact finding.

Marcello: These first two sections are fine, more or less. They make a clear overall case in favor of LR. Perhaps the second section should come first and then the intro after. The second section motivates looking into LR more closely, which is what we do in the rest of the chapter. The bit about irrelevant conjunction and LR is not clearly explained and needs further elaboration.

Marcello: Revised, check if you're happy with this now.

2. The risk of false positive and its impact

One context in which probabilities are extensively used is the use of DNA evidence. In testifying about the DNA match at trial, experts often assess the probability that a random person, unrelated to the crime, would coincidentally match the crime stain profile (random match probability). RMP is often an impressively low number, say 1 in 100 million or lower. Usually, such a match is taken to constitute strong evidence against the defendant. We will illustrate the utility of likelihood ratios by using to explain how this apparent strength of DNA evidence can be mitigated by the probability of a false positive.

First, observe that while DNA evidence seems as scientific as it gets, the risk of a false positive is not negligible (Shaer, 2016). For instance, Houston Police Department Crime Laboratory, a large public forensic center in Texas, handles around 500 cases a year. In 2016, KHOU 11, a local television station, sent dozens of profiles processed by the lab to independent experts. The results were not optimistic: police technicians quite systematically misinterpreted samples.

One notorious case involving a false positive is that of Josiah Sutton (then 16) and Gregory Adams (then 19), who were arrested for a rape of a 41-year-old woman. The victim was abducted in a parking lot and assaulted in a driving car (Ford Expedition). A few days after the incident, the victim spotted Sutton and Adams walking down a street, flagged down a patrol car, and accused them of the assault. Both Sutton and Adams had alibis, neither of them matched the victim's original description of the perpetrators. Sutton and Adams agreed to a DNA test to clear their names. A Houston lab analyst Christy Kim compared their results with DNA obtained from a vaginal swab, which contained a mixture of genetic material from at least three contributors, including the victim herself. The lab report did not report a match for Adams, but concluded that Sutton's DNA was consistent with the mixture DNA. In result, in 1999, Sutton was sentenced to 25 years in prison. Later on, a re-examination by prof. William

Thompson, indicated that the three DNA profiles typed by Kim (two from blood, one from saliva) varied, despite reportedly coming from a single source. Moreover, Kim failed to report that the DNA from the semen found on the car seat did not match that of Sutton. In effect, the DNA evidence was reprocessed, no DNA match was found, and in 2003 Sutton was released from prison.³

This is only one example of quite a few cases of DNA matching going awry, and the existing anecdotal evidence suggests there are quite a few potential sources of error (see Thompson, 2013 for a more exhaustive treatment and multiple examples):

- **Cross-contamination of samples.** For instance, in Dwayne Johnson (2003) samples were accidentally swapped. In Lukis Anderson (2012) the material has been carried over by the responding paramedics. In one case, German police invested a considerable amount of time and effort searching for the so-called Phantom of Heilbronn, whose DNA profile was found on evidence from a large variety of crimes. A bounty of 300k EUR was placed on her head. It turned out she was an innocent employee involved in the production of cotton swabs used across the country.
- **Mislabeling of samples.** For instance, in 2011 the Las Vegas Metropolitan Police Department acknowledged that samples of two men suspected of a 2001 robbery were switched, leading to the exclusion of the perpetrator and four years of incarceration of the other suspect. The mistake came to light only because the perpetrator was later on arrested for an unrelated crime. In a high-profile case of a serial rapist, the notorious Night Stalker who committed more than 140 sexual assaults in London, the actual perpetrator came to the attention of the police quite soon, but a DNA test excluded him (falsely so, because the samples had been mistakenly switched), and so his spree continued for months.
- **Misinterpretation of test results.** While single-source sample comparison is not too prone to this sort of error, the interpretation of mixtures—which is usually what is needed in sexual assault cases—is quite complicated. Here is an illustration of this fact. Dror & Hampikian (2011) re-examined a 2002 Georgia rape trial in which two forensic scientists had concluded that the defendant could not be excluded as a contributor to the mixture of sperm from inside the victim (the defendant was found guilty). The evidence—DNA mixture and the DNA profiles of the victim and three suspects together those pieces of information that were highly relevant (such as the DNA amplification conditions) was sent to 17 lab technicians for examination. One of them agreed that the defendant could not be excluded as a contributor. Twelve considered the DNA exclusionary, and four found it inconclusive. If the quantity of DNA is limited, there is uncertainty about the number of contributors and about whether any alleles are missing, determining which alleles to assign to which contributor to some extent involves educated guesses on the part of the analysts. This suggests there is an element of subjectivity in mixed DNA interpretation.

Moreover, such errors are not easy to detect. Since DNA evidence carries so much weight in the fact-finders mind, it is very unusual to proceed with additional time- and cost-consuming DNA tests. It is also unusual that the suspect or their family can on their own afford further tests. For instance, an additional test exonerated Timothy Durham, sentenced to 3000 years for the rape of a young girl in Oklahoma City. So far there are two more cases known in the US where re-testing exonerated the accused: Josiah Sutton, whose case we already mentioned, and Gilbert Alejandro. Even more troubling is that errors from contamination or mislabeling of samples often cannot be detected with further DNA testing, because they will simply replicate the same misidentification. Sometimes, a lab discovers their own error and reports it, but this is a rather unlikely turn of events (Thompson, 2013).

DNA identification is to some extent prone to errors which are not measured by the random match probability, and no serious attempts to systematically quantify error rates in DNA testing have been made. Anecdotal reports about false matches suggest that errors take place more often than RMP would entail, but how often we should expect them remains unclear (Thompson, 2013). Regular proficiency tests used in accredited DNA laboratories involve comparison of samples from known sources, but they are criticized for being unrealistically easy (yet, it happens that analysts fail them). Sometimes, corrective action files are made available, and then they aren't too impressive. For instance, the Santa Clara County district attorney's crime laboratory between 2003 and 2007 caught 14 instances of evidence cross-contamination with staff DNA, three of contamination by unknown person, and six of DNA contamination from other samples, three cases of DNA sample switch, one mistake in which the analyst reported an incorrect result, and three errors in the computation of the statistics to be reported. Of

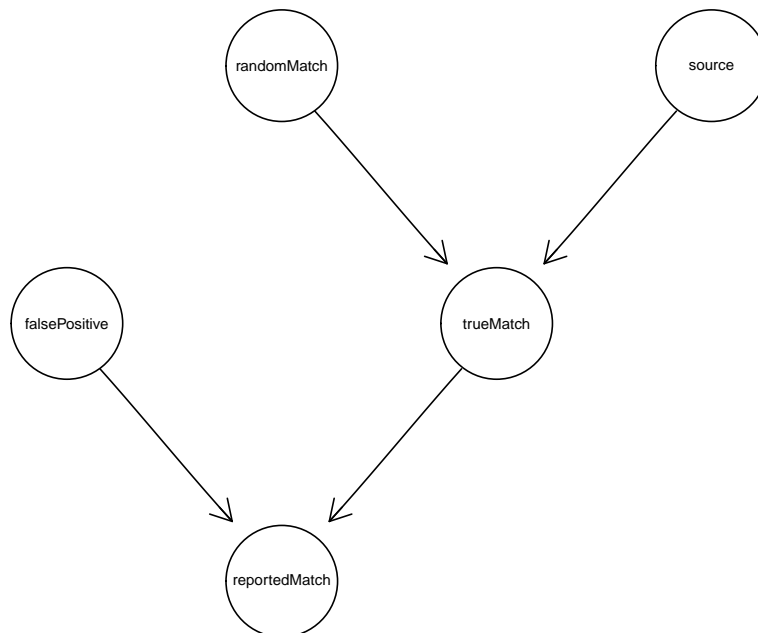
³Christy Kim later on sued her employer for her firing that resulted and won, her mistakes being attributed to systemic failures and inadequate supervision.

course, these are errors that were caught, and so one might argue that they show that labs are pretty good at catching their own errors. This, however, is an optimistic interpretation. These errors have been discovered due to unusual circumstances that led to the double-checking of the results. These circumstances, however, do not normally arise. It is not always the case that when a mistake is made the result implicates a staff member or an unknown person who was too young at the time of the crime to have committed it, for instance. Crucially, a match with a person whom the analyst might already know is a suspect is not an outcome that would raise an eyebrow and lead to a double-check.

Hopefully, having convinced the reader that the false positive probability is non-negligible, let us follow Aitken, Taroni, & Thompson (2003) in investigating its impact on the likelihood ratio of the DNA match. We just add a bit more details to the derivation they present for the sake of clarity. For simplicity, we still assume that the false negative probability is 0, that is, that if the match is real, it will be reported with certainty. We abbreviate:

- S The specimen comes from the suspect.
- R A match is reported.
- M There is a true match.

The difference between S and M is that a true match will exist not only if S true, but also if the profiles in fact are the same due to a random match. Similarly, a reported match might arise not only if there is a true match, but also if a false positive error has been made. The dependencies here are illustrated in Figure 2



Alicja: Something is off about this DAG representation, think about doing this better and more clearly. In terms of gammas?, random match is a conjunction, there is a dependence between source and random match...

Figure 2: Dependencies between variables in the false positive problem.

The formula we will end up with is:

$$LR(R, S, \neg S) = \frac{1}{RMP + [FPP \times (1 - RMP)]} \quad (\text{FPP-LR})$$

where RMP stands for the random match probability and FPP for the false positive probability. Re-written in terms of explicit probabilities the formula reads:

$$LR(R, S, \neg S) = \frac{1}{P(R|M)P(M|\neg S) + P(R|\neg M)P(\neg M|\neg S)}$$

The denominator mirrors the fact that there are two ways misleading evidence can arise: there is a true match and the suspect is not the source (so it's a random match), or there is no real match, and an error has been made in the identification process.

We will assume that whether a (lack of) match is reported is independent of whether it is coincidental,

$$\begin{aligned} P(R|M \wedge S) &= P(R|M \wedge \neg S) = P(R|M) \\ P(R|\neg M \wedge S) &= P(R|\neg M \wedge \neg S) = P(R|\neg M), \end{aligned} \quad (5)$$

that the probability of true match if the suspect is a source is 1,

$$P(M|S) = 1 \quad \text{so also} \quad P(\neg M|S) = 0, \quad (6)$$

and that the probability that a true match is reported,

$$P(R|M) = 1. \quad (7)$$

Here, for simplicity we take the probability of a false negative to be null; in fact, some of the reasons for taking false positives seriously are also reasons to take false negatives seriously, but let's deal with one problem at a time (and in the end, the impact of a false positive risk will be clear from the way the formula will be derived). Now, let us rewrite the numerator of the LR by extending the conversation, rewriting the probabilities of conjunctions in terms of conditional probability and simplifying:

$$\begin{aligned} P(R|S) &= \frac{P(R \wedge S)}{P(S)} \\ &= \frac{P(R \wedge M \wedge S) + P(R \wedge \neg M \wedge S)}{P(S)} \\ &= \frac{P(R|M \wedge S)P(M|S)P(S) + P(R|\neg M \wedge S)P(\neg M|S)P(S)}{P(S)} \\ &= P(R|M \wedge S)P(M|S) + P(R|\neg M \wedge S)P(\neg M|S) \end{aligned} \quad (8)$$

Analogously, we can rewrite the denominator:

$$P(R|\neg S) = P(R|M \wedge \neg S)P(M|\neg S) + P(R|\neg M \wedge \neg S)P(\neg M|\neg S) \quad (9)$$

Putting (8) and (9) together, we have that:

$$LR(R, S, \neg S) = \frac{P(R|M \wedge S)P(M|S) + P(R|\neg M \wedge S)P(\neg M|S)}{P(R|M \wedge \neg S)P(M|\neg S) + P(R|\neg M \wedge \neg S)P(\neg M|\neg S)} \quad (10)$$

Now, apply (5) in four places:

$$LR(R, S, \neg S) = \frac{P(R|M)P(M|S) + P(R|\neg M)P(\neg M|S)}{P(R|M)P(M|\neg S) + P(R|\neg M)P(\neg M|\neg S)} \quad (11)$$

Then, use (6) in the numerator:

$$LR(R, S, \neg S) = \frac{P(R|M) \times 1 + P(R|\neg M) \times 0}{P(R|M)P(M|\neg S) + P(R|\neg M)P(\neg M|\neg S)} \quad (12)$$

Finally, (7) yields:

$$LR(R, S, \neg S) = \frac{1}{P(R|M)P(M|\neg S) + P(R|\neg M)P(\neg M|\neg S)} \quad (13)$$

Once we abbreviate $P(M|\neg S)$ as RMP, $P(R|\neg M)$ as FPP and $P(\neg M|\neg S)$, we arrive at the desired formula.

In Figure 2 we illustrate this impact for the range of FPP between 0 and 0.05, for two values of RMP: 10^{-9} (often reported in the case of two single-source samples over ten or more loci) and 10^{-3} (sometimes obtained by means of less discriminating tests when the comparison involves a mixed sample). The upshot is that even a small increase in FPP can lower the likelihood ratio dramatically, which is yet another reason not to ignore FPP in DNA evidence evaluation. In our illustration, we look at two pieces of DNA evidence with the two RMP rates at $1 \cdot 10^8$ and 100, respectively. If the false positive probability reaches 0.02, they fall down to ≈ 49.99 and ≈ 33.55 , and they get fairly close to each other already at $FPP = 0.05$, where they are ≈ 20 and ≈ 16.8 .

Impact of false positive probability on likelihood ratio




Figure 3: Impact of the false positive probability on the likelihood ratio for two values of RMP. The horizontal reference line is at 16.8, the likelihood reached at $RMP=10^{-3}$ for $FPP=0.05$. At the same value of FPP, the LR for $RMP 10^{-9}$ is 20.

Interestingly, Buckleton, Bright, & Taylor (2018) give a seemingly different formula for the impact of errors on likelihood ratio. Since the derivation is simpler and it turns out that in fact this is simply a more general formula, of which (FPP-LR) is just a particular instance, it worth taking a look.

First, Buckleton et al. (2018) make the conceptual distinction between the probability that an error occurs (E) and the probability that a match is reported if it does. Mind your head: here E stands for error, not for evidence! In terms of our notation, we have:

$$e = P(E) = P(E|S) = P(E|\neg S)$$

That is, we denote the probability of error as e , and we assume it doesn't depend on whether the prosecution hypothesis is true (whether the suspect is the source).

Separately, the formula includes the probability of a reported match if an error occurs, also assumed to be independent of whether the prosecution hypothesis is true:

$$k = P(R|E, S) = P(R|E, \neg S)$$

M: I wonder if there is a deeper point here. That is, extremely low RMP values are meaningless when they cannot be accompanied by equally low FPP. See Alex's comment via email.

Marcello: Added a warning, I can rewrite the proof if you can think of a better abbreviation.

Further, it is assumed that the probability of false negatives is zero ($P(R|S, \neg E) = 1$) and the probability of reported match if no error occurs and the defense hypothesis is true is RMP ($P(R|\neg E, \neg S) = RMP$).
Now the derivation:

$$\begin{aligned}
 LR &= \frac{P(R|S)}{P(R|\neg S)} \\
 &= \frac{P(R|\neg E, S)P(\neg E|S) + P(R|E, S)P(E|S)}{P(R|\neg E, \neg S)P(\neg E|\neg S) + P(R|E, \neg S)P(E|\neg S)} \\
 &= \frac{1(1-e) + ke}{RMP(1-e) + ke} = \frac{1-e+ke}{RMP-e \times RMP + ke} \\
 &= \frac{1-(1-k)e}{RMP(1-e) + ke}
 \end{aligned}$$

Let's take it slow. First, the likelihood ratio is just the ratio of probabilities of a reported match (1) if the suspect is the source and (2) if the suspect is not the source.

Now, think of $P(R|S)$. This can be split into two possible scenarios: an error has not been made, or an error has been made. Accordingly, the numerator in the second line uses the law of total probability to split $P(R|S)$ into these two options.

Similarly, $P(R|\neg S)$ can be split into two cases: the suspect is not the source, but we are dealing with a random match, or the suspect is not the source, and an error has been made. An application of the law of total probability in the denominator mirrors this. The rest of the argument is just rewriting in terms of abbreviations, and algebraic manipulation.

Note now that if you think of an error as something that guarantees a mistaken identification, k becomes 1 and e becomes the false positive rate. On this assumption straightforward algebraic manipulation gives:

$$\begin{aligned}
 \frac{1-(1-k)e}{RMP(1-e) + ke} &= \frac{1-e+e}{RMP(1-e) + e} \\
 &= \frac{1}{RMP-e \times RMP + e} = \frac{1}{1+e(1-RMP)}
 \end{aligned}$$

which is the same as the formula obtained by Aitken et al. (2003) if we take e to be FPP, as we should on the assumption that $k = 1$.

An analogous reasoning can be used to study the impact of false negative probability on the value of exculpatory DNA evidence. Consider the probability of no match being reported if an error has been made, analogous to k above:

$$l = P(\neg R|E, S) = P(\neg R|E, \neg S) = P(\neg R|E)$$

Now, the likelihood ratio calculations, assuming $l = 1$, go as follows:

$$\begin{aligned}
 LR(\neg R, S, \neg S) &= \frac{P(\neg R|S)}{P(\neg R|\neg S)} \\
 &= \frac{P(\neg R|\neg E, S)P(\neg E|S) + P(\neg R|E, S)P(E|S)}{P(\neg R|\neg E, \neg S)P(\neg E|\neg S) + P(\neg R|E, \neg S)P(E|\neg S)} \\
 &= \frac{0(1-e) + le}{(1-RMP)(1-e) + le} \\
 &= \frac{le}{1-RMP-e+eRMP+le} = \frac{le}{1-RMP+e(l+RMP-1)} \\
 &= \frac{e}{1-RMP+eRMP} = \frac{e}{1+(e-1)RMP}
 \end{aligned}$$

Marcello: I added these three passages to explain the argument, is this clear now?

Marcello: Do you also want me to explain the manipulation here?

Alicja: make sure all the proofs are clear to you!

Now, this argument is very similar, do we really need an informal gloss here as well?

If the error rate is 0, then the numerator is 0 and so is the LR, as it should. In such a case, the evidence is completely exculpatory, the posterior probability that the suspect is the source will be also 0. If the error rate is not 0, the numerator simply is the probability of error, and the numerator takes values between $1 - RMP$ and 1, depending on the value of e . Quite crucially, 1 in the denominator is decreased by $(1 - e)RMP$, which with usually very low RMP in the case of DNA evidence is a very small change as compared to one, so the denominator stays very close to 1 even if e is very high, and the LR effectively simply is $\approx e/1 = e$. The lines in Figure 2, strictly speaking, do not overlap, but the difference between them (with RMP being fairly low) is negligible.




Figure 4: Impact of the false negative probability on the likelihood ratio of exculpatory evidence for two values of RMP.

Interestingly, there is a sense in which the situation is not symmetric when we compare FPP to FNP. With FPP, the LR is 50 when $e = 0.01$ for $RMP = 10^{-3}$, while for the same e and RMP it is 0.01 for FNP, an hundredfold decrease. This illustrates that the exculpatory value of DNA evidence is higher than its incriminating value, even if the error rates and random match probabilities are the same.

Some conceptual symmetry can be regained though. Suppose RMP is really low as compared to FPP and let's ignore it in our approximation. Then, the likelihood ratio of the incriminating evidence becomes $1/FPP$ and the likelihood ratio of exculpatory evidence becomes $FNP/1$. However, the change rate of these differ. While $d/dx(1/x) = d/dx(x^{-1}) = -1/x^2$, $d/dx(x/1) = 1$, and the derivatives look quite different (Figure 5)

The key lesson from this section is as follows. First, likelihood ratios are useful in the evaluation and comparison of the impact of positive and negative error rates: and it makes clear that this impact is not the same, contrary to what one may intuitively think. Second, the likelihood ratio analysis reveals that even a seemingly small error rate in a sense trumps random match probability: if you think there are good reasons to worry about random matches, the analysis shows that there are much better reasons to worry about error rates. While in fact error rates have not been properly studied, the above application of likelihood ratio to investigate the potential impact of hypothetical error rates is still useful in guiding further research, as it clearly suggests that error rates may have serious impact on the value of various types of evidence, and so require attention.

Marcello: Added some explanation of the apparent asymmetry does this help?

Marcello: Added this summary passage here, let me know what you think, including a bit on guiding further research




Figure 5: Derivatives of incriminatory and exculpatory likelihood ratios, range restricted to (-10,2).

3. Eyewitness identification and likelihood ratio

So far we paid attention to DNA evidence, which rather uncontroversially is quantitative. This might give the impression that likelihood ratio is useful only for thinking about clearly quantitative evidence of this sort. In this section, we discuss how likelihood ratio is still useful for reflecting on evidence which, at least seemingly, is not quantitative: eyewitness evidence and for combining various items of evidence.

We will argue that a quantitative perspective on eyewitness evidence is not only available, but also useful. First, it teaches us that intuitive evaluation of such evidence leads us astray more often than we tend to think. Second, it provides us with better tools of eyewitness evidence evaluation, as it allows us to study factors that impact its reliability. Next, we will sketch how such a quantitative perspective clears the path to a likelihood ratio treatment of such evidence: (i) in likelihood ratio evaluation of a stand-alone piece of eyewitness evidence, (ii) in combination of eyewitness evidence with a piece of quantitative incriminating evidence, and (iii) in adjudication when different pieces of evidence collide.

The perspective we take here is that there is no magical barrier between quantitative and qualitative evidence. A certain type of evidence can become numerical if sufficient amount of evidence about its reliability has been collected and statistically analyzed. Eyewitness testimony is not only no exception, but also a good example of this. We will now go over the main quantitative findings regarding the reliability of eyewitness evidence. Next, we will show how insights from such findings can be used in a likelihood ratio analysis of the impact of eyewitness evidence.

First of all, quantitative analyses might lead to a more sensible assessment of evidence than merely intuitive judgments. In the case of eyewitness testimony, this is crucial because eyewitness evidence tends to be overvalued, and it is the quantitative information that can and should be used to stop this madness. Field studies indicate filler identification rates of 20-24% (Klobuchar, Steblay, & Caligiuri, 2006) in eyewitness identifications. That is, around 20% of the time, an innocent presented to an eyewitness is going to be 'identified' by the eyewitness (the situation is a bit more complicated, read on for details).

To get a better perspective on the fallibility of eyewitness evidence, consider that 4.1% is a conservative estimate of false death sentence convictions in the United States, and those are based on much stronger and multiple pieces of evidence, not a single eyewitness' testimony (Gross, O'Brien, Hu, &

Kennedy, 2014). A study of 340 exonerations in years 1989-2003 indicates that around 90% of false convictions for rape (and pretty much all inter-racial rape mistaken convictions) are based on eyewitness misidentification, and that in 43% of false convictions for murder the defendant was misidentified by one or more eyewitnesses.

What else is quantitatively known about the reliability of eyewitness testimony? A study of line-ups in 1561 witnesses and 616 suspect in real cases in Greater London (Wright & McDaid, 1996) and of 689 identification attempts in 271 real identification cases in Sacramento (Behrman & Davey, 2001) suggest false positive rate of around 20%. Also in experimental setting (where the witnesses are less emotionally affected by the crime), eyewitnesses identify a filler in approximately twenty percent of all real criminal line-ups (Thompson, 2007).

Moreover, studies of cross-examination failed to show that it improves accuracy and that there is a clear relation between witness' confidence and accuracy. In a series of experiments subjects were asked to cross-examine eyewitnesses to determine whether witnesses made accurate or mistaken identifications. Subjects have shown little or no ability to make such discriminations (Wells & Olson, 2003). Another example of the unreliability of cross-examination is an experiment by Lindsay, Wells, & Rumpel (1981), in which a representative sample of 48 witnesses was cross-examined. Subjects ($n = 96$) viewing the cross-examinations showed no ability to distinguish accurate- from false-identification witnesses within conditions as measured by subjects' trust in witnesses.

Moreover, there are various factors that have impact on the reliability of eyewitness testimony, and here is where the quantitative analysis shines. Several of the eyewitness quality issues have been studied by means of experimental methods. Some of them are systemic variables (simultaneous/sequential lineups, showups,⁴ presence of prior identification, culprit present/culprit absent, frequency of witnesses per suspect), some of them—estimator variables – are not (the effect of delay, cross versus own-race effects, weapon focus effects, presence of violence)—see (Behrman & Davey, 2001) and (Wells & Olson, 2003) for an interesting discussion of such factors.

A fascinating meta-analysis by Wixted & Wells (2017) suggests an even more complicated picture, whose key points are as follows:

- In light of the empirical results that we've discussed and results similar to them the justice systems has grown more suspicious of eyewitness evidence over the last twenty years or so, especially doubting whether the eyewitness confidence is predictive of accuracy.
- Isolating cases in which the identification has been made *in pristine conditions* with high *initial* confidence focusing on cases which would go to trial, that is, in which the suspect indeed was identified with high confidence, Wixted & Wells (2017) argue that the data show that in such circumstances, the eyewitness confidence indeed is highly predictive of accuracy.
- Pristine conditions involve a double-blind experiment containing only one suspect, at least five fillers, with no fillers who don't look like the perpetrator at all, from among which the suspect doesn't stand out as obviously fitting the description which the eyewitness is familiar with, with no extreme coincidental resemblance of a filler to the suspect. The witness has to be cautioned that the offender might not be in the line-up and understand that they aren't failing if they don't indicate anyone, and the confidence statement needs to be collected at the time of first identification. Very few police departments are known to run their lineups in pristine conditions.
- Whether the conditions were pristine and whether the eyewitness initial confidence was high are factors which trump the role of the estimator variables.
- The extent to which initial high confidence in pristine conditions is predictive of identification accuracy depends on the base rate of target-present lineups. In lab studies it is usually fixed at around 50%, but it is expected to vary widely in real circumstances. The best estimate is around 35%, and it seems that at that base rate high-confidence witnesses (initial confidence, pristine conditions) are around 90% accurate.
- With time, and especially in the contexts in which a witness expects a cross-examination, witness confidence loses its predictive power. Preparations for the trial are known to inflate witness confidence, especially if they received a positive feedback from anyone after the initial identification.

Ideally, properly formatted data about all the factors we discussed could, with appropriate effort, be used to develop a multivariate model which would lead to quantitative eyewitness reliability estimates,

⁴Two basic types of identification procedures can be found in the literature, lineups, and field showups. A showup refers to the observation of a single suspect by a witness in the field, typically at the crime scene, whereas a lineup refers to the presentation of the suspect and several foils, either live or via photographs.

including some estimation of the uncertainty involved. While we are far from reaching the state of such maturity, it should be clear that even with the current state of knowledge, and expert in eyewitness testimony evaluation aware of the literature some of which we have cited, aware of the circumstances of the case and identification could express their evaluation of the witness' reliability quantitatively. Claiming that an untrained jury member can do better by just intuitively evaluating what the eyewitness said instead is at least hasty.

So consider two scenarios in which we numerically estimate the reliability of the eyewitness evidence in line with the research we just discussed. In one, such an expert recognizes the identification conditions as pristine and testifies: the probability of the testimony if the suspect in fact is the perpetrator ($P(E|H)$) is $.95 \pm .05$, and the false positive probability ($P(E|\neg H)$) is $.1 \pm .03$. In scenario two, the conditions were not pristine, and the expert testifies that $P(E|H) = .8 \pm .05$ and $P(E|\neg H) = .25 \pm .05$. We get two different *ranges* of likelihood ratios, lr_1 and lr_2 . In the first case, the minimum and the maximum are as follows:

$$\min(lr_1) = .95/.07 \approx 13.57 \quad \max(lr_1) = .85/.13 \approx 6.53$$

so the expert can, say, testify that the likelihood ratio is in the range of 6.5-13.5. In the second scenario, the minimum and the maximum are as follows:

$$\min(lr_1) = .85/.2 = 4.25 \quad \max(lr_1) = .75/.3 \approx 2.5$$

so the expert can testify that the likelihood ratio is in the range of 2.5-4.25.

Now, suppose further evidence is put forward to the effect that the suspect blood type matches the crime scene sample. Say the probability of a match if the suspect is the source is simply 1, while the probability of a random match is .05. The likelihood of this evidence alone is $1/.05 = 20$. Assuming independence, the joint likelihood for total evidence is obtained by multiplying separate likelihood ratios. Now, without any commitment to the priors, the impact of likelihood ratio ranges on any prior probability can be quantified and visualised, as in Figure 6.

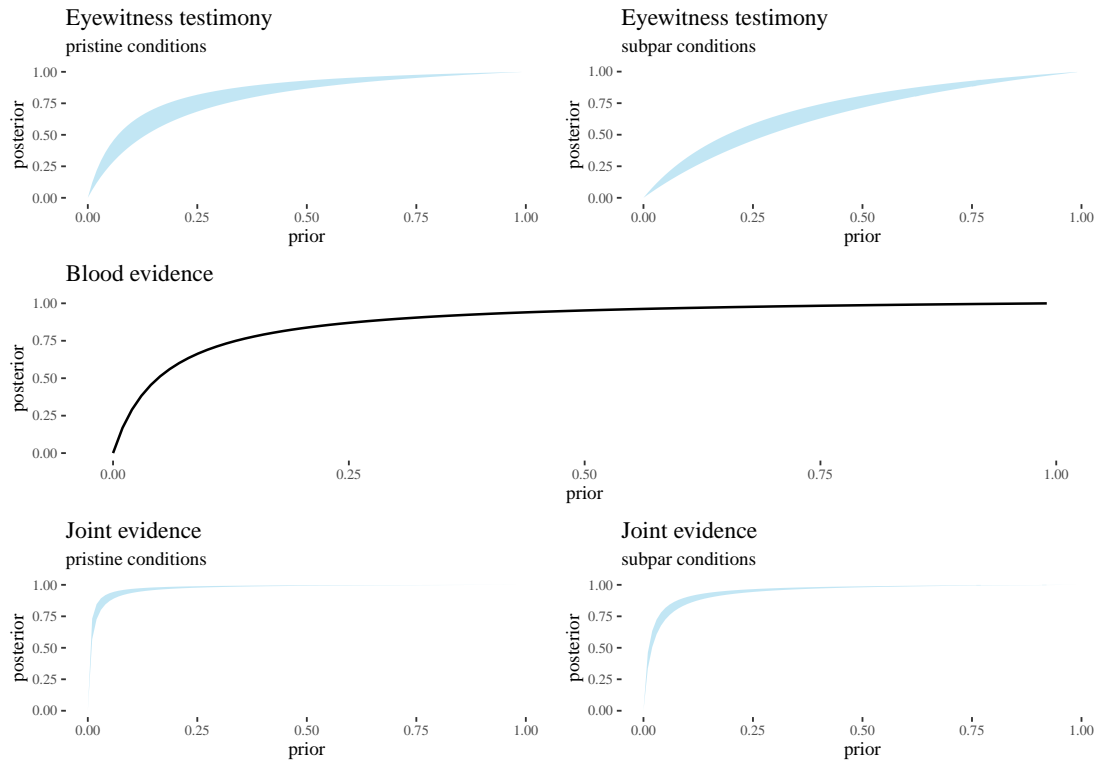


Figure 6: Impact of converging items of evidence on the posteriors.

What about conflicting evidence? Suppose this time the conditions were pristine, the expert's evaluation of the eyewitness evidence in pristine conditions is as before, but the witness identified someone else than the suspect (evidence E), while DNA evidence (evidence D) supports the prosecution hypothesis H with $LR = 300$ (this is, say, because we take the false positive probability seriously). Then, the range of plausible likelihood ratios for eyewitness evidence alone is:

$$\left[\frac{.07}{.95}, \frac{.13}{.85} \right] \approx [.073, .15]$$

In contrast, if the eyewitness is a friend of the suspect, so that you estimate $P(E|H)$ to be .6, while the conditions were suboptimal, so that $P(E|\neg H) = .8 \pm .05$, the likelihood ratio range is .7 – .8 and the impact of such eyewitness evidence is quite different.

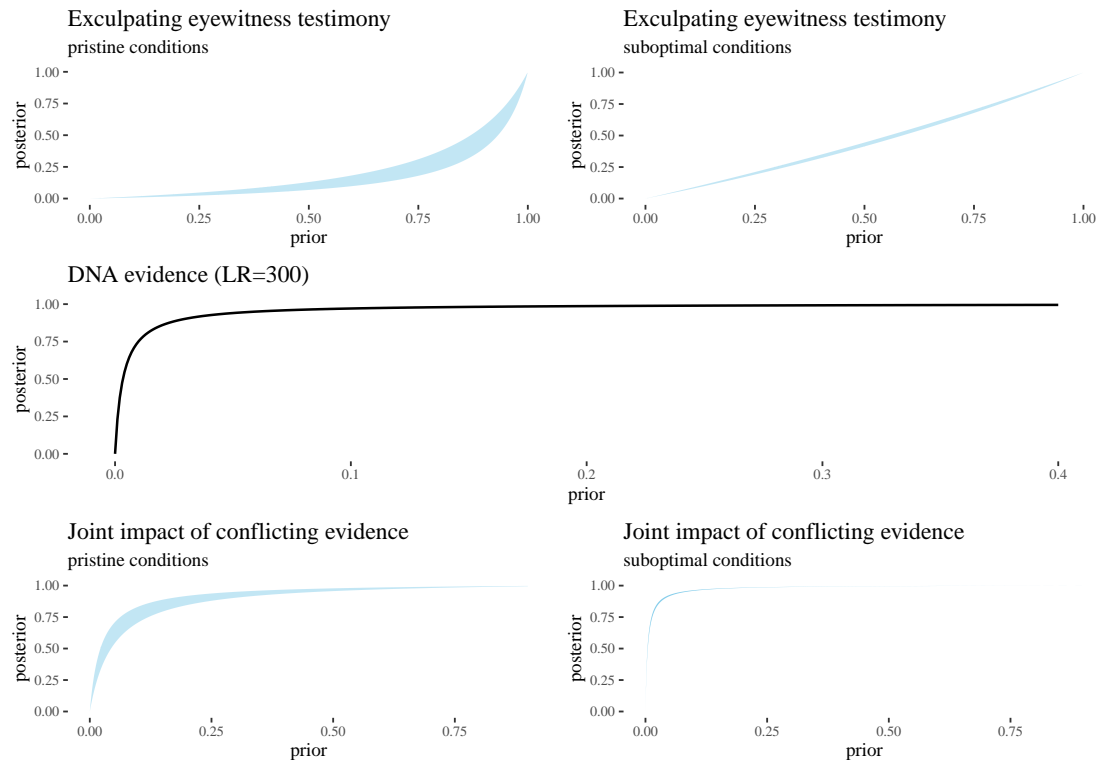


Figure 7: Impact of conflicting items of evidence on the posteriors.

The key observation here is that whether exculpatory eyewitness testimony can in fact be sufficient to acquit in light of incriminating DNA evidence depends on the particulars: both on the quality of the eyewitness testimony and on the realistic assessment of the likelihood ratio for the DNA evidence (incorporating the probability of false positives), and on what priors other pieces of evidence led us to so far prior to the evaluation of such items of evidence. The devil is in the detail, and it is not always the case that one should trump the other.

The reader might be disappointed: simply saying “it depends” does not seem to help much. However, keep in mind that there are well known items in the literature, on one hand, in psychology, illustrating that human jurors tend to value eyewitness evidence over statistical evidence (Niedermeier, Kerr, & Messeé, 1999; Wells, 1992; Wells & Olson, 2003), and, on the other hand, in epistemology, attempting to explain this intuition in a philosophically principled manner (Bolinger, 2020; Smith, 2017; Thomson, 1986). In line with a more balanced view (Redmayne, 2008), our approach suggests that while the psychological effect exists, there are multiple cases in which the preference for eyewitness evidence is mistaken, and relegates the identification and evaluation of factors that contribute to its proper evaluation to actual empirical research. If the situation is complicated, the best that can be done is take its complexity seriously, and quantitative methods are much more fit for this task than juror’s intuitions or general philosophical statements.

Marcello: We might need more references in this passage, think about this.

4. Hypothesis choice

The point so far is that the likelihood ratio is a fruitful conceptual framework for assessing the strength of the evidence. However, its use is not devoid of challenges, and we move to a discussion of how they arise.

Marcello: revised these two intro passages, check

Now, we turn to a difficulty with the use of likelihood ratio in the context of legal fact-finding: likelihood ratios is sensitive to the hypothesis choice. We discuss this problem in general in this section, (i) showing how an *ad hoc* hypothesis can always lead to a useless likelihood ratio, (ii) discussing a real case in which the choice of hypothesis was fairly confusing and had impact on evidence evaluation, (iii) indicating that the problem doesn't arise if the hypotheses are exclusive and exhaustive, and (iv) arguing that this theoretical resolution may be unhelpful in practice, as there seem often to be good reasons why experts do not use such pairs of hypotheses. Further, in Section 5 we further illustrate how large shifts in the likelihood ratio can take place even if the hypotheses chosen seem reasonable, and argue that another clear source of variability of likelihood ratio is the choice of hypothesis levels. The two sections taken together paint a difficulty with using likelihood ratio alone as a measure of evidential strength and suggest that a more elaborate framework, employing likelihood ratios, but not restricted to their reporting, would be more useful.

One major difficulty, however, is the choice of the hypotheses H and H' that should be compared. Generally speaking, the hypotheses should in some sense compete with one another—say, in a criminal trial, H is the hypothesis put forward by the prosecution and H' is the hypothesis put forward by the defense. Presumably, the two hypotheses should be something that the two parties disagree about. But this minimal constraint offers too little guidance and leaves open the possibility for manipulations and misinterpretations of the evidence. What follows outlines some of the main arguments in the literature on this topic.

Consider a stylized DNA evidence case. Suppose the prosecutor puts forward the hypothesis that the suspect left the traces found at the crime scene. This hypothesis is well supported by laboratory analyses showing that the defendant genetically matches the traces.

The defense, however, responds by putting forward the following *ad hoc* hypothesis: 'The crime stain was left by some unknown person who happened to have the same genotype as the suspect.' Since the probability of the DNA match given either hypothesis is 1, the likelihood ratio equals 1 (Evelt, Jackson, & Lambert, 2000). The problem generalizes. For any item of evidence and any given prosecutor's hypothesis H , there is an *ad hoc* competing hypothesis H^* such that $P(E|H)/P(E|H^*) = 1$.

Hypothesis H^* is simply a just-so hypothesis, one that is selected only because it explains the evidence just as well as hypothesis H does (Mayo, 2018). If no further constraints are placed on the choice of the competing hypotheses—it would seem—no evidence could ever incriminate a defendant. This is unsettling.

One reply might be that this phenomenon need not be so damning in practice. Judges and jurors, after all, will often recognize *ad hoc* hypotheses for what they are—artificial theories that should not be taken seriously. Perhaps, the reasonable expectations of the participants in a trial will suffice to constrain the choice of hypotheses in just the right way.

One issue with this reply is that it is not principled. Fair enough, maybe in practice fact-finders will avoid *ad hoc* hypotheses. But why should they avoid them, other than because otherwise the resulting likelihood ratio will be useless? Moreover, as we will soon see, hypothesis choice has impact on likelihood ratio even if no *ad hoc* hypothesis seems to be used. Moreover, real cases tend to be quite complex, and it is not always obvious whether a certain choice of competing hypotheses, none of which are obviously *ad hoc*, is legitimate or not. In what follows we will illustrate some of these issues, and argue that due to this sensitivity, likelihood ratio alone is not informative enough, if not supplied by careful considerations of hypothesis choice.

Here is an example that illustrates how even when the competing hypotheses are not obviously *ad hoc*, the absence of a clear rationale for their choice may create confusions in the assessment of the evidence. In *R. v. Barry George* (2007 EWCA Crim 2722). Barry George was accused of murdering TV celebrity Jill Dando. A key piece of evidence at play was:

- E* A single particle of firearm residue (FDR) was found one year later in George's coat pocket and it matched the residue from the crime scene. This was the key incriminating evidence against him.

The defense argued that, since it was only one particle, there must have been contamination. The experts for the prosecution, however, testified that it was not unusual that a single particle would be found on the person who fired the gun. George was convicted, and his first appeal was unsuccessful.

After the first appeal, Dr. Evett from the Forensic Science Service worried that the evidence had not been properly assessed at trial. The jurors were presented with the conditional probability $P(\text{residue}|H_d)$ of finding the firearm residue in George's coat given the defense hypothesis H_d that George *did not* fire the gun. This probability was estimated to be quite low, indicating that the evidence spoke against the defense's hypothesis. But the jurors were not presented with the conditional probability $P(\text{residue}|H_p)$ of finding the same evidence given the prosecutor's hypothesis H_p that George *did* fire the gun that shot Dando. An expert witness, Mr. Keeley, was asked to provide both conditional probabilities and estimated them to be $1/100$, which indicated that the firearm residue had no probative value. After new guidelines for reporting low level FDR in 2006, the FSS re-assessed the evidence and concluded that it was irrelevant. George appealed again in 2007, and relying on Keeley's estimates, won the appeal.

At first, this case seems a good illustration of how likelihood ratios help to correctly assess the value of the evidence presented at trial. But this reading of the case would be overly optimistic. In fact, a close study of the trial transcript shows that Keeley's choice of hypotheses was not systematic and the likelihood ratio based on them was therefore really hard to interpret (Fenton, Berger, Lagnado, Neil, & Hsu, 2014). For instance, Mr Keeley is reported to have said:

It was necessary to balance the likelihood that the particle came from a gun fired by the appellant and the likelihood that it came from some other source. Both were unlikely but both were possible.

Keeley compared the hypothesis that the particle found in George's pocket came from a gun fired by George himself, and the alternative hypothesis that the particle came from another source. In line with the quotation, Keeley said that the prior probabilities of both hypotheses should be low. But this is mathematically impossible if they were exhaustive and exclusive.

On another occasion, Keeley took the prosecutor's hypothesis to be 'The particle found in George's pocket came from the gun that killed Dando' and the defense hypothesis to be 'The particle on George's pocket was inserted by contamination.' The problem is that the evidence is a logical consequence of either of them, so the conditional probability of the evidence given each of these hypothesis is one. Crucially, they are therefore useless for the evaluation of the weight of evidence, because in such case the likelihood ratio will always be one for trivial reasons. The most charitable reading of the trial transcript suggests that the expert had in mind the hypotheses 'George was the man who shot Dando' and 'The integrity of George's coat was corrupted.' But these hypotheses are neither exhaustive nor exclusive, and Keeley gave no clear criterion for why these hypotheses should be compared in the likelihood ratio (see Fenton et al., 2014 for further details).

The confusion in the Barry George case is attributable to the absence of clear rules for choosing the hypotheses in the likelihood ratio. One such rule could be: pick competing hypotheses that are exclusive (they cannot be both true) and exhaustive (they cannot be both false). In this way, the parties would not be able to pick *ad hoc* hypotheses and skew the assessment of the evidence in their own favor.

Besides blocking partisan interpretations of the evidence, there are other principled reasons to follow the exclusive-and-exhaustive rule, specifically, the fact that when the hypotheses are not exclusive or exhaustive, the likelihood ratio might deliver counterintuitive results and cause confusion in the assessment of the strength of the evidence. If two competing hypotheses, H_p and H_d are not mutually exclusive, it is possible that they both make the evidence equally likely (the likelihood ratio is one), and yet the posterior probabilities of the hypotheses given the evidence are higher than their prior probabilities.

For instance, let H_p stand for 'The defendant is guilty' and H_d for 'The defendant was not at the crime scene'. Both hypotheses might be true. Let E stand for 'Ten minutes before the crime took place the defendant—seen at a different location—was overheard on the phone saying *go ahead and kill him*.' It is conceivable that the likelihood ratio should equal one in this context, yet the posterior probabilities of each hypothesis, given E , should be higher than the prior probability. So, intuitively, the evidence should positively support each hypothesis, contrary to what the likelihood ratio would suggest.

Further, when the two competing hypotheses are not exhaustive, the likelihood ratio may once again

clash with our intuitions. The likelihood ratio might then equal one even though the evidence lowers their posterior probability. For example, suppose Fred and Bill attempted to rob a man. The victim resisted, was struck on the head and died. Say H_p stand for 'Fred struck the fatal blow' and H_d stand for 'Bill struck the fatal blow.' The hypotheses are not exhaustive. A missing hypothesis is 'The man did not die from the blow.' Suppose E is the information that the victim had a heart attack six months earlier. The likelihood ratio $P(E|H_p)/P(E|H_d)$ equals one since $P(E|H_p) = P(E|H_d)$. Yet E reduces the probability of both H_p and H_d . So, in this case, the evidence should negatively support each hypothesis, contrary to what the likelihood ratio suggests.

Whether the exhaustive-and-exclusive rule would be a good guiding principle, however, is not clear-cut. Requiring that the hypotheses be always exclusive and exhaustive hypotheses is not without complications either. For consider an expert who decides to formulate the defense hypothesis by negating the prosecution hypothesis, say, 'the defendant did not hit the victim in the head.' This choice of defense hypothesis can be unhelpful in assessing the evidence, because the required probabilities are hard to estimate. For instance, what is the probability that the suspect would carry such and such blood stain if he did not hit the victim in the head? This depends on whether he was present at the scene, what he was doing at the time and many other circumstances.

(Reader warning: this passage will discuss hypothesis choice in a rape case.) Similarly, in a rape case, it is hard to estimate the probability of the matching evidence if the suspect did not have the intercourse with the victim. Instead, what is considered is the hypothesis that someone else, unrelated to the suspect, had intercourse with the victim. As (Evet et al., 2000) point out, in many real life rape cases the choice of a particular hypothesis to be used by the expert in the evaluation of the strength of the evidence (of, say, the lack of semen in a rape case), will depend on contextual factors. Sometimes it will be 'intercourse did not take place,' sometimes it will be 'the intercourse took place, but the complainant used a vagina douche,' or sometimes 'another sexual act took place'. More often than not, the hypotheses chosen will not be mutually exclusive.

Moreover, comparing exclusive and exhaustive hypotheses can also be unhelpful for jurors or judges making a decision at trial. In a paternity case, for example, the expert should not compare the hypotheses 'The accused is the father of the child' and its negation, but rather, 'The accused is the father of the child' and 'The father of the child is a man unrelated to the putative father' (Biedermann, Hicks, Taroni, Champod, & Aitken, 2014). The choice of the latter pair of competing hypotheses is preferable. Even though the relatives of the accused are potential fathers, considering such a far-fetched possibility would make the assessment of the evidence more difficult than needed. At the same time, if the defense hypothesis is too specific, *ad hoc* and entails the evidence, it won't be of much use. For example, take 'The crime stain was left by some unknown person who happened to have the same genotype as the suspect.' The probability of a DNA match given this hypothesis would be 1. But usually the probability of the DNA match given the prosecution's hypothesis, say 'The crime stain was left by the suspect,' is also 1. This would result in a rather uninformative likelihood ratio of 1. Another feature of such specific explanations is that it's hard to reasonably estimate their prior probability, and so hard to use them in arguments between opposing sides. (Evet et al., 2000).

So, it seems, the choice of competing hypotheses lies between two extremes. Exclusive and exhaustive hypotheses guard against *ad hoc* hypotheses, arbitrary comparisons and ensure a more objective assessment of the evidence. Unfortunately, exhaustive and exclusive hypothesis cover the entire space of possibilities, and sifting through this space may be cognitively unfeasible. So, in this respect, comparing more circumscribed hypotheses is preferable. The danger of doing so, however, is slipping into arbitrariness as likelihood ratios heavily depend on the hypotheses that are compared. The more latitude in the choice of the hypotheses, the more variable the likelihood ratio as a measure of evidentiary value. This is a particularly troubling phenomenon, as competing hypotheses can concern any factual dispute, from minute details such as whether the cloth used to suffocate the victim was red or blue, to ultimate questions such as whether the defendant stabbed the victim.

To add another complication, the likelihood ratio varies across hypotheses formulated at different levels of granularity: offense, activity and source level hypotheses. It is even possible that, at the source level, the likelihood ratio favors one side, say the prosecution, but at the offence level, the likelihood ratio favors the other side, say the defense, even though the hypotheses at the two levels are quite similar. Further, a likelihood ratio that equals 1 when source level hypotheses are compared may tip in favor of one side or the other when offence level hypotheses are compared (Fenton et al., 2014). This variability

Marcello: You keep mentioning this dispute between Taroni research group and Fenton research group about LR, but I never got the impression of deep disagreement. Maybe I missed something, do you have any specific papers in mind?

makes the likelihood ratio a seemingly arbitrary—and easily manipulable—measure of evidentiary value.

The likelihood ratio can be misleading, but this risk is mitigated when its assessment is accompanied by a careful discussion of a number of issues, such as: which hypotheses are being compared; how they are formulated; their level of granularity (that is, source, activity and offense level); why the hypotheses are (or are not) exclusive and exhaustive, and why other hypotheses are ruled out as unworthy of consideration.

This is our first warning sign when it comes to relying on likelihood ratios. In real-life applications, multiple hypotheses and pairs thereof are at play, and their choice does matter. For this reason, merely reporting likelihood ratios might be misleading,⁵ and such a presentation of evidence is best accompanied by a careful discussion of which hypotheses were considered, which were chosen and why, and how they all are related to the evidence. Later on we will argue that a clear treatment of such issues is immensely assisted by the use of Bayesian networks in evidence presentation and evaluation. This will be a recurring point in this book.

Marcello: Added this summary passage, check.

5. Levels of Hypotheses and the two-stain problem

Since we want to put likelihood ratios in their proper place when it comes to evidence evaluation and reporting, we would like to obtain a deeper understanding of the factors that might lead to their variation depending on the hypothesis choice. let's take a closer look at one systematic dimension of hypothesis choice sensitivity: likelihood ratios change with the levels of hypothesis under consideration. In this section we take a closer look at this issue.

Marcello: added intro to improve flow

To some approximation, hypotheses can be divided into three levels: offence, activity, and source level hypotheses. At the offence level, the issue is one of guilt or innocence, as in the statement 'Smith intentionally attacked the victim with a knife'. At the activity level, hypotheses do not include information about intent but simply describe what happened and what those involved did or did not do. An example of activity level hypothesis is 'Smith bled at the scene.' Finally, source level hypotheses describe the source of the traces, such as 'Smith left the stains at the crime scene,' without specifying how the traces got there. Overlooking differences in hypothesis level can lead to serious confusions. To illustrate, consider a case in which a DNA match is the primary incriminating evidence. In testifying about the DNA match at trial, experts will often assess the probability that a random person, unrelated to the crime, would coincidentally match the crime stain profile⁶ The random match probability is often an impressively low number, say 1 in 100 million or lower, at least excluding the possibility that relatives or identical twins would coincidentally match (Donnelly, 1995). This should count as a strong evidence against the suspect. But how exactly? RMP—the probability that a random person from the population matches the crime stain profile—is taken to be the probability that the suspect is a match if in fact he is innocent and is usually estimated as the frequency of a given profile in the relevant population. As we already know from the chapter in which we discussed probabilistic fallacies, RMP is not the posterior probability of innocence, since $P(\text{match}|\text{innocence})$ should not be confused with $P(\text{innocence}|\text{match})$. To confuse the two would be to commit the prosecutor's fallacy. Further, it is tempting to equate the random match probability to $P(\text{match}|\text{innocence})$ and together with the prior $P(\text{innocence})$ use Bayes' theorem to calculate the posterior probability of innocence $P(\text{innocence}|\text{match})$. But this also might be a mistake. Equating the random match probability with $P(\text{match}|\text{innocence})$ overlooks the difference between offense, activity and source level hypothesis. It is hasty to assume that, in one way or another, a DNA match can speak directly to the question of guilt or innocence. Even if the suspect actually left the genetic material at the scene—source level proposition—the match does not establish guilt. Even if the defendant did visit the scene and came into contact with the victim, it does not follow that he committed the crime he was accused of. It is true, that in many circumstances the random match probability and the posterior probability of innocence given a match would both be very low, but such issues need to be considered and took into consideration in DNA evidence evaluation.

Few forms of evidence can speak directly to offense level hypotheses. Circumstantial evidence that is more amenable to a probabilistic quantification, such as DNA matches and other trace evidence, does

⁵Similarly, BF would be sensitive to the choice of the prosecution hypothesis, but since BF faces other problems we already discussed, we will not get into this issue.

⁶For a survey of developments and complications of this model, see (Foreman et al., 2003).

not. Eyewitness testimony may speak more directly to offense level hypotheses, but it is also less easily amenable to a probabilistic quantification. This makes it difficult to assign probabilities to offense level hypotheses. Experts are usually not supposed to comment directly on offense level hypotheses, but they often comment on activity level and source level hypotheses. In moving from source to activity level, however, additional sources of uncertainty come into play. The assessment of activity level hypotheses depends on additional variables other than those on which the assessment of source level hypotheses depends. For example, the probability of finding such and such quantity of matching glass if the suspect smashed the window depends on how the window was smashed, when it was smashed, and what the suspect did after the action. Another problem arises due to recent improvements in DNA profiling technology. Since today investigators are able to obtain profiles from minimal amounts of genetic material, transfer probabilities become more difficult to assess as more opportunities of transfer arise. If small traces such as dust speckles can be used as evidence, the possibility that the traces were brought to the scene accidentally becomes more likely. For this reason, moving beyond source level hypotheses requires a close collaboration between scientists, investigators and attorneys (see Cook, Evett, Jackson, & Jones, 1998 for a discussion). The choice and formulation of the hypotheses are up for revision as new evidence is obtained or facts about what happened are accepted (Evett et al., 2000). A case study that further illustrates both advantages and limitations of the likelihood ratio as a measure of evidentiary strength is the two-stain problem, originally formulated by Evett (1987). The key limitation is due to the combination of two circumstances: first, that likelihood ratios vary depending on the choice of hypotheses being compared; second, that it is not always clear which hypotheses should be compared. To illustrate what is at stake, what follows begins with Evett's original version of the two-stain problem (which does not pose any challenge to the likelihood ratio) and then turns to a more complex version (which suggests that likelihood ratios, in and of themselves, are insufficiently informative). Suppose two stains from two different sources were left at the crime scene, and the suspect's blood matches one of them. More precisely, the two items of evidence are as follows:

- E_1 The blood stains at the crime scene are of types γ_1 and γ_2 of estimated frequencies q_1 and q_2 respectively.
- E_2 The suspect's blood type is γ_1 .

Let the first hypothesis be that the suspect was one of the two men who committed the crime and the second hypothesis the negation of the first.

- H_p The suspect was one of the two men who committed the crime.
- H_d The suspect was not one of the two men who committed the crime.

Evett (1987) shows that the likelihood ratio of the match relative to these two hypotheses is $1/2q_1$ where q_1 is the estimated frequency of the characteristics of the first stain. Surprisingly, the likelihood ratio does not depend on the frequency associated with the second stain. To understand Evett's argument, consider first the likelihood ratio:

$$\frac{P(E_1 \wedge E_2 | H_p)}{P(E_1 \wedge E_2 | H_d)} = \frac{P(E_1 | E_2 \wedge H_p)}{P(E_1 | E_2 \wedge H_d)} \times \frac{P(E_2 | H_p)}{P(E_2 | H_d)}.$$

Notice that the suspect's blood type as reported in E_2 is independent of whether or not he participated in the crime, that is, $P(E_2 | H_p) = P(E_2 | H_d)$. So the likelihood reduces to:

$$\frac{P(E_1 \wedge E_2 | H_p)}{P(E_1 \wedge E_2 | H_d)} = \frac{P(E_1 | E_2 \wedge H_p)}{P(E_1 | E_2 \wedge H_d)}.$$

The numerator $P(E_1 | E_2 \wedge H_p)$ is the probability that one of the stains is γ_1 and the other γ_2 given that the suspect is guilty and has profile γ_1 . The probability that one of the stains is γ_1 is simply 1, and assuming blood type does not affect someone's propensity to commit a crime, the probability that the second stain is γ_2 equals its relative frequency in the population, q_2 . So the numerator is $1 \times q_2 = q_2$. % Next, consider the denominator $P(E_1 | E_2 \wedge H_d)$. If H_d is true, the fact that the suspect has profile γ_1 is irrelevant for the crime scene profiles. the crime was committed by two randomly selected men with profiles γ_1 and γ_2 , who can be seen as two random samples from the general population as far as their blood profiles are concerned. There are two ways of picking two men with such profiles (γ_1, γ_2

and γ_2, γ_1), each having probability $q_1 q_2$. So the denominator equals $2q_1 q$. By putting numerator and denominator together, we have:

$$\frac{q_2}{2q_1 q_2} = \frac{1}{2q_1}.$$

which completes the argument. In general, if there are n bloodstains of different phenotypes, the likelihood ratio is $1/nq_1$, or in other words, the likelihood ratio depends on the number of stains but not on the frequency of the other characteristics.

Consider now a more complex two-stain scenario. Suppose a crime was committed by two people, who left two stains at the crime scene: one on a pillow and another on a sheet. John Smith, who was arrested for a different reason, genetically matches the DNA on the pillow, but not the one on the sheet. What likelihood ratio should we assign to the DNA match in question? Meester & Sjerps (2004a) argue that there are three plausible pairs of hypotheses associated with numerically different likelihood ratios (see their paper for the derivations). The three options are listed below, where R is the random match probability of Smith's genetic profile and δ the prior probability that Smith was one of the crime scene donors.

H_p	H_d	LR
Smith was one of the crime scene donors.	Smith was not one of the crime scene donors.	$R/2$
Smith was the pillow stain donor.	Smith was not one of the crime scene donors.	R
Smith was the pillow stain donor.	Smith was not the pillow stain donor.	$R(2-\delta)/2(1-\delta)$

Two facts are worth noting here. First, even though the likelihood ratios associated with the hypotheses in the table above are numerically different, the hypotheses are in fact equivalent conditional on the evidence. After all, Smith was one of the crime scene donors just in case he was the pillow stain donor, because he is excluded as the stain sheet donor. Smith was not one of the crime scene donors just in case he was not the pillow stain donor, because he is excluded as the sheet stain donor. Second, the example illustrates that sometimes the likelihood ratio is sensitive to the prior probability (after all, δ occurs in the third likelihood ratio in the table).

In addition, even though the likelihood ratios are numerically different, their posterior probabilities given the evidence are the same.

To see why, note that the prior odds of the three H_p 's in the table should be written in terms of δ . Following Meester & Sjerps (2004a), the prior odds of the first hypothesis in the table are $\delta/(1-\delta)$. The prior odds of the second hypothesis are $(\delta/2)/(1-\delta)$. The prior odds of the third hypothesis are $(\delta/2)/(1-(\delta/2))$. In each case, the posterior odds — the result of multiplying the prior odds by the likelihood ratio — are the same: $R \times \delta/2(1-\delta)$. So despite differences in the likelihood ratio, the posterior odds of equivalent hypotheses are the same so long as the priors are appropriately related (this point holds generally).

Dawid (2004) cautions that the equivalence of hypotheses, conditional on the evidence, does not imply that they can all be presented in court. He argues that the only natural hypothesis for the two-stain problem is that Smith is guilty as charged. Meester & Sjerps (2004b) reply that focusing on the guilt hypothesis is beyond the competence of expert witnesses who should rather select pairs of hypotheses on which they are competent to comment. Some such pairs of hypotheses, however, will not be exclusive and exhaustive. When this happens, as seen earlier, the selection of hypotheses is prone to arbitrariness. To avoid this problem, Meester & Sjerps (2004a) recommend that the likelihood ratio should be accompanied by a tabular account of how a choice of prior odds (or prior probabilities) will impact the posterior odds, for a sensible range of priors

In this way, the impact of the likelihood ratio is made clear, no matter the hypotheses chosen. This strategy concedes that likelihood ratios, in and of themselves, are insufficiently informative, and that they should be combined with other information, such as a range of priors, to allow for an adequate assessment of the evidence.⁷

commented out a crossref, not sure if it is still needed

The sensitivity of the likelihood ratio to the choice of hypotheses is not confined to the two-stain problem or alike scenarios. Recall our discussion of DNA matches in cold-hit cases. When the suspect is identified through a database search of different profiles, Taroni, Biedermann, Bozza, Garbolino, & Aitken (2014) and Balding & Donnelly (1996) have argued that the likelihood ratio of the match—which

⁷The reference class problem is lurking in the background. Balding (2004) argues that, in order to calculate the probability of a match given the evidence, the class of possible culprits should be identified, and different choices of such a class might lead to different likelihood ratios. On the problem of priors see. On the reference class problem, see Chapter XXX.

usually equals $1/\gamma$ where γ is the random match probability—should be adjusted by the database search ratio. This proposal tacitly assumes that the hypothesis of interest is something like ‘the defendant is the true source of the crime traces.’ This assumption is eminently plausible but not uncontroversial.

The National Research Council (NRC II) recommended in 1996 that the likelihood ratio of the match $1/\gamma$ be divided by the size of the database. In defending this proposal, Stockmarr (1999) argues the likelihood ratio of the match in cold-hit cases should be divided by the size of the database. Stockmarr believes we should evaluate the likelihood ratio using hypotheses that can be formulated prior to the database search, such as ‘The true source of the crime traces is among the suspects in the database,’ while others insist on using ‘The defendant is the true source of the crime traces.’ Now, interestingly, while these approaches lead to different LR evaluations, they are equivalent conditional on the evidence: given the same evidence, they read to the same posterior. This is another example of how likelihood ratios on their own might be insufficiently informative to allow for an adequate assessment of the evidence. We will discuss these issues in more depth later on. First, we will look at an argument against likelihood ratio being useful as a measure of evidential relevance, as dealing with it is conceptually more straightforward.

This section emphasizes the warning message we left you with at the end of the previous section: even if likelihood ratios are a legitimate tool of evidence evaluation, they should be used with care, and accompanied with an explicit discussion how they fit into a larger picture of various items of evidence and hypotheses at play. Again, we will argue that a prominent way to do so is to employ Bayesian networks in evidence presentation and evaluation.

Marcello: added this summary passage

6. Relevance and the small-town murder scenario

Another objection to likelihood ratios is that they might suggest evidence is irrelevant even if reported with respect to a pair of exclusive and exhaustive hypotheses. The objection indeed sometimes applies when the key hypotheses are the main guilt hypothesis and its negation. But we will argue that this is not an objection to the use of likelihood ratios, but rather against reporting them in isolation. When justice is done to the complexities at play, appropriate likelihood ratios still evaluate such evidence as relevant. Again, the point is that there is no escape from using a more elaborate evidence presentation tool than just reporting likelihood ratios in isolation.

Marcello: added this intro, check

The U.S. Federal Rules of Evidence define relevant evidence as one that has ‘any tendency to make the existence of any fact that is of consequence to the determination of the action more probable or less probable than it would be without the evidence’ (rule 401). This definition is formulated in a probabilistic language. Legal probabilists interpret it by relying on the likelihood ratio, a standard probabilistic measure of evidential relevance [Lempert (1977); Lyon 1996 relevance; aitken 2004 statistics; aitken 2010 fundamentals; Sullivan 2016 Likelihood Story Theory]. The likelihood ratio is the probability of observing the evidence given that the prosecutor’s or plaintiff’s hypothesis is true, divided by the probability of observing the same evidence given that the defense’s hypothesis is true. Let E be the evidence, H the prosecutor’s or plaintiff’s hypothesis, and H' the defense’s hypothesis. Recall that the likelihood ratio, $LR(E, H, H')$, is defined as follows:

$$LR(E, H, H') = \frac{PE|H}{PE|H'}$$

On this interpretation, relevance depends on the choice of the competing hypotheses. H_p and H_d are used as examples, but other competing hypotheses H and H' could also be used. When there are no ambiguities, $LR(E, H_p, H_d)$ will be shortened into the less cumbersome $LR(E)$. On the approach under consideration, a piece of evidence is relevant—in relation to a pair of hypotheses H and H' —provided the likelihood ratio $LR(E, H, H')$ is different from one and irrelevant otherwise. For example, the bloody knife found in the suspect’s home is relevant evidence in favor of the prosecutor’s hypothesis because we think it is far more likely to find such evidence if the suspect committed the crime (prosecutor’s hypothesis) than if he did not (defense’s hypothesis) (Finkelstein, 2009). In general, for values greater than one, $LR(E, H, H') > 1$, the evidence supports the prosecutor’s or plaintiff’s hypothesis H , and for values below one, $LR(E, H, H') < 1$, the evidence supports the defense’s hypothesis H' . If the evidence is equally likely under either hypothesis, $LR(E, H, H') = 1$, the evidence is considered irrelevant.

This account of relevance has been challenged by cases in which the evidence is intuitively relevant and yet its likelihood ratio, arguably, equals one. Here is one of them (The difficulty has been formulated by Ronald Allen, see the multi-authored discussion in Park et al., 2010):

Small Town Murder. A person accused of murder in a small town was seen driving to the small town at a time prior to the murder. The prosecution's theory is that he was driving there to commit the murder. The defense theory is an alibi: he was driving to the town because his mother lives there to visit her. The probability of this evidence if he is guilty equals that if he is innocent, and thus the likelihood ratio is 1 . . . , and under what is suggested as the "Bayesian" analysis, it is therefore irrelevant. Yet, every judge in every trial courtroom of the country would admit it [as relevant evidence] . . . and I think everyone on this list would say it is relevant. And so we have a puzzle.

Seeming counterexamples of this sort abound, here are a few of them:

- Suppose a prisoner and two guards had an altercation because the prisoner refused to return a food tray. The prisoner had not received a package sent to him by his family and kept the tray in protest. According to the defense, the prisoner was attacked by the guards, but according to the prosecution, he attacked the guards. The information about the package sent to the prisoner and the withholding of the tray fails to favor either version of the facts, yet it is relevant evidence (Pardo, 2013).
- In response to an eyewitness testimony the defendant claims that his identical twin is the culprit. The testimony is unable to favor any of the two options and yet is considered relevant.
- Suppose the evidence at issue is that a fight occurred and the only dispute is over who started it.
- Or suppose the defendant was stopped because of speeding three minutes after an aborted bank robbery and 1/2 a mile away from the site. The prosecution says this is evidence of guilt: it shows the defendant was escaping. The defense responds that this is evidence of innocence: no bank robber would speed and attract attention.
- Or, in a murder case, the defendant is the victim's son. Is that relevant to show he's guilty? Is it relevant to show he's innocent? The answer seems to be yes, to both questions (this example is due to Samuel Gross and is discussed in (Park et al., 2010)).

In general, there seem to be numerous examples in which evidence is, intuitively relevant, and the evidence supports neither side's theory over the other side's theory. How is such evidence to be judged relevant from the probabilist perspective?

In response (inspired by the ideas put forward in the discussion by David Kaye, Bruce Hay and Roger Park), note that it is true that if a piece of evidence E fits equally well with two competing hypotheses H and H' , then $P(E|H) = P(E|H')$ and thus $LR(E, H, H')$ will equal 1. But the likelihood ratio may change depending on the selection of hypotheses. Rule 401 makes clear that relevant evidence should have 'any tendency to make the existence of *any fact that is of consequence* [emphasis ours] to the determination of the action more probable or less probable'. So the range of hypotheses to compare should be quite broad. Just because the likelihood ratio equals one for a specific selection of H and H' , it does not follow that it equals one for *any* selection of H and H' which are of consequence to the determination of what happened. In *Small Town Murder*, whether the suspect was in town at all is surely of consequence for determining what happened (if he was not in town, he could not have committed the crime). The fact that he was seen driving is helpful information for establishing whether or not he was in town.

But if the range of hypotheses H and H' to compare in the likelihood ratio $LR(E, H, H')$ is quite broad, this may raise another concern. The choice of hypotheses needed to determine the relevance of an item of evidence might depend on other items of evidence, and so it might be difficult to determine relevance until one has heard all the evidence. This fact—Ronald Allen and Samuel Gross argue in (Park et al., 2010)—makes the probabilistic account of relevance impractical. But, in response, David Kaye points out that deciding whether a reasonable juror would find evidence E helpful requires only looking at what hypotheses or stories the juror would reasonably consider. Since the juror will rely on several clues about which stories are reasonable, this task is computationally easier than going over all possible combinations of hypotheses (Park et al., 2010).

Legal probabilists can also offer a more principled response to *Small Town Murder* and related problems based on Bayesian networks. Let H_p be the prosecutor's hypothesis that the defendant committed the murder, and H_d the defense's hypothesis that the defendant was visiting his mother. Let E be the fact that the defendant was seen driving to the town prior to the murder. Further, suppose the

prior probabilities of H_d and H_p are .5, and the conditional probability of E on each of those hypotheses is .7 (nothing of what will be said depends on this particular choice of values). Crucially, while indeed the evidence supports both hypotheses, this example is based on a pair of hypotheses that are neither mutually exclusive nor exhaustive. A Bayesian network can be used to calculate other likelihood ratios for hypotheses that are exclusive and exhaustive.



Figure 8: Graphic model of Small Town Murder, with the probability distribution of E .

Following the calculations in (de Zoete, Fenton, Noguchi, & Lagnado, 2019), for exclusive and exhaustive hypotheses, $LR(E, H_d, \neg H_d) = 1.75$, and similarly, $LR(E, H_p, \neg H_p) = 1.75$, since $P(E|H_d) = 0.7$ and $P(E|\neg H_d) = 0.4$. The likelihood ratio of the evidence, if it is measured against exclusive and exhaustive hypotheses, is not equal to one. (de Zoete et al., 2019 offer a slightly different solution to the problem. They construct a Bayesian network with three hypotheses, also exhaustive and exclusive: in town to visit mother, in town to murder, out of town.) Such considerations should also generalize to other paradoxes of relevance.

For instance, in the twins problem, the LR is 1 if the hypotheses are: ‘the suspect committed the crime’, and ‘the suspect’s twin brother committed the crime’, but is not 1 if we consider the fairly natural hypothesis that the defendant is innocent.

Similarly, in the food tray example, Bayesian network analysis shows that the value of the evidence ‘prisoner withholds tray’ for the question who started the fight depends on a range of uncertain events and other pieces of evidence (such as whether indeed a parcel he was supposed to obtain was withheld; whether the prisoner inquired about this; whether and how this inquiry was answered). Considered in this context, the piece of evidence will not have a likelihood ratio of one with respect to at least some choice of sensible hypotheses.

The general problem with the paradoxes of relevance is that in complex situations there is no single likelihood ratio that corresponds to a single piece of evidence. The problematic scenarios focus on a single likelihood ratio based on non-exclusive or non-exhaustive hypotheses. However, evidence can be relevant so long as it has a probabilistic impact on a sub-hypothesis involved in the case, even without having a recognizable probabilistic impact on the prosecutor’s or defense’s ultimate hypotheses. When this happens, it is relevant, in agreement with Rule 401 of the Federal Rules of Evidence. Bayesian networks help to see how pieces of evidence can increase or decrease the probability of different sub-hypotheses (de Zoete et al., 2019). Thus, while likelihood ratios are useful, they need to be properly presented in the context.

A. The cold-hit confusion

We have been quite negative in the last two sections. What we are after, though, is a balanced perspective. We are not saying that likelihood ratios should be disposed of: rather that they should be reported for individual pieces of evidence and used to evaluate the impact of combined evidence, but only when proper attention is paid to the structure of the problem. To re-emphasize the utility of likelihood ratios when care is paid to both the hypothesis formulation, and to what the evidence in fact is, we illustrate it by showing how it can be used to resolve a lengthy debate in legal evidence studies: the value of a cold-hit DNA match. First, we spend some time explaining the problem in Appendix A, and then show how LR can be deployed to resolve it in Appendix B.

First, we explain what a cold hit is and give an example of case based on a cold hit, the murder of Diana Sylvester. In that case, instead of random match probability, the defense focused on another measure, database match probability (DMP). We explain what it is and how it was recommended the the

Added this intro here.

Marcello: added this signposting bit, to make the appendix structure more transparent, is this good enough?

second National Research Council on DNA Technology in Forensic Science report. Then we explain how NRC tried to motivate their recommendation, and why these recommendation are not convincing. This sets up the stage for the next appendix: if DMP is not the right tool to evaluate cold-hit evidence, what is?

DNA evidence is one of the most widely used forms of quantitative evidence currently in use. It may be used to corroborate other evidence in a case, or as the primary incriminating evidence. For example, suppose different investigative leads point to an individual, Mark Smith, as the perpetrator. The investigators also find several traces at the crime scene left by the perpetrator. Laboratory analyses show that the genetic profile associated with the traces matches Smith. In this scenario, the DNA match corroborates the other evidence against Smith. In contrast, suppose the police has no other investigative lead except the traces left at the crime scene. Hoping to find the perpetrator, the police run the genetic profile associated with the traces through a database of profiles and find a match, a so-called **cold-hit**. Cold-hit DNA matches have been the focus of intense discussion in recent years. Since in cold-hit cases there is little or no other evidence, cold-hit matches are often the primary item of evidence against the defendant. Some believe that this circumstance weakens the case. Others disagree. This debate illustrates how probability theory—in particular, the likelihood ratio—can help to assess the strength of evidence at trial. What follows examines some of the main arguments.

For concreteness, consider the California rape and murder case of Diana Sylvester. In 2008, many years after the crime, John Puckett was identified as a unique 9-loci match through a database search of 338,000 profiles. He was the only individual in the database who matched the traces collected from the victim Diana Sylvester in 1972. According to an expert witness, the particular pattern of alleles present in the material was (conservatively) expected to occur randomly among Caucasian men with a frequency of 1 in 1.1 million. This is the **random match probability (RMP)**. The random match probability—often interpreted as the probability that someone who is not the source would coincidentally match, $P(\text{match}|\neg\text{source})$ —is a common measure of the strength of a DNA match. The lower the RMP, the more strongly incriminating the match. The rationale here is that a low random match probability suggests that it is unlikely that two people would share the same DNA profile. In line with what we already discussed, strictly speaking, a match is a strong evidence that the defendant is the source only if the probability that the person who left the traces (the ‘source’) would match is significantly greater than RMP. In practice, when it comes to DNA evidence, it is often assumed that $P(\text{match}|\text{source})$ is very high.

Although clearly 1 in 1.1 million should not be confused with the probability of Puckett’s innocence, the small figure indicates it is very unlikely that a random person unrelated to the crime would match. The match is therefore strong evidence of Puckett’s guilt. Assuming that the probability of a match if Puckett indeed was the source was (practically) 1, the likelihood ratio is simply 1.1×10^6 .

During the pretrial hearing, however, Bicka Barlow, the DNA expert for the defense, pointed out that this was a cold-hit case. No evidence tied Puckett to the crime other than the cold-hit match, Puckett’s previous rape convictions, and the fact that he was in the area at the time of the murder. In order to correctly assess the probative value of the cold-hit match, Barlow argued, the random match probability should be multiplied by the size of the database. The result of such a multiplication is called the **database match probability (DMP)**. In Puckett’s case, the multiplication of $1/1.1 \times 10^6$ by 338,000 resulted in a database match probability of approximately .3.

which is a less impressive number than the original RMP (the likelihood ratio for the DMP is approximately 3.25). According to this calculation, it was no longer very unlikely that an unrelated person from the database would match, and so the cold-hit DNA match was no longer strong evidence of guilt. At least, this was Barlow’s argument.

Barlow followed a 1996 report by the National Research Council called NRC II (National Research Council, 1996), preceded by an earlier report on DNA evidence called NRC I (National Research Council, 1992). NRC II recommended precisely what Barlow did: that in cold-hit cases RMP should be multiplied by the database size, yielding DMP. The underlying idea was that the larger the size of the dataset, the higher the database match probability, and the lower the strength of the match. This correction was meant to guard against the heightened risk of mistaken matches for the innocent people in the database. To see, however, if this was sound advice, we need to look under the hood.

The NRC formed the Committee on DNA Technology in Forensic Science, which issued its first report in 1992. In that report they advised against using cold hit results as evidence, and insisted that only the frequencies related to loci not used in the original identification should be presented at trial,

that is, that the evidence used to identify the suspect should not be used as evidence against the suspect.

This recommendation has been criticized by many because it underestimates the value of cold-hit matches. The problem was, given a certain amount of evidence the expert, prior to suspect identification, had to make a somewhat subjective decision of how to divide the evidence into two items: one to be used only in the suspect identification, and one to be used only in the trial itself as evidence against the suspect. This overly limited the utility of the evidence and introduced an unnecessary element of subjectivity.⁸

NRC II withdrew the earlier recommendation. However, the contrast between low RMP and the frequency of DNA matches in actual database searches was indeed stark. For instance, the Arizona Department of Public Safety searched for matching profiles in a database comprising 65,000 individuals. The search found 122 pairs of people whose DNA partially matched at 9 out of 13 loci; 20 pairs people who matched at 10 loci; and one pair of people who matches at 12 loci. So it is not that unlikely to find two people in a database who share the same genetic profiles (examples of fairly high counts of DNA matches in database searches was actually used by Barlow in the Diana Sylvester case). In light of this contrast, NRC II recommended the use of DMP rather than RMP. NRC II recommended also that in cold-hit cases the likelihood ratio R associated with the DNA match should be divided by $d + 1$. Their first recommendation was about a correction of the random match probability, and this second recommendation is about the likelihood ratio.

One argument by NRC employed an analogy involving coin tosses. If you toss several different coins at once and all show heads on the first attempt, this seems strong evidence that the coins are biased. If, however, you repeat this experiment sufficiently many times, it is almost certain that at some point all coins will land heads. This outcome should not count as evidence that the coins are biased. According to NRC II, repeating the coin toss experiment multiple times is analogous to trying to find a match by searching through a database of profiles. As the size of the database increases, so does the number of attempts at finding a match, and it is more likely that someone in the database who had nothing to do with the crime would match.

Another argument provided by NRC II compared a database trawl to multiple hypothesis testing, and multiple hypothesis testing should be avoided if possible in light of classical statistical methods.

Third, NRC II was concerned with the fact that in cold-hit cases the identification of a particular defendant occurs after testing several individuals. This concern has to do with the data-dependency of one's hypothesis: seemingly, the hypothesis 'at least one person in a given database matches the DNA profile in question' changes its content with the choice of the database.

We will start with the coin analogy. It is in fact unclear how the analogy with coin tossing translates to cold-hit cases. Searching a larger database no doubt increases the probability of finding a match at some point, but is the increase as fast as the Arizona Department of Public Safety examples and the coin analogy suggest? Quite crucially, following (Donnelly & Friedman, 1999) we need to pay attention to what hypotheses are tested, what probabilistic methods the context recommends, and what exactly the evidence we obtained is. For instance, one hypothesis of interest is what we will call a *general match hypothesis*:

(General match hypothesis) At least one of the profiles in the database of size n matches the crime sample.

The general match hypothesis is what NRC II seems to have been concerned with. If for each data point $RMP = \gamma$ were held constant, and if random matches with different data points $match_1, match_2, \dots, match_d$ excluded each other, the probability of there being at least one random match would be the same as the probability of their disjunction and could be calculated by the additivity axiom:

$$\begin{aligned} P(\text{at least one match}) &= P(\text{match}_1 \vee \text{match}_2 \vee \dots \vee \text{match}_d) \\ &= \sum_i^d P(\text{match}_i) = \gamma \times d \end{aligned}$$

⁸It also opened the gate for multiple testing with various evidence division points, and multiple testing leads to its own statistical problems. But let's put this issue aside.

Rafal: M: A lot of this is not about LR, so the reader now is confused and impatient since you promised you would talk about the virtue of LR in DNA evidence evaluation. Need to say upfront this section is NOT about LR but about some other metrics, or else the reader would be lost. The next section is about LR, right?

Marcello: now the first passage in the appendix says it introduces the problem and that the next one deploys LR to handle it, is this ok now?

This calculation would result in the outcome recommended by NRC II, if the value of the evidence were to be a function of the probability of (General match hypothesis).

The first question is, whether a directly additive calculation should be applied to database matches. Notice that in applications DMP does not really behave like probability. Take a simple example. Suppose a given profile frequency is .1 and you search for this particular profile in a database of size 10. Does the probability of a match equal $.1 \times 10 = 1$? The answer is clearly negative. Multiplication by database size would make sense if we thought of it as addition of individual match probabilities, provided matches exclude each other and so are not independent. Here is a coin analogy. Suppose I toss a die, and my database contains $n = \text{three different numbers: } 1, 2 \text{ and } 3$. Then, for each element of the database, the probability p of each particular match is $1/6$, and the probability of *at least one* match is $1/6 + 1/6 + 1/6 = 1/6 \times 3 = n \times p = 1/2$. We could use addition in such a situation because each match excludes the other matches, a condition that is not satisfied in the database scenario.

Another reason why DMP is problematic can be seen by taking a limiting case. Suppose everyone in the world is recorded in the database. In this case, a unique cold-hit match would be extremely strong evidence of guilt, since everybody except for one matching individual would be excluded as a suspect. But if RMP were to be multiplied by the size of the database, the probative value of the match as measured by DMP should be extremely low. This is highly counter-intuitive.

Even without a world database, the NRC II proposal remains problematic, since it sets up a way for the defendant to arbitrarily weaken the weight of cold-hit DNA matches. It is enough to make more tests against more profiles in more databases. Even if all the additional profiles are excluded (intuitively, pointing even more clearly to the defendant as the perpetrator), the NRC II recommendation would require to devalue the cold-hit match even further. This, again, is highly counter-intuitive.

Perhaps a somewhat more sensible answer is obtained by assuming the independence of nomatch for the members of the database and deploying a solution similar to the one used in the birthday problem. Here, the idea would be—assuming matches for different data points are independent and have constant RMP — to calculate:

$$\begin{aligned} P(\text{match}) &= 1 - P(\text{nomatch}) \\ &= 1 - (1 - \gamma)^d \end{aligned}$$

where γ is RMP, and d is the database size. This would be in line with using the binomial distribution to calculate the probability of no match:


$$\begin{aligned} \text{dbinom}(0, d, \gamma) &= \binom{n}{0} \gamma^0 (1 - \gamma)^{d-0} \\ &= 1 \times 1 \times (1 - \gamma)^d \end{aligned}$$

Now, assuming indeed that γ is constant and that matches between data points are independent, the dependence of the probability of at least one match on the database size can be pictured as in Figure A.

If we use the RMP and database size used in the Puckett case, the calculated probability of at least one match is 0.2645501. Not exactly the DMP postulated by the defendant, but pretty close. The question is, should this number be the probability used to evaluate the evidential impact of the cold hit?

One problem is, whether the independence assumption is satisfied in the database search problem is unclear. After all, if you are informed about the match frequencies in the database, and they teach you that since two arbitrary database points quite likely do not match, if the sample matches one of them, it is less likely to match the other one. And the independence assumption is not benign. We will illustrate it with a somewhat distand, but a very striking example, coming from (Barnett, 2020). Suppose you consider whether your effort of casting a vote in the upcoming election is worth it in a context where there are 500k other voters. One of the probabilities you might be interested in is the probability that your vote would make a difference. If we apply the binomial model to the problem, the probability that a candidate will receive exactly k votes if n people vote is supposed to be $\binom{n}{k} p^k (1 - p)^{(n-k)}$, where votes of the population members are supposed to be independent and estimated to have the same probability p of being for the candidate. For instance, if 500k people vote and $p = 0.5$, the probability that the candidate will receive exactly 250k votes is 0.0011284, which is around $1/886$ and much higher than $1/n$. This fairly high chance made some claim that the chance that your voice is decisive if the chances are equal is

fairly high in such circumstances. However, note that if $p = .505$, the probability that the candidate will receive exactly 250k votes is $1.5651281 \times 10^{-14}$, which is less than one in a trillion. This lead some (Brennan, 2012; *Democracy and decision*, 1993) to claim that outside of the very specific circumstances, decision-theoretic arguments for the rationality of voting are hopeless. Barnett, however, points out that such a sensitivity to success probability simply makes the binomial model inappropriate for the voting context, observing that its calculations also disagree with empirical estimates which are not too far from $1/n$ (Gelman, Katz, & Tuerlinckx, 2002; Gelman, King, & Boscardin, 1998; Mulligan & Hunter, 2003). This sensitivity arises, because within the binomial model the more trials (voters) there are, the more tightly the results will tend to cluster around the probability of success. To observe how unrealistic that is, keep $p = .505$ and ask yourself how probable it is that the voting result will be between 50.4% and 50.6%. Sure, this outcome might be quite likely, but the binomial certainly overestimates it at 0.8427212. Another unrealistic estimate obtained by the binomial model is the estimate of the probability of an upset (that the leading candidate will lose). With $p = .505$ this is $\text{pbinom}(249999, 500000, .505)$, which turns out to be extremely and unrealistically low: $7.5994495 \times 10^{-13}$.



Rafal: M: Pretty interesting, but in the voting model, it seems that feedback mechanisms between voters affect behaviour, but not so in the genetic case. Can you spell out the analogy more clearly?

Marcello: Not sure what you mean, we should talk about this.

Figure 9: Binomial model of the database search problem. The probability of at least one match depending on the database size, assuming independence and constant RMP used in the Puckett case as compared with the binomial estimate for $p=2/1.1e6$ (dashed line). The actual database size marked with a vertical line."

Coming back to our original problem, the binomial estimate of the probability of a match is also quite sensitive to RMP, as illustrated in Figure A. The bottom line is that if we have reasons to think the independence assumption is not satisfied, the binomial model is not appropriate. So, it seems, it is not appropriate for the database match problem either.

The binomial model, however, is useful, in its simplicity, for illustrating an important distinction whose conflation underlies one of the involved arguments. You might have been surprised learning that while the expert testified that RMP on 9 loci for Puckett was 1 in 1.1 million, the Arizona Department of Public Safety found 122 9-loci matches among 65,000 individuals. After all, $122/65000$ is 0.0018769, which is much higher than the reported RMP.

Crucially, notice that there is a difference between having a sample and looking for a match in a database of size n and taking a database of size n and checking all pairs that occur within it for a match. In the former case, you are making n comparisons. In the latter case, the number of comparisons is $\binom{n}{2}$, which is much higher. If $n = 65000$, there are 2.1124675×10^9 pairs to compare, so while the binomial estimate of the probability of at least one match for n comparisons (the former case) is 0.057379, it is approximately 1 for $\binom{n}{2}$ comparisons. For the impact it has on the Arizona Department of Public Safety statistics, consider the binomial estimate of the probability of at least 122 matches among all pairs as a

function of the database size, even for relatively low database size range (up to 50000), as illustrated in Figure 10.

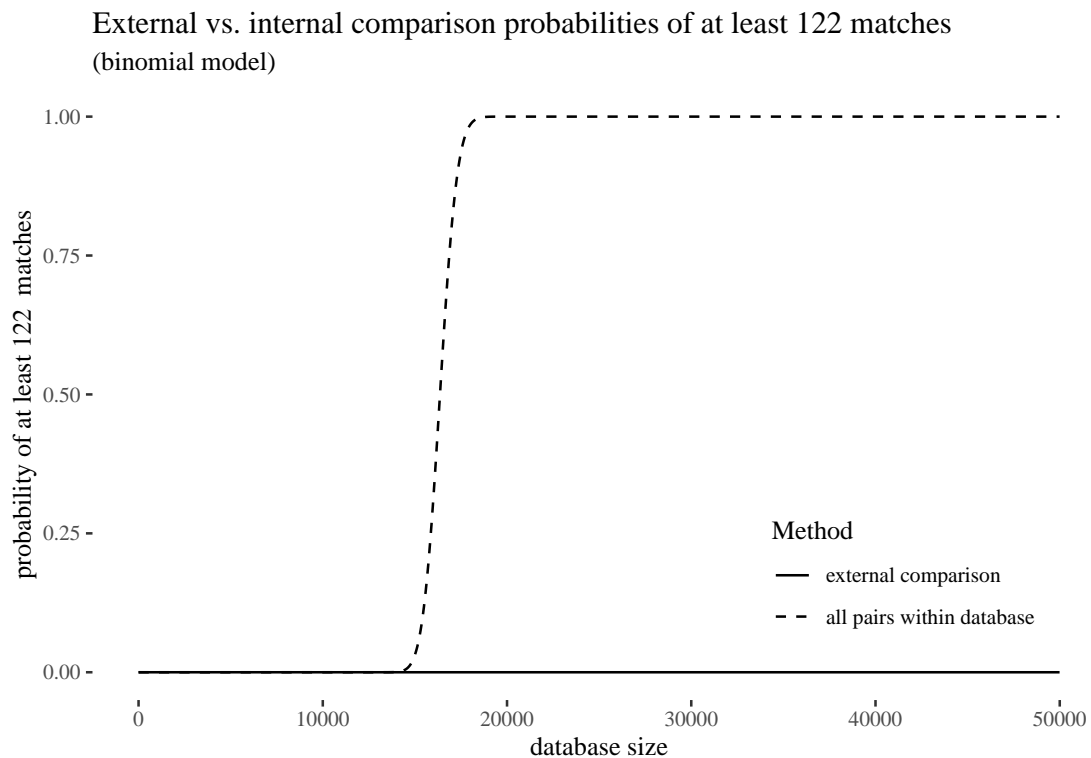


Figure 10: Binomial model of the database search problem. The probability of at least 122 matches depending on the database size for n comparisons with an external sample, and for all possible pairs among n datapoints (dashed), assuming $RMP=1/1.1e6$.

From this perspective, it is no surprise there were so many matching pairs among all the pairs from the database. Unfortunately, this frequency does not estimate the probability of at least one match in the set-up we are actually interested in. After all, in a cold-hit scenario we do have a sample outside of the database and make n comparisons, instead of testing all possible pairs from the database for a match.

Before we move on, it is worthwhile to think about how the example at hand actually constitutes evidence against a certain probabilistic model that one might initially want to deploy in this context. Note how the Arizona statistics constitute some empirical evidence against the adequacy of the binomial model. While the binomial estimate probability of at least 122 matches with an external sample for $n = 6500$ is pretty much 0, let us look at the most likely number of matches if we test $\binom{6500}{2}$ pairs, as estimated by the binomial model. We illustrate it with an 89% highest density interval in Figure 11. So, if the expert's estimate and the binomial model are both adequate, we indeed should be surprised by the presence of 122 matches. But this is because this number is surprisingly low: instead we should expect a much higher number, around 2000 of them! Of course, it is unlikely that in fact all possible pairs have been compared, and it is hard to evaluate this evidence against the binomial model unless we know the exact number of comparisons made.

Now that we used the imperfect binomial model to clear up at least one confusion, let us put it aside, and focus on an even deeper problem with using the probability of (General match hypothesis) instead of RMP in evidence evaluation. Probabilistic epistemology recommends that once we obtain new evidence, our new degrees of belief should be the probabilities obtained by conditionalizing on this evidence. Crucially, we should update on the total evidence we obtained rather than only on a part of it. Here the question is, does (General match hypothesis) exhaust what we have learned from our database match?

To get us started thinking about this question, consider a coin analogy which Donnelly & Friedman (1999, p. 950) found more adequate than the one proposed by the NRC. Imagine a biased coin whose

Rafal: M: In general this section needs more signposting for guiding the reader. This last argument seems to be an add-on.

Marcello: added appendix structure in the beginning and revised the next sentence to improve flow




Figure 11: Binomial probability density of n matches in pairwise comparison within a database of size 65000, assuming $p = 1/1.1e6$, with 89% highest density interval = (1849, 1990) shaded in blue.

physical appearance is indistinguishable from the fair coins in a piggy bank. The biased coin is the perpetrator and the piggy bank is the database containing innocent people. After the biased coin is thrown into the bank with the other coins, someone picks out a handful of coins at random and flips each of them twenty times. Each coin lands heads approximately ten times—except for one coin, which lands heads on all twenty flips. The fact that other coins seem unbiased makes the claim that this one is biased better supported.

Coming back to DNA matches, think about the following scenario: first, you identified the suspect by some means other than a database trawl. Then, it turned out his DNA profile matches the crime scene stain. Fine, here it seems uncontroversial that this constitutes strong further incriminating evidence. Now, imagine a further database search in a database not containing this suspect finds no matches. Would you think that this information supports the guilt hypothesis? If your answer is yes, then you do have the intuition that the lack of matches with other people (whose profiles, in this particular case, happen to be in a database) strengthens the evidence.

The key lesson here (and in the complete-world database scenario we already discussed) is that we not only learned that there was a match in the database of size n , but also that in $n - 1$ cases there was no match, and this information also has evidential value. In line with this, contrary to NRC II, Donnelly & Friedman (1999) argues that if potential suspects in the database are excluded as sources, this should increase, not decrease, the probability that the defendant who matches the crime traces is the source. A cold-hit match, then, is stronger and not weaker evidence of guilty than ordinary DNA matches.

Before moving on to a better model that captures how this could be, let us look at another argument put forward by NRC, an analogy to multiple hypothesis testing. NRC claimed that there is an analogy between searching for a match in a database and multiple hypothesis testing, which is a dubious research practice. In classical hypothesis testing, if the probability of type I error in a single test of a hypothesis is 5%, the probability of at least one type I error will increase by testing the same hypothesis multiple times. In analogy—the argument goes—we need to correct for the increased risk of type I error, and just as the Bonferroni correction requires that the p -value threshold be divided by the number of tests, NRC II requires that the estimated probability of a random match should be multiplied by the number of comparisons.

This analogy with multiple testing, however, is misplaced. As Balding (2002) points out, multiple testing consists in testing the *same* hypothesis multiple times against new evidence. In cold-hit cases, no such multiple testing is involved. Rather, multiple hypotheses—each concerning a different individual in the database—are tested only once and then excluded if the test is negative. From this perspective, for each $1 < i < n$, the following hypothesis is tested:

(Particular match hypothesis) Profile i in the database matches the crime sample.

and the hypothesis that the defendant is the source was one of the many hypotheses subject to testing. The cold-hit match supports that hypothesis and rules out multiple other hypotheses.

B. Likelihood ratio and cold-hit DNA matches

To find a better path towards the resolution of the database search problem, let us look at another recommendation of NRC II, that in cold-hit cases the likelihood ratio R associated with the DNA match should be divided by $n + 1$, where n is the database size. This approach has been defended by (Stockmarr, 1999), who points out that since the suspect is identified after the databases search, the hypothesis is formulated *ad hoc*. Without the correction, then, the likelihood ratio would be data-dependent. For instance, he insists that hypotheses such as *JS was one of the crime scene donors* are evidence-dependent in the case of database search, “since we had no way of knowing prior to the search that Smith would be the person that matched” (p. 672). Instead, Stockmarr claims, we should evaluate LR using hypotheses that can be formulated prior to the search, such as *the true perpetrator is among the suspects identified from the database*. And indeed, the likelihood of this hypothesis is as NRC II suggests, k/np , where k is the number of matching profiles, n the database size, and p the random match probability (see Stockmarr, 1999 for a derivation).

Dawid, in a discussion with Stockmarr (Dawid & Stockmarr, 2001) points out that Stockmarr’s hypotheses, while not depending on the result of the search, depend on the data themselves (because they change with the database size). More importantly, he also indicates that Stockmarr’s hypotheses are composite and the assessment of LR therefore requires additional assumptions about the priors. Once these are used with Stockmarr’s own LR, the posterior is the same as the one obtained using the methods proposed by the critics of NCR II. This is a particular case of a general phenomenon that we will discuss later on: different hypotheses might result in different LR, but be equivalent conditional on the evidence, and so result in the same posterior probabilities. This general phenomenon indicates how LR on its own might be insufficiently informative.⁹

Putting Stockmarr’s defense and its problems aside, the NRC II recommendation is questionable on more principled grounds. Suppose R is not too high, say because the identified profile is common since the crime scene DNA is degraded and only a few markers could be used. Then, $n + 1$ can be greater than R , so $R/(n + 1) < 1$. The match would then be exculpatory, a very counter-intuitive result. Moreover, if the defendant on trial is the source, the probability that he would match is practically 1. If he is not, the probability that he would still match equals the random match probability. Neither of these probabilities change because other suspects have been tested in the database search. In fact, if potential suspects are excluded as potential sources, this should increase, not decrease, the probability that the defendant who matches the crime traces is the source.

A more principled way to assess cold-hit matches based on the likelihood ratio, exists. The proposal draws from the literature on the so-called **island problem**, studied by (Dawid, 1994; Dawid & Mortera, 1996; Eggleston, 1978). Let the prosecutor’s hypothesis H_p be ‘The suspect is the source of the crime traces’ and the defense’s hypothesis H_d be ‘The suspect is not the source of the crime traces’. Let E be the DNA match between the crime stain and the suspect (included in the database) and D the information that none of the $n - 1$ profiles in the database matches the crime stain. The likelihood ratio associated with E and D should be (Balding & Donnelly, 1996; Taroni et al., 2014):

$$V = \frac{P(E, D|H_p)}{P(E, D|H_d)}.$$

Since $P(A \wedge B) = P(A|B)P(B)$, for any statement A and B , this ratio can be rewritten as:

⁹See however, Stockmarr’s own reply (*ibidem*).

$$V = \frac{P(E|H_p, D)}{P(E|H_d, D)} \times \frac{P(D|H_p)}{P(D|H_d)}. \quad (14)$$

The first ratio in (14) is roughly $1/\gamma$, where γ is the random match probability. The second ratio—call it the **database search ratio**—requires some more work. Consider first the denominator $P(D|H_d)$. If the suspect is not the source (H_d), someone else is, either someone who is in the database or someone not in the database. Let S stand for The source is someone in the database. By the law of total probability,

$$P(D|H_d) = P(D|S, H_d)P(S|H_d) + P(D|\neg S, H_d)P(\neg S|H_d). \quad (15)$$

If the source is someone in the database (S) and the suspect is not the source (H_d), it is very unlikely that no one in the database would match (D), so $P(D|S, H_d) \approx 0$. The equality in (15) therefore simplifies to:

$$P(D|H_d) = P(D|\neg S, H_d)P(\neg S|H_d),$$

The database search ratio would therefore be:

$$\frac{P(D|H_p)}{P(D|H_d)} = \frac{P(D|H_p)}{P(D|\neg S, H_d)P(\neg S|H_d)}.$$

Note that $P(D|H_p) = P(D|\neg S, H_d)$ because whether the suspect is the source (H_p) or not (H_d) does not affect whether there is a match in a database that does not contain the source ($\neg S$). Let the probability that no person in the database other than the suspect would match (D), assuming the suspect was in fact the source, be ψ_{n-1} . Notice that $P(D|\neg S, H_d)$ is the probability that no one other than the suspect matches in the database that does not contain the real source, if the suspect is not the source. So this conditional probability can also be estimated as ψ_{n-1} .¹⁰ Let $P(S|H_d) = \varphi$. The database search ratio then would reduce to

$$\frac{P(D|H_p)}{P(D|H_d)} = \frac{1}{1 - \varphi}.$$

As the database gets larger, φ increases and the database search ratio also increases. This ratio equals one only if no one in the database could be the source, that is, $\varphi = 0$.

Since the likelihood ratio V of the cold-hit match results by multiplying the likelihood ratio of the DNA match and the database search ratio, V will always be greater than the mere likelihood ratio of the match (except for the unrealistic case in which $\varphi = 0$). Thus, a cold-hit DNA match should count as stronger evidence than a DNA match of a previously identified suspect.

Dawid & Mortera (1996) study different database search strategies and consider the possibility that information about the match is itself uncertain, but the general point remains. Under reasonable assumptions, ignoring the database search would give a conservative assessment of the evidentiary strength of the cold-hit match. Donnelly & Friedman (1999), with slightly different assumptions, derived the formula $R \times [1 + md/N]$, where $R = 1/\gamma$, d is the database size, N the number of people in population not in database, and m is an optional multiplier reflecting how much more likely persons in the database are thought to be the source when compared to the rest of the population. The expression cannot be less than γ . If no other profile has been tested, $d = 0$ and LR is simply the regular DNA match LR. If N is zero, that is, everyone in population is in the database, the result is infinitely large.

This proposal is able to accommodate different apparently competing intuitions. First, consider the intuition that as the size of the database grows, it is more likely that someone in the database would match. This intuition is captured by the fact that φ increases proportionally to the size of the database even though this increase does not imply that the evidential value of the cold-hit match should decrease. Second, there is intuitive resistance to basing a conviction on a cold-hit match, although this resistance is less strong in case of an ordinary match (more on this later in Section ??). This preference for convictions

Fix crossref later

¹⁰If the prior probability that the perpetrator is in the database was high, the calculations would need to be different. But normally, this prior is not too high.

based on an ordinary DNA match seems in tension with the claim that a cold-hit match is stronger evidence of guilt than an ordinary match. There is a way to make sense of this, though. The key is to keep in mind that the evidentiary strength—measured by the likelihood ratio—should not be confused with the posterior probability of guilt given the evidence. Even if a cold-hit match is stronger evidence of guilty, this fact does not imply that the posterior probability of the defendant’s guilt should be higher. If the cold-hit match is the only evidence of guilt, the posterior probability of guilt may well be lower compared to cases in which other evidence, such as investigative leads, supplements the DNA match. This lower posterior probability would justify the intuitive resistance towards convictions in cold-hit cases, despite the fact that a cold-hit match alone is stronger evidence than a dna match obtained otherwise and taken on its own. Moreover, it is possible that the intuitive assesment of cold-hit evidence takes to some extent the impact of false positive probability into account.

C. Confirmation measures

At this point a philosophically minded reader might recall that there is an important notion in the vicinity—that of confirmation—and that there is a vast philosophical literature on probabilistic explication of that notion. Natural questions arise: how is the notion of confirmation related to the notion of evidence strength, and why almost none of the probabilistic explications of confirmation have not been deployed in legal probabilism?

The key question behind the enterprise we are going to take a look at is: when does a piece of evidence confirm a theory, and how are these requirements to be explicated probabilistically to agree with both successful scientific practice and sensible philosophical principles? The hope is that having answered these questions would facilitate both rational reconstructions of various developments in the history of science, and a critical evaluation of various ongoing scientific investigations.

The underlying probabilistic idea is that the level of confirmation of a theory (T) by a piece of evidence (E) is a function of an agent’s degrees of belief. The first stab might be, let’s simply identify the confirmation level with $P(T|E)$. This, however, is way too quick. Multiple factors come into the assessment of this conditional probability, and two agents can agree on the extent to which E confirms T without agreeing on the posterior probability of T (identified with $P(T|E)$), because the agents might disagree about the prior probability of T and this might have an impact on the posterior.

Still, some requirements on confirmation measures can be formulated in terms of probabilities. One usual assumption (Sprenger & Hartmann, 2019) is that the level of confirmation is to be a continuous function of $P(T)$ and $P(E|T)$ which is non-decreasing in the first argument and non-increasing in the second argument. That is, increasing the prior, should not lower the confirmation level, and increasing the likelihood should not increase the confirmation level. Let’s call this condition the *prior-posterior dependence*.¹¹

One consequence of the prior-posterior-dependence—called *Final Probability Incrementality* is that confirmation of T by E , $c(T, E)$ should track the posterior order-wise, that is $c(T, E) > c(T, E')$ just in case $P(T|E) > P(T|E')$.

Another requirement is that there should be a neutral point n such that E confirms (disconfirms) T just in case $c(T, E) > n$ ($c(T, E) < n$) and is neutral exactly at n . This is called the *qualitative-quantitative bridge*.

Yet another requirement is *local equivalence*. Theories that are logically equivalent given the evidence should receive equal confirmation from this evidence. Interestingly, all confirmation measures which satisfy prior-posterior dependence, qualitative-quantitative bridge, and local equivalence are strictly increasing functions of $P(H|E)$. Such measures are said to explicate confirmation as *firmness of belief*. Moreover, all functions satisfying these three conditions are ordinally equivalent.¹²

However, another notion of confirmation seems often at play. For instance, even if the posterior T is low, one might still think that a given experiment still speaks strongly in favor of T . And relatedly,

Rafal: M: We need a proper conclusion here. Now it feels as though LR are a good thing, at least for cold-hits. But is this just an exception? Do LR have a very narrow applicability, say only for DNA evidence? So where do we stand exactly? And what about BF and Bayesian networks and priors? Need general morals here.

Marcello: I think I now reiterated our point too many times throughout the chapter, we’ll discuss this.

¹¹ Some formulations (Crupi, 2015) are a bit more general and include background knowledge K . In that setting, the corresponding requirement is called *Formality* and takes the confirmation to be a function of $P(H \wedge E|K)$, $P(H|K)$ and $P(E|K)$. For the sake of simplicity, we will suppress the reference to K , unless required by the context.

¹² Measure c is ordinally equivalent to measure c' just in case always $c(E, T) \geq c(E', T')$ iff $c'(E, T) \geq c'(E', T')$.

E can lower the posterior of T while still leading the posterior to be sufficiently high for the firmness confirmation measure to be above the neutrality threshold. Another feature of confirmation as firmness is that if, in this sense, T confirms H , then for any H' that is excluded by H , T disconfirms H' . But now think of the small town murder scenario we already discussed: the fact that the suspect was seen in town seems to support both the prosecution hypothesis that he committed the murder, and the defense hypothesis, that he was in town to visit his mother. Confirmation as firmness cannot capture such intuitions, as relevance cannot be captured as a function of the posterior alone.

For such reasons, following the second edition of (Carnap, 1962), it is customary to distinguish another notion in the vicinity: confirmation as increase in firmness of belief. If we replace local equivalence with tautological equivalence $c(T, \top) = c(T', \top)$, where \top is a logical tautology—the idea being that hypotheses are equally supported by empty evidence—we end up with another class of confirmation measures, those meant to capture *probabilistic relevance*. On this approach, E confirms (disconfirms) T just in case $P(H|E) > P(H)$ ($P(H|E) < P(H)$).

Here is a list of key confirmation measures available on the market (Sprenger & Hartmann, 2019), normalized so that they all have neutral points at 0:

$$D(T, E) = P(T|E) - P(T) \quad (\text{Difference})$$

$$Lr(T, E) = \log \left(\frac{P(T|E)}{P(T)} \right) \quad (\text{Log-ratio})$$

$$LL(T, E) = \log \left(\frac{P(E|T)}{P(E|\neg T)} \right) \quad (\text{Log-likelihood})$$

$$K(T, E) = \frac{P(E|T) - P(E|\neg T)}{P(E|T) + P(E|\neg T)} \quad (\text{Kemeny-Oppenheim})$$

$$Z(T, E) = \begin{cases} \frac{P(T|E) - P(T)}{1 - P(T)} & \text{if } P(T|E) \geq P(T) \\ \frac{P(T|E) - P(T)}{P(T)} & \text{if } P(T|E) < P(T) \end{cases} \quad (\text{Generalized entailment})$$

$$S(T, E) = P(T|E) - P(T|\neg E) \quad (\text{Christensen-Joyce})$$

$$C(T, E) = P(E)(P(T|E) - P(T)) \quad (\text{Carnap})$$

$$R(T|E) = 1 - \frac{P(\neg T|E)}{P(\neg T)} \quad (\text{Rips})$$

(Log-likelihood), our good old likelihood ratio, and (Kemeny–Oppenheim) are ordinally equivalent (and no other pair on the list is). Further grouping and assessment of the confirmation measures for a given purpose is facilitated by the following facts:

- One might require that E always confirms the disjunction of excluding hypotheses more than one of them just in case it also confirms the other one (*disjunction of alternative hypotheses*). This can happen only if the confirmation measure is a strictly increasing function of the difference measure. Whether this is an intuitive requirement in our context is unclear.
- One might require that confirmation should track likelihood— $c(T, E) > c(T, E')$ (*Law of likelihood*)—just in case $P(E|T) > P(E'|T)$. This can happen only if the measure is a strictly increasing function of the Bayes factor. At least in legal applications, the law of likelihood is suspicious, as our example with rocking child abuse victims discussed on page 1 indicates.
- You might wish that confirmation be *contrapositive* ($c(T, E) = c(\neg E, \neg T)$) and *commutative* ($c(H, E) = c(E, H)$). The only measures that satisfy both are relative distance measures, that is, they are strictly increasing functions of the generalized entailment measure.
- One might require that if E and E' are conditionally independent given T and $\neg T$, then $c(T, E)$ should be identical with $c(T, E|E')$ (the confirmation obtained when E' is added to the background knowledge). This condition is called *modularity*. This condition holds only if a confirmation measure is a strictly increasing function of the likelihood ratio.

Moreover, if you require strict additivity: $c(H, E \wedge E') = c(H, E) + c(H, E'|E)$, the only measure that satisfies the disjunction of alternative hypotheses is the difference measure, the only measure that satisfies the law of likelihood is the log-ratio measure, and the only one that satisfies modularity is the log-likelihood measure.

Some unity can be brought into the picture (Crupi, Tentori, & Gonzalez, 2007) by normalizing by

what happens with a measure where logical consequence or exclusion is involved. For instance, if $E|T$, $D(E, T) = P(\neg T)$ and if $E|\neg T$, $D(E, T) = -P(T)$. So the normalized version has the form:

$$D_n(E, T) = \begin{cases} D(E, T)/P(\neg T) & \text{if } P(T|E) \geq P(T) \\ D(E, T)/P(H) & \text{otherwise.} \end{cases}$$

Interestingly, analogous normalization of measures other than (Generalized entailment) leads to the same single new Bayesian measure of confirmation: (Generalized entailment). Another reason one might have to like this measure is as follows. Take any $k > 0$ and say $v(E, T) = k$ iff $E \models T$, $v(E, T) = -k$ iff $E \models \neg T$ and $v(E, T) = 0$ otherwise. The *logical closure requirement* is that if $v(E, T) > v(E', T')$, then $c(E, T) > c(E', T')$. It turns out that all measures ordinally equivalent to the listed measures other than (Log-likelihood), likelihood ratio, (Generalized entailment) fail to satisfy this condition and Z, likelihood ratio, (Kemeny-Oppenheim) and (Kemeny-Oppenheim) succeed at satisfying it.

Now, what reasons do we have to not use some of the measures we introduced? First, some insights are obtained by considering the abstract requirements. Crucially, (1) final probability incrementality with prior-posterior dependence exclude (Carnap) and (Christensen-Joyce), (2) (Carnap) and (Log-ratio) have the unintuitive consequence that $C(T, E) = C(E, T)$ (call this *symmetry*), (3) (Difference), (Generalized entailment), (Log-ratio), (Carnap) and (Rips) depend on the prior of T , and (4) logical closure requirement excludes many of the measures.

Moreover, there is an issue with Z (Fitelson, 2021). Say E and E' are confirmationally independent regarding H just in case both $c(T, E|E') = c(T, E)$ and $c(T, E'|E) = c(T, E')$. Say E and E' are conflicting evidence regarding T iff $P(T|E) > P(T)$ while $P(T|E') < P(T)$. Fitelson (2021) has proven that any measure ordinally equivalent with Z, however, excludes the fairly intuitive possibility of the existence of confirmationally independent and yet conflicting evidence (he also gives a clear example of such a case).

Last but not least, for legal applications it seems that dependence on the prior probability is undesirable. We propose that at least two conceptual takes on confirmation is available. On one hand, say a scientific community pretty much agrees on the status of a given theory prior to an experiment. Then, after the experiment, it is a legitimate question what impact the experiment has on the status of that theory, and perhaps it makes sense that the prior status of that theory plays a role. On the other hand, in legal context, we would like (1) the expert's assessment not to depend on the expert's prior convictions about the hypothesis, and (2) the expert's statement to mean the same for various agents involved in the fact-finding process, even if they assign different priors to the hypothesis. For this reason, we propose that dependence on priors in legal evidence evaluation is an undesirable feature of a confirmation measure.

Thus, the general picture obtained, pictured in Table 1, seems to suggest that the likelihood ratio is a decent choice for our applications. This, of course, also applies to (Log likelihood) and (Kemeny-Oppenheim), which are ordinally equivalent to likelihood ratio, but the reasons to not use them in a legal context are that (Kemeny-Oppenheim) is conceptually more complex than likelihood ratio, and that thinking in terms of logarithms is not very natural for human agents.

References

- Aitken, C., & Taroni, F. (2008). Fundamentals of statistical evidence - a primer for legal professionals. *The International Journal of Evidence and Proof*, 12(3), 181–207. SAGE Publications.
- Aitken, C., Roberts, P., & Jackson, G. (2010). Fundamentals of probability and statistical evidence in criminal proceedings (Practitioner Guide No. 1), Guidance for judges, lawyers, forensic scientists and expert witnesses. *Royal Statistical Society's Working Group on Statistics and the Law*.
- Aitken, C., Taroni, F., & Thompson, W. (2003). How the probability of a false positive affects the value of dna evidence. *Journal of Forensic Science*, 48(1), 1–8. ASTM International.
- Balding, D. J. (2002). The DNA Database Search Controversy. *Biometrics*, 58(1), 241–244.
- Balding, D. J. (2004). Comment on: Why the effect of prior odds should accompany the likelihood ratio when reporting dna evidence. *Law, Probability and Risk*, 3(1), 63–64. Oxford Univ Press.
- Balding, D. J., & Donnelly, P. (1996). Evaluating DNA Profile Evidence When the Suspect Is

Measure	Reason not to use
(Difference)	dependence on priors, logical closure failure
(Log-ratio) and (Bayes factor)	satisfies law of likelihood, symmetry, dependence on priors, failure to satisfy logical closure
(Generalized entailment)	dependence on priors, independent conflicting evidence
(Christensen-Joyce)	excluded by final probability incrementality with prior-posterior dependence
(Carnap)	excluded by final probability incrementality with prior-posterior dependence, symmetry, logical closure failure
(Rips)	dependence on priors, failure of logical closure
(Christensen-Joyce)	excluded by final probability incrementality with prior-posterior dependence
(Kemeny-Oppenheim)	none of the above, but unnecessarily complex
(Log likelihood)	none of the above, but logarithms are hard for humans
(Likelihood ratio)	none of the above

Table 1: Reasons not to use various confirmation measures in legal fact-finding applications.

Identified Through a Database Search. *Journal of Forensic Sciences*, 41(4), 13961J.

Barnett, Z. (2020). Why you should vote to change the outcome. *Philosophy and Public Affairs*, 48(4), 422–446.

Behrman, B. W., & Davey, S. L. (2001). Eyewitness identification in actual criminal cases: An archival analysis. *Law and Human Behavior*, 25(5), 475–491.

Biedermann, A., Hicks, T., Taroni, F., Champod, C., & Aitken, C. (2014). On the use of the likelihood ratio for forensic evaluation: Response to Fenton et al. *Science & Justice*, 54(4), 316–318.

Bolinger, R. J. (2020). Explaining the justificatory asymmetry between statistical and individualized evidence. In *The social epistemology of legal trials*. Routledge.

Brennan, J. (2012). *The ethics of voting*. Princeton University Press.

Buckleton, J. S., Bright, J.-A., & Taylor, D. (2018). *Forensic dna evidence interpretation*. CRC press.

Carnap, R. (1962). *Logical foundations of probability*. Citeseer.

Cook, R., Evett, I., Jackson, G., & Jones, P. (1998). A hierarchy of propositions: Deciding which level to address in casework. *Science & Justice*, 38(4), 231–239.

Crupi, V. (2015). Confirmation. In E. N. Zalta (Ed.), *Stanford encyclopedia of philosophy*.

Crupi, V., Tentori, K., & Gonzalez, M. (2007). On Bayesian Measures of Evidential Support: Theoretical and Empirical Issues. *Philosophy of Science*, 74(2), 229–252.

Dawid, A. P. (1994). The island problem: Coherent use of identification evidence. In P. Freeman & A. Smith (Eds.), *Aspects of uncertainty: A tribute to D. V. Lindley* (pp. 159–170). John Wiley & Sons, New York.

Dawid, A. P. (2004). Which likelihood ratio? (Comment on “Why the effect of prior odds should accompany the likelihood ratio when reporting DNA evidence”, by Ronald Meester and Marjan Sjerps). *Law, probability and risk*, 3(1), 65–71. Oxford Univ Press.

Dawid, A. P., & Mortera, J. (1996). Coherent Analysis of Forensic Identification Evidence. *Journal of the Royal Statistical Society. Series B (Methodological)*, 58(2), 425–443.

Dawid, A. P., & Stockmarr, A. (2001). Comment on Stockmarr’s “Likelihood Ratios for Evaluating DNA Evidence When the Suspect Is Found through a Database Search”. *Biometrics*, 57(3), 976–980.

de Zoete, J. C., Fenton, N., Noguchi, T., & Lagnado, D. (2019). Resolving the so-called “probabilistic paradoxes in legal reasoning” with Bayesian networks. *Science & Justice*, 59(4), 367–379.

Democracy and decision: The pure theory of electoral preference. (1993).. Cambridge University Press.

Donnelly, P. (1995). Nonindependence of matches at different loci in DNA profiles: Quantifying the effect of close relatives on the match probability. *Heredity*, 75(1), 26–34.

Donnelly, P., & Friedman, R. D. (1999). DNA Database Searches and the Legal Consumption of Scientific Evidence. *Michigan Law Review*, 97(4), 931.

Dror, I. E., & Hampikian, G. (2011). Subjectivity and bias in forensic DNA mixture interpretation. *Science & Justice*, 51(4), 204–208. Elsevier BV. Retrieved from <https://doi.org/10.1016/j.scijus.2011.08>.

- Eggleston, R. (1978). *Evidence, proof and probability* (Vol. 2). Weidenfeld; Nicolson London.
- ENFSI. (2015). *Guidelines for evaluative reporting in forensic sciences*.
- Evetts, I. (1987). On meaningful questions: A two-trace transfer problem. *Journal of the Forensic Science Society*, 27(6), 375–381. Elsevier BV. Retrieved from [https://doi.org/10.1016/s0015-7368\(87\)72785-6](https://doi.org/10.1016/s0015-7368(87)72785-6)
- Evetts, I., Jackson, G., & Lambert, J. (2000). More on the hierarchy of propositions: Exploring the distinction between explanations and propositions. *Science & Justice*, 40(1), 3–10.
- Fenton, N., Berger, D., Lagnado, D., Neil, M., & Hsu, A. (2014). When “neutral” evidence still has probative value (with implications from the Barry George Case). *Science & Justice*, 54(4), 274–287.
- Finkelstein, M. (2009). *Basic concepts of probability and statistics in the law*. Springer.
- Fitelson, B. (1999). The plurality of bayesian measures of confirmation and the problem of measure sensitivity. *Philosophy of Science*, 66, S362–S378. University of Chicago Press. Retrieved from <https://doi.org/10.1086/392738>
- Fitelson, B. (2021). *A problem for confirmation measure Z*. [online manuscript].
- Foreman, L., Champod, C., Evetts, I. W., Lambert, J., Pope, S., & others. (2003). Interpreting dna evidence: A review. *International Statistical Review*, 71(3), 473–495. International Statistical Institute.
- Gelman, A., Katz, J. N., & Tuerlinckx, F. (2002). The mathematics and statistics of voting power. *Statistical Science*, 420–435. JSTOR.
- Gelman, A., King, G., & Boscardin, W. J. (1998). Estimating the probability of events that have never occurred: When is your vote decisive? *Journal of the American Statistical Association*, 93(441), 1–9. Taylor & Francis.
- Gillies, D. (1986). In defense of the popper-miller argument. *Philosophy of Science*, 53(1), 110–113. University of Chicago Press. Retrieved from <https://doi.org/10.1086/289295>
- Gross, S. R., O’Brien, B., Hu, C., & Kennedy, E. H. (2014). Rate of false conviction of criminal defendants who are sentenced to death. *Proceedings of the National Academy of Sciences*, 111(20), 7230–7235.
- Klobuchar, A., Steblay, N. K. M., & Caligiuri, H. L. (2006). Improving eyewitness identifications: Hennepin county’s blind sequential lineup pilot project. *Cardozo Pub. L. Pol’y & Ethics J.*, 4, 381–413. HeinOnline.
- Lempert, R. O. (1977). Modeling relevance. *Michigan Law Review*, 75, 1021–1057. JSTOR.
- Lindsay, R. C. L., Wells, G. L., & Rumpel, C. M. (1981). Can people detect eyewitness-identification accuracy within and across situations? *Journal of Applied Psychology*, 66(1), 79–89.
- Mayo, D. (2018). *Statistical inference as severe testing*. Cambridge University Press.
- Meester, R., & Sjerps, M. (2004a). Why the effect of prior odds should accompany the likelihood ratio when reporting DNA evidence. *Law, Probability and Risk*, 3(1), 51–62.
- Meester, R., & Sjerps, M. (2004b). Response to Dawid, Balding, Triggs and Buckleton. *Law, Probability and Risk*, 3(1), 83–86.
- Mulligan, C. B., & Hunter, C. G. (2003). The empirical frequency of a pivotal vote. *Public Choice*, 116(1), 31–54. Springer.
- National Research Council. (1992). *DNA technology in forensic science* [NRC I]. Committee on DNA technology in Forensic Science, National Research Council.
- National Research Council. (1996). *The evaluation of forensic DNA evidence* [NRC II]. Committee on DNA technology in Forensic Science, National Research Council.
- Niedermeier, K. E., Kerr, N. L., & Messeé, L. A. (1999). Jurors’ use of naked statistical evidence: Exploring bases and implications of the Wells effect. *Journal of Personality and Social Psychology*, 76(4), 533–542.
- Pardo, M. S. (2013). The Nature and Purpose of Evidence Theory. *Vanderbilt Law Review*, 66, 547–613.
- Park, R. C., Tillers, P., Moss, F. C., Risinger, D. M., Kaye, D. H., Allen, R. J., Gross, S. R., et al. (2010). Bayes Wars Redivivus – An Exchange. *International Commentary on Evidence*, 8(1).
- Redmayne, M. (2008). Exploring the proof paradoxes. *Legal Theory*, 14(4), 281–309. Cambridge University Press.
- Royall, R. M. (1997). *Statistical evidence: A likelihood paradigm*. Chapman; Hall/CRC.
- Shaer, M. (2016). The false promise of dna testing. *The Atlantic*, (June). Retrieved from <https://www.theatlantic.com/science/archive/2016/06/dna-testing/250111/>

- Smith, M. (2017). When does evidence suffice for conviction? *Mind*.
- Sprenger, J., & Hartmann, S. (2019). *Bayesian philosophy of science*. Oxford University Press.
- Stockmarr, A. (1999). Likelihood Ratios for Evaluating DNA Evidence When the Suspect is Found Through a Database Search. *Biometrics*, 55(3), 671–677.
- Taroni, F., Biedermann, A., Bozza, S., Garbolino, P., & Aitken, C. (2014). *Bayesian networks for probabilistic inference and decision analysis in forensic science* (2nd ed.). John Wiley & Sons.
- Thompson, S. G. (2007). Beyond a reasonable doubt-reconsidering uncorroborated eyewitness identification testimony. *UC Davis L. Rev.*, 41, 1487–1545. HeinOnline.
- Thompson, W. C. (2013). Forensic dna evidence: The myth of infallibility. In S. Krinsky & J. Gruber (Eds.), *Genetic explanations: Sense and nonsense* (pp. 227–347). Harvard University Press.
- Thomson, J. J. (1986). Liability and individualized evidence. *Law and Contemporary Problems*, 49(3), 199–219. JSTOR.
- Triggs, C. M., & Buckleton, J. S. (2004). Comment on: Why the effect of prior odds should accompany the likelihood ratio when reporting DNA evidence. *Law, Probability and Risk*, 3, 73–82.
- Wells, G. (1992). Naked statistical evidence of liability: Is subjective probability enough? *Journal of Personality and Social Psychology*, 62(5), 739–752. American Psychological Association.
- Wells, G. L., & Olson, E. A. (2003). Eyewitness Testimony. *Annual Review of Psychology*, 54(1), 277–295.
- Wixted, J. T., & Wells, G. L. (2017). The Relationship Between Eyewitness Confidence and Identification Accuracy: A New Synthesis. *Psychological Science in the Public Interest*, 18(1), 10–65.
- Wright, D., & McDaid, A. (1996). Comparing system and estimator variables using data from real line-ups. *Applied Cognitive Psychology*, 10, 75–84.