Chapter 5: Assessing evidential strength with the likelihood ratio

Marcello Di Bello and Rafal Urbaniak

The fallacies we considered in Chapter 3—base rate, prosecutor's and defense attorney's fallacy—show how the posterior probability of a hypothesis can be overestimated or underestimated. The posterior probability should reflect the evidence, but should not be identified with the probative value or strength of the evidence. A hypothesis may have a low posterior probability given the evidence, even though the evidence increases the probability of the hypothesis substantially.[1] So measuring evidential strength or probative value solely by posterior probabilities leaves out something crucial.

Another distinction worth making is between the global and local value of the evidence (Di Bello & Verheij, 2018). Let $E_1, E_2, \ldots, E_k$ be the total evidence presented at trial and H the ultimate hypothesis, say that the defendant is guilty of insider trading. The ultimate hypothesis is usually complex and can be thought of as the conjunction of several sub-hypotheses $H_1, H_2, \ldots H_k$. The total evidence bearing on the ultimate hypothesis should guide the final decision. But, as a preliminary step, it is useful to locally evaluate the impact of an individual piece of evidence E_i on the probability of a specific hypothesis H_i. When lay witnesses and experts testify at trial, the assessment of the evidential value of their individual testimonies should precede their aggregation into a whole, complex body of evidence.

This chapter articulates a probabilistic account of probative value or evidential strength that is incremental (it tracks changes in probability) and local (it is limited to individual pieces of evidence and specific hypotheses). We will argue that the likelihood ratio serves these purposes well. Section 1 explains why it fares better than another popular measure of evidential strength, the Bayes factor. (Appendix A broadens the discussion to probabilistic measures of confirmation and reaches a similar conclusion.) We then offer a couple of illustrations of how the likelihood ratio can be fruitfully deployed. Section 2 shows that it allows for a nuanced assessment of the strength of quantitative evidence, DNA match evidence. Section 3 examines how it can help to evaluate eyewitness testimony This should dispel the impression that the likelihood ratio is only suited for explicitly quantitative evidence.

Despite its versatility, however, the likelihood ratio should be deployed with care. It can be hard to interpret in practice, as we discuss in Section 4. In Section 5 we discuss evidential relevance, a topic closely related to that of evidential value. The likelihood ratio may categorize an item of evidence as irrelevant while intuitively the item is relevant. We explain this problem away by insisting that the likelihood ratio is *local* measure whose meaning is relative to specific hypotheses. To be sure, the value of an item of evidence is to be established both locally (relative to specific hypotheses) and globally (relative to the case as a whole). This suggests the need of formulating a more complex theory. We undertake this task in Part III of the book.

1. The likelihood ratio is better than the Bayes factor

A popular measure of evidential strength is the *Bayes factor*, corresponding to the likelihood of the evidence—the probability of the evidence given the hypothesis of interest, $P(E|H)$—divided by the

---

[1] Here is a more concrete example. Suppose an expert testifies that the blood found at the crime scene matches the defendant's and it is .05 probable that a person unrelated to the crime would match by coincidence. Absent other evidence to the contrary, it should initially be very likely that the defendant, as anyone else, had little to do with the crime. Say, for illustrative purposes, that the prior probability of the source hypothesis is .01, and let the probability of a match if the suspect is the source be approximately 1. By running Bayes' theorem, the posterior probability that the defendant is the source comes out to be roughly .17. While the match did not make it very likely that the defendant was the source of the traces, the posterior probability is seventeen times larger than the prior.

probability of the evidence P(E):

$$\text{BF}(E,H) = \frac{\text{P}(E|H)}{\text{P}(E)}.$$

It is a plausible measure as it appropriately deviates from one, its point of neutrality. Since, by Bayes' theorem, $\text{P}(H|E) = \text{BF}(H,E) \times \text{P}(H)$, the Bayes factor is greater than one if and only if the posterior probability $\text{P}(H|E)$ is higher than the prior probability $\text{P}(H)$. The greater the Bayes factor (for values above one), the greater the upward shift from prior to posterior probability, the more strongly E positively supports H. Conversely, the smaller the Bayes factor (for values below one), the greater the downward shift from prior to posterior probability, the more strongly E negatively supports H. If $\text{P}(H) = \text{P}(H|E)$, the Bayes factor equals one and the evidence has no impact on H.

The posterior probability of H given E could still be low even when the Bayes factor is significantly above one, indicating that the evidence is strongly probative of H despite the low posterior probability. So, as desired, this measure captures a dimension of the value of evidence that is not reflected in the posterior probability. Unfortunately, the Bayes factor suffers from three shortcomings that make it unsuitable for applications in trial proceedings.

The first shortcoming is a dependency on the prior probability of the hypothesis of interest. To see why, consider the denominator P(E). It can be unpacked by the law of total probability:

$$\text{P}(E) = \text{P}(E|H)\text{P}(H) + \text{P}(E|\neg H)\text{P}(\neg H). \tag{1}$$

So the Bayes factor can be written in a longer form:

$$\text{BF}(E,H) = \frac{\text{P}(E|H)}{\text{P}(E|H)\text{P}(H) + \text{P}(E|\neg H)\text{P}(\neg H)}. \tag{2}$$

What should be clear from this formulation is the dependency on the prior probabilities $\text{P}(H)$ and $\text{P}(\neg H)$. Indeed, suppose $\text{P}(E|H) = 1$ and $\text{P}(E|\neg H) = .1$. If $\text{P}(H) = .1$, $\text{P}(E)$, the denominator, is $.19$, and so the Bayes factor is approximately 5.26. If, however, $\text{P}(H) = .2$, the denominator is $.28$ and the Bayes factor is approximately 3.57. In fact, a more general look (Figure 1) shows that the prior probability can have larger impact on the Bayes factor than the likelihood $\text{P}(E|\neg H)$.

This is a strike against adopting this measure of evidential strength in legal fact-finding. For suppose an expert who is testifying in court is tasked with assessing the value of an item of evidence, say a DNA or fingerprint match. This assessment should not depend on the expert's prior convictions about the plausibility of the hypothesis. Further, judges and lay jurors should be in a position to understand the expert's assessment in the same way, even if they assign different prior probabilities to the hypothesis.[2]

A second reason to worry about the Bayes factor is its complexity. To see this, the catch-all alternative hypothesis $\neg H$ in the denominator can be replaced by a more fine-grained set of alternatives, $H_1, H_2, \ldots H_k$, provided H and these alternatives are exclusive and cover the entire space of possibilities (that is, they form a partition). The denominator becomes:

$$\text{P}(E) = \text{P}(E|H)\text{P}(H) + \sum_{i=1}^{k} \text{P}(E|H_i)\text{P}(H_i). \tag{3}$$

Assessing P(E) now looks quite difficult. It would require one to sift through the entire space of possibilities, as well as coming up with a sensible selection of priors probabilities for the several alternative hypotheses on hand. Whoever is tasked with assessing the strength of evidence—lay jurors, judges, or expert witnesses—might face too great a cognitive burden.

A third reason to hesitate about the Bayes factor comes from the problem of irrelevant conjuncts (Branden Fitelson, 1999; Gillies, 1986). Consider the hypothesis H = 'the suspect is guilty of murder' and suppose it is a fact that E = 'the suspect killed the victim.' Fact E does not establish guilt with certainty since guilt requires both *actus reus*, the killing, and *mens rea*, the intention. But clearly E provides positive support for H. Now consider a composite hypothesis H' = 'the suspect is guilty of murder *and* we live in a simulation built by aliens.' Presumably, the support E provides for H' should

---

[2]The requirement of prior independence is also in line with an objectivity requirement that the strength of evidence should not vary from one researcher to another (Bickel, 2012).

Bayes factor as a function of prior and P(E|~H).


Figure 1: Impact of the prior and likelihood of E given H for probabilities in (0, 0.05) and Bayes Factor restricted to (0, 250) for visibility.

be weaker than the support it provides for H. After all, the addition of a far-fetched hypothesis should weaken evidential support. This weakening, however, cannot be captured by the Bayes factor since both H and H' deductively entail E. In general, suppose $H \models E$ (and so, also, $H \wedge X \models E$). Then both $P(E|H)$ and $P(E|H \wedge X)$ equal 1. But this means that $\mathsf{BF}(H, E) = \mathsf{BF}(H \wedge X, E) = 1/P(E)$. So, contrary to what one would expect, the Bayes factors for the two support relations are equal.[3]

The discussion so far suggests that a good measure of evidential strength should satisfy three desiderata: (i) it does not depend on priors, (ii) places no unreasonably heavy cognitive requirements, and (iii) does not fall prey to the problem of irrelevant conjuncts. We will argue that a measure that satisfies these desiderata is the *likelihood ratio*. It is defined as

$$\mathsf{LR}(E, H, H') = \frac{P(E|H)}{P(E|H')},$$

where H' is a hypothesis that is a competing alternative to H. In the most straightforward case, H' is just the negation of H. But the competing hypotheses H and H' need not be one the negation of the other, a point to which we will return. As with the Bayes factor, support levels correspond to deviations from one. If the evidence is more likely given H than H', the ratio would be above one, and if the evidence is more likely given H' than H, the ratio would be below one. The greater the likelihood ratio (for values above one), the stronger the evidence in favor of H as contrasted with H'. The smaller the likelihood ratio (for values below one), the stronger the evidence in favor of the competing hypothesis H' as contrasted with H.

Note that both conditional probabilities in the numerator and denominator are needed for a correct assessment of the value of the evidence. For suppose we only relied on $P(E|H)$. In some cases, this conditional probability will be close to one. For instance, the probability that the blood from the crime matches the accused, if the accused is the source, $P(\mathsf{blood\ match}|\mathsf{source})$, may be close to one. Similarly, the probability that the DNA from the crime scene matches the accused, again if the accused is the source, $P(\mathsf{DNA\ match}|\mathsf{source})$, may also be close to one. Now, the DNA match should be stronger incriminating evidence than the blood type match because a specific genetic profile typically is less

---

[3]The same point can be made using irrelevant hypotheses that are not so far-fetched. For instance, suppose one hypothesis of interest is whether the victim was running in the park on a certain night, and the relevant piece of evidence is her footprints in the park. Perhaps, another hypothesis is whether she had wine at dinner later on. Clearly, whether she did is not obviously relevant to whether she was running in the park beforehand. However, one should be very hesitant to say that the evidential strength of the presence of footprints is the same relative to 'She was running in the park' and 'She was running in the park and had wine at dinner later on.' But this is what the Bayes factor would commit one to.

common than a specific blood type. Yet the quantity P($E|H$), by itself, makes no distinction here. The difference can instead be captured by the other conditional probability P($E|\neg H$). If the accused is *not* the source, the probability of a blood type match, while relatively small, should be higher than the probability of a DNA profile match. The likelihood ratio tracks both conditional probabilities, and thus it would be higher for the DNA match than the blood match, as desired.[4] For similar reasons, the strength of evidence cannot be measured by the probability P($E|\neg H$) alone. Consider an example by Triggs & Buckleton (2004). In a child abuse case, the prosecutor offers evidence that a couple's child rocks and that only 3% of non-abused children rock, P(child rocks|¬abuse) = .3. If it is unlikely that a child who is not abused would rock, that this child rocks might seem evidence of abuse. But this interpretation is mistaken. It could also be that 3% of abused children rock, P(child rocks|abuse) = .3. If rocking is equally unlikely under either hypothesis, rocking cannot count as evidence of abuse.

So, both the probability of the evidence given the hypothesis and the probability of the evidence given an alternative hypothesis should be part of any good measure of evidential strength (ENFSI, 2015; Royall, 1997; Triggs & Buckleton, 2004). The Bayes factor includes both probabilities, but—as seen before—it falls prey to three difficulties. The likelihood ratio tracks both conditional probabilities without falling prey to these difficulties.

First, unlike the Bayes factor, the likelihood ratio does not depend on the prior probability of the hypothesis. This is is apparent from the odds version of Bayes' theorem:

$$\frac{P(H|E)}{P(H'|E)} = \frac{P(E|H)}{P(E|H')} \times \frac{P(H)}{P(H')}. \tag{4}$$

If the likelihood ratio is greater (lower) than one, the posterior odds will be greater (lower) than the prior odds of H. The likelihood ratio, then, is a measure of the upward or downward impact of the evidence on the prior odds of two hypotheses H and H'. This fits nicely with the division of labor common in legal fact-finding between experts and decision-makers, judges or lay jurors. A prominent forensic scientist recommends that 'in criminal adjudication, the values of the prior odds and the posterior odds are matters for the judge and jury ' (Colin Aitken & Taroni, 2008, p. 194). Other scholars recommend that experts should 'not trespass on the province of the jury . . . and should generally confine their testimony to presenting the likelihood of their evidence under competing propositions' (CGG Aitken, Roberts, & Jackson, 2010, p. 42). The meaning of the likelihood ratio can be made more perspicuous by supplementing it with a graph that conveys visually the extent to which the evidence changes the probability of the hypothesis of interest (See Figure 2).

Second, the likelihood ratio is less cognitively burdensome than the Bayes factor. It does not require one to think about the probability of the evidence in general, P(E), say, the probability of a blood match under any possible scenarios. Direct and reliable estimation of this probability is difficult. From equation (1), it would require, besides an assessment of the conditional probabilities P($E|H$) and P($E|\neg H$), an assessment of the prior probabilities of P(H) and P($\neg H$). Instead, the likelihood ratio only requires an assessment of the conditional probabilities. In this sense, its calculation requires less information. This simplicity makes the likelihood ratio well-suited for presentation in trial proceedings. An expert, for instance, may testify that the blood-staining on the jacket of the defendant is ten times more likely to be seen if the wearer of the jacket hit the victim (prosecutor's hypothesis) rather than if he did not (defense's hypothesis) (CGG Aitken, Roberts, & Jackson, 2010, p. 38).

Finally, unlike the Bayes factor, the likelihood ratio is not susceptible to the problem of irrelevant hypotheses. For suppose P($E|H$) = P($E|H \wedge X$) = 1, where X is an additional hypothesis that is irrelevant to H. Note that LR(E, H) = 1/P($E|\neg H$), while LR($E, H \wedge X$) = 1/P($E|\neg H \vee \neg X$). Here the two denominators might differ. For example, suppose a fair coin is tossed three times. Let H = 'two first tosses resulted in two heads,' E = 'at least one of the two first tosses resulted in a head,' and X = 'the third toss resulted in heads.' Then P($E|H$) = 1, P($E|\neg H$) = 2/3, LR(E, H) = $\frac{1}{2/3}$ = 1.5. However, P($E|H \wedge X$) = 1, P($E|\neg(H \wedge X)$) ≈ .71, so LR($E, H \wedge X$) ≈ $\frac{1}{.71}$ = 1.4. Thus, the support, as measured by the likelihood ratio, can drop by adding a conjunct that is probabilistically irrelevant to the original hypothesis. In fact, this weakening of evidential support by adding an irrelevant conjunct holds in general for the likelihood ratio given sensible assumptions.[5]

---

[4]Specifically, P(DNA match|source)/P(DNA match|¬source) > P(blood match|source)/P(blood match|¬source).

[5]Branden Fitelson (2002) proved a general claim about irrelevant conjunctions. Hawthorne & Fitelson (2004) later strengthened this claim. The claim is that, if LR($E, H, \neg H$) > 1, P($E|X \wedge H$) = P($E|H$), and P($X|H$) ≠ 1, then LR($E, H, \neg H$) > LR($E, H \wedge X, \neg(H \wedge X)$). Crupi & Tentori (2010) raised a related problem. They point out that if LR(E, H) ≤ 1 and X is confirmationally

Figure 2: This graphical representation can supplement the likelihood ratio to convey visually the extent to which the evidence changes the probability of the hypothesis. This representation assumes that the two hypotheses in the likelihood ratio are one the negation of the other. In this case, the posterior probability of H equals $PO/1+PO$, where PO are the posterior odds.

All in all, the likelihood ratio outperforms the Bayes factor on several respects. But, of course, there could be other measures of evidential strength that fare even better. Other measures worth considering come from the literature in formal epistemology on confirmation theory. The expression 'confirmation' is more common in this literature than 'strength' (or value, support). A discussion of these measures, however, would detract us from the main task at hand. We therefore relegate it to Appendix A. In it, we show that the likelihood ratio is still, all things considered, the best measure on offer. A more general consideration to keep in mind is that there may well be two different questions here: (1) To what extend does a piece of evidence confirm our beliefs about a given hypothesis? (2) What is the strength (value, support) of a piece of evidence relative to a hypothesis? The two questions overlap to some extent. But the difference is that confirmation can depend on prior probabilities, while evidential strength should be kept separate from prior probabilities. Some confirmation measures may be seen as concerned with (1) rather than (2), and thus they are not always suitable for the evaluation of evidence in trial proceedings.

2. Match evidence and error probabilities

The two conditional probabilities that make up the likelihood ratio—for example, $P(E|H)$ and $P(E|\neg H)$—should be used in the evaluation of any form of evidence, both quantitative and non-quantitative. This section examines how a DNA match, a widely used form of quantitative evidence, should be evaluated by means of the likelihood ratio. The argument formulated here can be generalized to any 'match evidence.' The match can be between genetic profiles, fingerprints, blood types, bite marks, etc.

---

irrelevant conjunct to H with regard to E, then E will have the same negative or null impact on $H \wedge X$, that is $LR(E, H \wedge X) \leq LR(E, H)$. They find this counter-intuitive and argue that this can be avoided by switching to the Z confirmation measure (Crupi, Tentori, & Gonzalez, 2007). As we argue in Appendix A, the Z measure is prior-sensitive and therefore not fit for our purpose. Further, the phenomenon might not be deeply troubling either. If the likelihood ratio tracks how strongly the evidence supports a hypothesis, it should be no surprise that a more complex hypothesis—one obtained by adding an irrelevant proposition—-enjoys a lower support from the same evidence.

Consider an expert testimony that there is a genetic, DNA match between the traces at the crime scene and a sample from the defendant. This statement is evidence that the defendant was the *source* of the traces—that the materials found at the scene originated from the defendant. The match can also be evidence that the defendant was present at the scene or committed the crime, but these claims are more questionable, as the chain of inferences is weaker. For simplicity, let us focus on the source hypothesis. The likelihood ratio we should be concerned with is therefore the following:

$$\frac{P(\text{match}|\text{source})}{P(\text{match}|\neg\text{source})}.$$

How strongly does a match favor the source hypothesis? When experts testify about a DNA match, they often only provide the so-called *random match probability* as an indicator of evidential strength. This quantity expresses the probability that a random person, unrelated to the crime, would coincidentally match the crime scene profile. The random match probability coincides (roughly) with the denominator of the likelihood ratio. But what about the numerator? The hidden assumption often made is that the numerator must be close to one. So the likelihood ratio is approximated by a simple formula:

$$\frac{P(\text{match}|\text{source})}{P(\text{match}|\neg\text{source})} \approx \frac{1}{\text{random match probability}}.$$

Since the random match probability is usually an impressively low number, say 1 in 500 million, this is enough to ensure that the above ratio is significantly grater than one.

This analysis, though simple and elegant, lacks precision in at least two respects. First, it assumes that the numerator $P(\text{match}|\text{source})$ is close to one. But a DNA match need not track with 100% probability the fact that the suspect is the source. There could be false negative matches. In addition—and more importantly—equating the denominator $P(\text{match}|\neg\text{source})$ with the random match probability ignores the risk of false positive matches. This risk is not negligible (Shaer, 2016). The denominator, in fact, should depend on two sources of error: a false positive match and a coincidental match. These errors are quite distinct. For suppose two individuals—say the perpetrator and the defendant—happen to share the same DNA profile by coincidence. If an expert states that the crime scene sample and the defendant's sample match, this would be a coincidental match, not a false positive match. This risk of error is captured by the random match probability. But if the two samples do not actually match, and yet the expert says that they do, this would count as a false positive match, not a coincidental match. This risk of error is not captured by the random match probability.

Unlike a coincidental match, a false positive match is often caused by a human error in a number of circumstances (see (W. C. Thompson, 2013) for a more exhaustive treatment and multiple examples):

- **Cross-contamination of samples.** For instance, in Dwayne Johnson (2003) samples were accidentally swapped. In Lukis Anderson (2012), the genetic material was carried over by the paramedics. In one case, German police invested a considerable amount of time and effort searching for the so-called Phantom of Heilbronn, whose DNA profile was associated with many crimes. A bounty of EUR 300,000 was placed on her head. It turned out she was an innocent employee involved in the production of cotton swabs used across the country.
- **Mislabeling of samples.** For instance, in 2011 the Las Vegas Metropolitan Police Department acknowledged that samples of two men suspected of a 2001 robbery were switched, leading to the exclusion of the perpetrator and four years of incarceration for the other suspect. The mistake came to light only because the perpetrator was later arrested for another crime.
- **Misinterpretation of test results.** Single-source sample comparison is not easily prone to misrepresentation, but evidence mixtures—often needed in sexual assault cases—are complicated to interpret. For example, Dror & Hampikian (2011) re-examined a 2002 Georgia rape trial in which two forensic scientists had concluded that the defendant could not be excluded as a contributor of the crime traces. The evidence was sent to 17 lab technicians for re-examination. One of them agreed that the defendant could not be excluded as a contributor. Twelve considered the DNA exclusionary, and four found it inconclusive. If the quantity of DNA is limited, there is uncertainty about the number of contributors and about whether any alleles are missing. Ultimately, there is an element of subjectivity in mixed DNA interpretation.
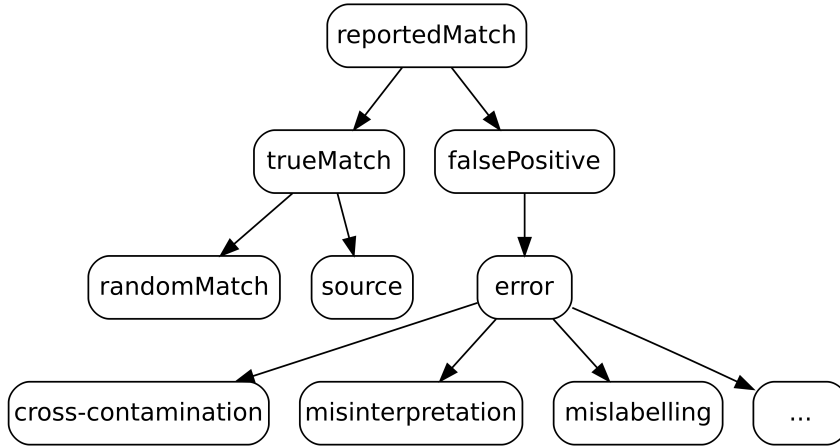
Figure 3: Dependencies between variables in the false positive problem.

The moral is that a careful evaluation of DNA evidence should take into account, besides the random match probability, the risks of false positive and false negative matches. This can be done using the likelihood ratio (Colin Aitken, Taroni, & Thompson, 2003). False positives are usually more worrisome than false negatives as they increase the risk of a mistaken conviction. That is also why we devoted to them more space in the foregoing discussion. But, for the sake of completeness, we will consider both.

This more refined analysis begins by making a conceptual distinction between true match and reported match. A true match is the fact that two samples actually carry the same genetic profile, while a reported match is a statement made by an expert that two samples match. A true match will exist not only if the suspect is the source, but also if, even though the suspect is not the source, the profiles are in fact the same due to a random, coincidental match. Similarly, a reported match might arise not only if there is a true match, but also if a false positive error has been made.[6] These possibilities are represented in Figure 3. For ease of reference, we will use the following abbreviations:

S The specimen comes from the suspect (source).
R A match is reported (reported match).
M There is a true match (true match).

In this set-up, the evidence to be assessed is the *reported* match relative to the pair of hypotheses S and $\neg S$. So the likelihood ratio we are after has the form:

$$\frac{P(R|S)}{P(R|\neg S)}.$$

With a few manipulations and assumptions in place, the likelihood ratio can be written as:[7]

$$\frac{P(R|S)}{P(R|\neg S)} = \frac{P(R|M)}{P(R|M)P(M|\neg S) + P(R|\neg M)P(\neg M|\neg S)} \qquad (5)$$

---

[6]The notion of a reported match is a simplification. The expert's opinion may be more fine-grained. A reported match could be conclusive or merely probable. A non-reported match could be a definite exclusion or a probable exclusion. The expert could also testify that the laboratory analyses were inconclusive. So, instead of a binary report, match versus non-match, the expert could testify about a conclusive match, probable match, inconclusive laboratory results, probable exclusion, definite exclusion. This complexity would further complicate the analysis we present in this section, but would not invalidate the conceptual framework.

[7]By the law of total probability, the denominator $P(R|\neg S)$ can be unpacked as $P(R \wedge M|\neg S) + P(R \wedge \neg M|\neg S)$. The latter, by the chain rule, is equivalent to $P(R|M \wedge \neg S)P(M|\neg S) + P(R|\neg M \wedge \neg S)P(\neg M|\neg S)$. A similar reasoning applies to the numerator.

Note that, as intended, the numerator $P(R|M)$ can be different from one, since a false negative reported match can occur (or, which is the same, a true match need not always occur). The denominator reflects the fact that there are two ways misleading evidence can arise: there is a true match and the suspect is not the source (because of a random, coincidental match), or there is no true match, and a false positive error has been made in the identification process.

To make the different sources of error more salient—false negative matches, false positive matches and random or coincidental matches—the likelihood ratio can be written, as follows:

$$\frac{P(R|S)}{P(R|\neg S)} = \frac{1 - FNP}{[(1 - FNP) \times RMP] + [FPP \times (1 - RMP)]} \tag{6}$$

This formula is the same as the earlier one. The expression FNP stands for the false negative probability $P(\neg R|M)$, so $1 - FNP$ equals the true positive probability $P(R|M)$. The expression FPP stands for the false positive probability $P(R|\neg M)$. The expression RMP stands for the random match probability $P(M|\neg S)$. A false positive or false negative probability track a human error, the possibility that a match may be reported (R) even without a true match ($\neg M$) or that a match may *not* be reported ($\neg R$) even with a true match (M). The random match probability, instead, tracks a coincidence of nature, the possibility that someone who is not the source ($\neg S$) could still be—coincidentally—a true match (M). If we set FNP to 0, we obtain the same formula derived by Colin Aitken, Taroni, & Thompson (2003), who did not consider false negatives. Their simpler formula reads:

$$\frac{P(R|S)}{P(R|\neg S)} = \frac{1}{RMP + [FPP \times (1 - RMP)]}$$

Let's now examine the impact of the error probabilities FNP and FPP on the likelihood ratio, holding fixed certain values of the random match probability. Figure 4 shows the impact of error rates (for values between 0 and .05). Random match probabilities are assumed to be in the order of 10^{-9} (often reported in the case of two single-source samples over ten or more loci). A small increase in the false positive probability can lower the likelihood ratio dramatically. For instance, with FNP set at .05, the likelihood ratio drops from 10^8 to 19 as FPP goes from 0 to .05. Interestingly, however, the impact of the false negative probability FNP (for values between 0 and .05) is rather negligible. For instance, if FNP is .05, the likelihood ratio goes from 20 to 19 as FPP moves from 0 to .05.

A similar analysis can be used to study the impact of error probabilities on the value of exculpatory DNA evidence, corresponding to a *negative* (reported) match $\neg R$. By replacing R with $\neg R$ in formula (5), the likelihood ratio becomes:

$$\frac{P(\neg R|S)}{P(\neg R|\neg S)} = \frac{P(\neg R|M)}{P(\neg R|M)P(M|\neg S) + P(\neg R|\neg M)P(\neg M|\neg S)} \tag{7}$$

$$= \frac{FNP}{FNP \times RMP + [(1 - FPP) \times (1 - RMP)]} \tag{8}$$

Keep in mind that the negative reported match $\neg R$ is evidence *against* the source hypothesis S so long as the likelihood ratio is below one. At the extreme, if the false negative probability FNP is zero, the numerator is zero. Thus, the likelihood ratio will be zero, as it should. In such a case, the negative match is completely exculpatory, and the posterior probability that the suspect is the source will also be zero. If the false negative probability is not zero, the greater the likelihood ratio (for values between 0 and 1),

So we have:

$$\frac{P(R|S)}{P(R|\neg S)} = \frac{P(R|M \wedge S)P(M|S) + P(R|\neg M \wedge S)P(\neg M|S)}{P(R|M \wedge \neg S)P(M|\neg S) + P(R|\neg M \wedge \neg S)P(\neg M|\neg S)}$$

Both numerator and denominator can be simplified because a reported match (R), given a true match obtains (M), is independent of whether the suspect is the source (S):

$$P(R|M \wedge S) = P(R|M \wedge \neg S) = P(R|M)$$

$$P(R|\neg M \wedge S) = P(R|\neg M \wedge \neg S) = P(R|\neg M)$$

Finally, in the numerator, let the probability of a true match if the suspect is the source be one:

$$P(M|S) = 1 \quad \text{so also} \quad P(\neg M|S) = 0.$$

This assumption holds in virtue of the meaning of the statements involved. That the suspect is the source of the crime sample entails, almost analytically, that the two samples must carry the same genetic profile.

Likelihood ratio of incriminating DNA evidence

Figure 4: Impact of the error probabilities on the likelihood ratio of incriminatory DNA evidence for RMP=10^-9 (grid approximation). The impact of FNP is minor.

the weaker the value of the exculpatory match. As Figure 5 shows, the likelihood ratio progressively moves away from zero as the false negative error probability increases. For instance, with FPP fixed at 0.05, it goes from 0 to .052. Interestingly, however, the impact of the false *positive* probability FPP on the likelihood ratio of exculpatory evidence is essentially null. For instance, if FNP is fixed at .05, the likelihood ratio moves from .05 to .052 as FPP goes from 0 to .05.

So, we have seen that even a seemingly small error probability outweighs the random match probability. If there are good reasons to worry about random matches, there are even better reasons to worry about error probabilities. But the impact of the error probabilities is not uniform, contrary to what one might intuitively think. The false positive probability has a marked impact on the value of incriminating DNA evidence (positive matches), while the false negative probability has a marked impact on the value of exculpatory DNA evidence (negative matches). So this analysis shows which error probabilities we should be concerned about in which circumstances.

What actual numbers should we use for false positive and false negative probabilities? Unfortunately, no serious attempt has been made to systematically quantify the relevant error probabilities. Sometimes, a lab discovers its own errors and reports them, but this is rare (W. C. Thompson, 2013). Anecdotal information suggest that false positive matches take place more often than coincidental matches would entail, but how often remains unclear. Regular proficiency tests used in accredited DNA laboratories involve comparison of samples from known sources, but they are criticized for being unrealistically easy (yet, it happens that analysts fail them). Sometimes, corrective action files are made available. They usually show relatively few false positive errors.[8] But because of the fragmentary data available, it is premature to conclude there is no reason for concern.

More data should be collected to plug in the right values of false positive and false negative probabilities in the likelihood ratio. Quite likely generic error frequencies will not be reliable enough, and relevant factors should be identified and used as predictors. To this end, error rate should be based on data that are fine-grained enough to document the error probabilities, FPP and FNP, corresponding to scenarios in which specific procedures, safeguards, or protocols are followed. As experts who testify about a match (or lack thereof) are cross-examined at trial, they could also testify about the procedures, safeguards and protocols that the laboratory technicians followed in the specific case. This case-specific

---

[8]For instance, the Santa Clara County district attorney's crime laboratory between 2003 and 2007 caught 14 instances of evidence cross-contamination with staff DNA, three of contamination by unknown person, and six of DNA contamination from other samples, three cases of DNA sample switch, one mistake in which the analyst reported an incorrect result, and three errors in the computation of the statistics to be reported.

Figure 5: Impact of the error probabilities on the likelihood ratio of exculpatory DNA evidence, assuming RMP=$10^{-}9$ (grid approximation). The impact of FPP is minor.

information could then be used to estimate error probabilities under different scenarios. This would yield a more individualized assessment of the value of match evidence.

We conclude this section by noting that other proposals exist in the literature for formulating the likelihood ratio of a genetic match which also incorporate error probabilities. Another, more general proposal is due to Buckleton, Bright, & Taylor (2018). But, interestingly, the likelihood ratio of a DNA match used in equations (5) and (7) agrees with this other proposal given certain assumptions. This convergence is encouraging. Here we briefly go through the derivation.

First, Buckleton and co-authors make the conceptual distinction between the probability that an error occurs, P(ERR), and the probability that a match is reported if an error occurs, P(R|ERR). Let err denote the probability of error. The intuition here is that a laboratory error is an event in which the identification fails to be proper for technical or procedural reasons, for instance, the samples were switched, or some equipment failed and produced an erroneous reading. Whether such a laboratory error occurs should not depend on whether the source hypothesis holds. So the derivation starts with the assumption that err is probabilistically independent of the source hypothesis S:

$$\text{err} = P(\text{ERR}) = P(\text{ERR}|S) = P(\text{ERR}|\neg S)$$

Separately, let k denote the probability of a reported match if an error occurs, also assumed to be independent of whether the source hypothesis is true:

$$k = P(R|\text{ERR}) = P(R|\text{ERR}, S) = P(R|\text{ERR}, \neg S)$$

Intuitively, even if normally a reported match R tracks to some extent the source of hypothesis S, say P($R|S$) > P($R|\neg S$), this connection breaks down once an error err occurs. Now the derivation:

$$LR = \frac{P(R|S)}{P(R|\neg S)} = \frac{P(R|\neg \text{ERR}, S)P(\neg \text{ERR}|S) + P(R|\text{ERR}, S)P(\text{ERR}|S)}{P(R|\neg \text{ERR}, \neg S)P(\neg \text{ERR}|\neg S) + P(R|\text{ERR}, \neg S)P(\text{ERR}|\neg S)}$$

$$= \frac{1 \times (1 - \text{err}) + k \times \text{err}}{RMP \times (1 - \text{err}) + k \times \text{err}} = \frac{1 - \text{err} + k \times \text{err}}{RMP \times (1 - \text{err}) + k \times \text{err}}$$

As before, the likelihood ratio is the ratio of the probabilities of a reported match if the suspect is the source and if the suspect is not the source. The numerator P($R|S$) can be split into two possible scenarios:

10

an error has not been made, or an error has been made. Accordingly, the numerator in the first line uses the law of total probability to split $P(R|S)$ into these two scenarios. Similarly, the numerator $P(R|\neg S)$ can be split into two cases: the suspect is not the source, but we are dealing with a random match, or the suspect is not the source, and an error has been made. An application of the law of total probability in the denominator mirrors this. The rest of the argument is just rewriting in terms of abbreviations, and algebraic manipulation.[9]

What is the connection of this formula to the one derived by Colin Aitken, Taroni, & Thompson (2003)? If an error guarantees a mistaken reported match, k becomes 1, err becomes the false positive rate, FPP. On this assumption, straightforward algebraic manipulation gives:

$$\frac{1 - FPP + 1 \times FPP}{RMP \times (1 - FPP) + 1 \times FPP} = \frac{1}{RMP \times (1 - FPP) + FPP}$$

$$= \frac{1}{RMP - FPP \times RMP + FPP}$$

It takes a straightforward algebraic manipulation to show that this formula is identical to the one derived by Colin Aitken, Taroni, & Thompson (2003):

$$\frac{1}{RMP + FPP \times (1 - RMP)}$$

3. Eyewitness identification and likelihood ratio

So far we paid attention to DNA evidence, a form of quantitative evidence widely used in trial proceedings. But the question arises whether evidence that is not explicitly quantitative, for example, eyewitness testimony, can also be evaluated by means of the likelihood ratio. We show that the answer is affirmative. In fact, there is no sharp divide between quantitative and non-quantitative evidence. A statement in court by an expert witness that the crime scene DNA matches the defendant is qualitative information, but its value is best assessed by means of numerical information, such as error rates and the random match probability. Similarly, a witness testifying 'I saw him!' is qualitative information, but again, its value is best assessed by means of numerical information about the risks of a false identification. As we will show, the likelihood ratio can be of service here.

To start, there is plenty of quantitative data about the risks of a false eyewitness identification. Consider first the statistics about false convictions. The rate of false convictions in death penalty cases in the United States is estimated at about 4%. How much of this can be attributed to false eyewitness identifications is hard to say exactly, but presumably quite a bit. In fact, a study of 340 exonerations in the years 1989-2003 showed that around 90% of false convictions in rape cases resulted from a false eyewitness identification. This percentage comes close to 100% in rape convictions in which the victim and the defendant were of difference races. In murder cases, 43% of false convictions resulted from a false identification by one or multiple eyewitnesses (Gross, O'Brien, Hu, & Kennedy, 2014). The innocence project reports that 56% of DNA exonerations in the United States involved mistaken eyewitness identification.

Field studies offer a more precise picture. These studies show that, in line-up identifications, eyewitnesses select filler individuals at a rate of 20-24% (Klobuchar, Steblay, & Caligiuri, 2006) . That is, around 20-24% of the time, an innocent person in a police line-up is incorrectly identified as the perpetrator. Similarly, a field study in Greater London (Wright & McDaid, 1996) and another study in Sacramento, California (Behrman & Davey, 2001) indicate that the false identification rate is around 20%. Even in experimental settings—where witnesses are less emotionally taxed—eyewitnesses identify a filler individual from a line-up in approximately 20% of the cases (S. G. Thompson, 2007).

In light of these empirical results and similar studies, the justice system has grown suspicious of eyewitness evidence over the last twenty years. This skepticism is welcome, but should not lead to discounting relevant evidence. Some may caution that the risks of mistaken identification should be assessed in the individual circumstances, and that blanketed statements that eyewitnesses are unreliable— even when backed up by well-researched statistics—are unhelpful. After all, judges and jurors should make determinations about the reliability of a specific eyewitness identification, not in general.

---

[9]The probability of a reported match R if no error occurs *and* the source hypothesis is false is the random match probability RMP, so $P(R|\neg ERR, \neg S) = RMP$. The probability that a reported match occurs when the source hypothesis is true and no error has made is assumed to be one, so $P(R|S, \neg ERR) = 1$.

Cross-examination is often thought to be the tool for this individualized assessment of the risk of error. The evidence law scholar Henry Wigmore in his monumental treatise on the law of evidence famously asserted that 'cross-examination is the greatest legal engine ever invented for the discovery of truth.' This assertion, however, has been subject to little empirical testing. In fact, empirical studies suggest that cross-examination is ineffective at detecting false identifications. In a series of experiments, subjects were asked to cross-examine eyewitnesses to determine whether they made accurate or mistaken identifications. Subjects showed little or no ability to make such discrimination (Wells & Olson, 2003, p. 285). In another experiment, a representative sample of 48 witnesses was cross-examined. Subjects ($n = 96$) viewing the cross-examination showed little ability to distinguish accurate from false identifications (Lindsay, Wells, & Rumpel, 1981).

The ineffectiveness of cross-examination at detecting errors might stem from its reliance on an intuitive, folk assessment of the risks of error, not on well-researched quantitative data. To remedy this, empirical research has identified a few canonical factors that affect the ability of a witness to correctly identify faces. These factors are typically divided into system variables and estimator variables (Behrman & Davey, 2001; Wells & Olson, 2003). The former refer to how the identification took place in a regimented setting, say whether it was a line-up or a show-up;[10] whether the line-up was simultaneous or sequential; whether the witness identified someone prior to the line-up. Instead, estimator variables refer to environmental conditions. The ability of a witness to make correct identifications is impaired by brief exposure, poor visibility (bad lighting or long distance) and a long interval between the first exposure and the moment of recollection. Other estimator variables include race (cross-racial identifications tend to be less reliable), stress (high stress can lead to worse memory), and weapon focus (the presence of a weapon weakens one's ability to make a correct identification).

A few examples can illustrate how data about estimator and system variables can be incorporated in the likelihood ratio. Consider a recent study about the correlation between distance and eyewitness reliability. It shows that the ratio of correct identifications (hits) to false identifications (false alarms) is 75% to 15% at 0 yard distance; 70% to 20% at 10 yard distance; 65% to 25% at 20 yard distance; 60% to 30% at 30 yard distance; 55% to 35% at 40 yard distance (**lampinen2014?**). These numbers can be used to fill in the conditional probabilities in the likelihood ratio for an eyewitness identification. Let id be the statement that the witness identifies the defendant as the person present at the scene, and let presence denote the fact that the witness was actually at the scene. The likelihood ratio of interest would look like this:

$$\frac{\mathsf{P(id|presence)}}{\mathsf{P(id|\neg presence)}}$$

Depending on distance, different numbers can be plugged in the numerator and the denominator. If, for example, the distance is 10 yards, the numerator should be .7 and the denominator .2. Thus, the likelihood ratio would be 3.5, an indication that the eyewitness identification is probative but its value is somewhat limited. This analysts is, of course, rather elementary and it is offered as merely illustrative. More-fine grained quantitative data are needed so that the other estimator variables besides distance can be taken into consideration, such as lighting, stress level, weapon focus, etc.

Research has also tackled the impact of system variables, especially in connection to another factor, witness confidence in the identification. It is a point of contention whether or not confidence is positively correlated with accuracy. A meta-analysis by Wixted & Wells (2017) offers a nuanced but overall optimistic picture. Their analysis focuses on identifications under *pristine conditions*, which require a certain set-up of system variables. Pristine conditions require, for example, a double-blind line-up containing one suspect and at least five fillers with no resemblance to the suspect. The witness is cautioned that the offender might not be in the line-up and there is no expectation that they identify someone. Needless to say, very few police departments run their line-ups in pristine conditions. But, as it turns out, if the identification occurs under pristine conditions, the high confidence of the witness is strongly predictive of an accurate identification. Most interestingly, witnesses with high confidence under pristine conditions should be around 90% accurate, a rather encouraging figure.[11]

Consider now two scenarios in which an expert is tasked with assessing the value of an eyewitness identification made during a police line-up. In one scenario, suppose the identification conditions are pristine. In line with the research we just discussed, the expert should testify: the probability of the

---

[10]A show-up refers to the observation of a single suspect by a witness in the field, typically at the crime scene. A line-up refers to the presentation of the suspect and several foils, either live or via photographs.

[11]The extent to which initial high confidence under pristine conditions is indicative of accuracy depends on the base rate of target-present lineups. In lab studies the base rate is about 50%, but in real-life circumstances, the best estimate is about 35%.

testimony if the suspect was present at the scene—P(id|presence)—is $.9 \pm .05$, and the probability of a false identification—P(id|¬presence)—is $.1 \pm .03$. In the second scenario, suppose the conditions are not pristine. The expert should testify that the probability of a correct identification is $.8 \pm .05$ and false identification is $.2 \pm .05$. These numbers are taken from the earlier research we cited about a 20-24% filler identification rate in police line-ups. We get two different *ranges* of likelihood ratios, lr_1 and lr_2. In the first scenario, the minimum and the maximum are as follows:

M: check and explain these numebers

$$\min(lr_1) = {}^{.95}\!/\!.07 \approx 13.57 \quad \max(lr_1) = {}^{.85}\!/\!.13 \approx 6.53$$

The likelihood ratio is in the range of 6.5-13.5. In the second scenario, the minimum and the maximum are as follows:

$$\min(lr_1) = {}^{.85}\!/\!.2 = 4.25 \quad \max(lr_1) = {}^{.75}\!/\!.3 \approx 2.5$$

So the likelihood ratio is in the range of 2.5-4.25. Figure 6 gives a compact illustration of the strength of an eyewitness identification under the two scenarios.


Figure 6: Impact of an eyewitnss identification on the prior probability of the presence hypothesis under pristine and non-pristine conditions.

The examples could continue, but the general point is this. A case-specific assessment of the risks of a false eyewitness identification can be carried out by taking into account estimator and system variables, as well as other relevant factors such as eyewitness confidence. Ideally, properly formatted data about all the factors we discussed could be used to develop a multivariate model. This model would allow an assessment of the risks of error under a variety of relatively well-specified circumstances. We are far from reaching this state of maturity, but even with the current research, an expert who is aware of the literature we have cited and of the specific circumstances of the case could assess the risks of a false identification quantitatively. An untrained jury member unaided by numerical data—even after viewing cross-examination—is unlikely to arrive at a better assessment of the risks of error involved.

We emphasize that the approach we are advocating here is continuous with the practice of cross-examination. Questions that are routine during cross-examination about time of day, distance, stress level, lighting conditions etc. elicit information about estimator variables as well as system variables.

Empirical research will help to determine how much poor visibility, long distance, heightened stress, etc. increase the risk of a false identification. As noted before, the reason why cross-examination is ineffective at detecting errors might well be that it does not rely on data about the magnitude of the risks of error under various circumstances. So, by combing numerical information in the likelihood ratio, cross-examination can better serve its purpose.[12]

We conclude this section by addressing the question whether an eyewitness identification is stronger or weaker evidence than a DNA match. This is not only a question that is often a source of confusion, but also a good opportunity to illustrate the applications of what we have learned about match evidence (previous section) as well as eyewitness testimony (this section).

Suppose eyewitness testimony in a case is exculpatory but there is also an incriminating DNA match. Which should prevail? It is tempting to take a clear-cut position. For example, one could reason that a DNA match is more reliable and thus should trump eyewitness testimony. Interestingly, an appellate court in New York took a similar stance in People v. Rush, 165 Misc. 2d 821 (N.Y. Sup. Ct. 1995) . A rape victim identified the defendant from a photograph a few weeks after the event. She also identified the defendant in a line-up two weeks later. At trial, however, she identified a spectator as her assailant, not the defendant. So, the eyewitness evidence was mixed but ultimately exculpatory. On the other hand, the defendant was seen in the vicinity of the crime scene three days prior. A DNA expert testified to a positive match between the crime scene genetic material and the defendant. The expert also testified that the random match probability was 1 in 500 million. Confronted with the question whether to believe the exculpatory eyewitness testimony or the incriminating DNA match, the court wrote:

> There can be little doubt, however, that the perils of eyewitness identification testimony far exceed those presented by DNA expert testimony ... This court is, therefore, satisfied that the testimony of even one DNA expert that there is a genetic match between the semen recovered from the victim of a rape and the blood of the defendant, a total stranger, and the statistical probability that anyone else was the source of that semen are 1 in 500 million is legally sufficient to support a guilty verdict.

We think such blanketed statements are unwarranted. The court framed the question as one of whether a DNA match *alone* can sustain a conviction. It gave no weight to the exculpatory testimony. But, even under this assumption, a random match probability of 1 over 500 million corresponds to a likelihood ratio of 20 (if the false positive error probability is .05) and 99 (if the false positive error probability is .01).[13] Are these numbers enough to convict? Assume the DNA match speaks directly to the hypothesis of guilt, a generous assumption in favor of the prosecution. By Bayes' theorem, if the prior probability of the guilt hypothesis is about .01, the posterior guilt probability would be .5 (if the false positive probability is .01) and .16 (if the false positive probability reaches .05). These low numbers do not even take into account the exculpatory testimony. Say the exculpatory testimony is assigned a likelihood ratio of .5. This would bring the posterior guilt probability to .33 and .09 respectively, a far cry from what is needed to convict an individual.[14]

M: Check calculations

These numbers are merely illustrative, but they convey an important lesson. The random match probability was fairly low, but to exclusively rely on it was a mistake. As the previous section has shown, the false positive probability can have a significant, negative impact on the value of an incriminating match. In addition, it is true that eyewitness testimony is fraught with problems. Its value might be low, especially if the identification was not conducted under pristine conditions. But without further details about estimator or system variables, it is hasty to make sweeping statements that any eyewitness testimony is worthless. This cautionary remark applies even more so to *exculpatory* testimony. The court reasoned that since the balance of the evidence tipped markedly against the accused—because

---

[12]A word of caution is necessary, however. The mechanics of cross-examination is quite complex. While in general it is meant to elicit further information that should guide the evaluation of the evidence, there are a few different ways the situation might develop, and several argumentative moves that can be made in such a context. On one hand, the additional information obtained can *rebutt* the eyewitness evidence, if it supports a hypothesis incompatible with the original one, as might happen when the eyewitness makes a statement contradicting something they have said previously. On the other, the new information might *udercut* the eyewitness testimony if it leads to the re-evaluation of the witness's reliability without leading to a new hypothesis, as might happen when the eyewitness provides new evidence about the lightning conditions. This complexity cannot be satisfactorily modeled by using likelihood ratios only. It requires a more sophisticated formalism, Bayesian networks (Di Bello, 2021). We will examine the mechanics of cross-examination in greater detail in Chapter 8.

[13]We applied formula (6), assuming a zero false negative probability FNP.

[14]We multiplied the DNA match likelihood ratio by the exculpatory testimony likelihood ratio and then computed the posterior probability with Bayes' theorem. We examine how to combine different lines of evidence together in Part III when we introduce Bayesian networks.

(incriminating) DNA evidence outweighs (exculpatory) eyewitness evidence—this was enough to convict. This is not quite right, however. Even if it has low value, exculpatory eyewitness testimony can in principle weaken an incriminating DNA match and create a reasonable doubt.[15] The value of eyewitness testimony should neither be exaggerated nor underestimated. As suggested in this section, a balanced assessment can be achieved by incorporating—as part of the likelihood ratio—numerical data about the risks of error under different circumstances.

4. Hypothesis choice

The argument so far has shown that the likelihood ratio is a fruitful conceptual framework for assessing the strength of the evidence, whether it is a genetic match or eyewitness testimony. However, its deployment is not devoid of challenges, and we now turn to a discussion of how they arise. One major challenge is the choice of the hypotheses H and H' that should be compared in the likelihood ratio. Generally speaking, they should 'compete' with one another—say, in a criminal trial, H is put forward by the prosecution and H' by the defense. But the two hypotheses need not be one the negation of the other, and thus there is leeway in their selection. This leeway makes the likelihood ratio quite versatile, but can also generate confusions and misunderstandings in its interpretation.

The first confusion stems from using *ad hoc* hypotheses. Suppose the prosecution argues that the defendant is the source of the traces found at the crime scene. This claim is well supported by laboratory analyses showing that the defendant genetically matches the traces. The defense, however, responds by putting forward the following *ad hoc* hypothesis: 'The crime stain was left by an unknown individual who happened to have the same genotype as the defendant'. Since the probability of the DNA match given either hypothesis is one, the likelihood ratio equals one (Evett, Jackson, & Lambert, 2000). The problem generalizes. For any item of evidence and any hypothesis H, there is an *ad hoc* competing hypothesis H^* that explains the evidence just as well as H does, so $P(E|H)/P(E|H^*) = 1$ (Mayo, 2018). If no further constraints are placed on the choice of the competing hypotheses—it would seem—no evidence could ever incriminate anyone.

The confusion here consists in thinking that, since the likelihood ratio equals one, the evidence must be worthless. But this is a mistake. To be sure, a genetic match cannot discriminate between hypothesis H (the crime stain came from the defendant) and H^* (the crime stain came from someone else who has the same genetic profile as the defendant). They are both equally well-supported by the match. This is just a fact about the world and a fact about how evidence works. The likelihood ratio correctly tracks this fact. But since the likelihood ratio is relative to a pair of hypotheses, the resulting assessment of the value of the evidence is also relative. Even if the match has no evidential value relative to H and H^*, it may have—and most likely does have—value relative to other hypotheses. It is a mistake to think that a single likelihood ratio must be associated with a piece of evidence. There are instead multiple likelihood ratios, each relative to a pair of competing hypotheses.

A second confusion stems from failing to clearly articulate the competing hypotheses, which then makes the assessment of the value of the evidence meaningless. A good illustration of this point is the case R. v. George (2007 EWCA Crim 2722). Barry George was accused of murdering TV celebrity Jill Dando. The key piece of incriminating evidence in the case was:

residue   A single particle of firearm residue was found one year later in George's coat pocket and matched the residue from the crime scene.

The defense argued that, since it was only one particle, there must have been contamination. The expert for the prosecution, however, testified that it was not unusual that a single particle would be found on the person who fired the gun. George was convicted, and his first appeal was unsuccessful.

After the first appeal, Dr. Evett from the Forensic Science Service in the United Kingdom worried that the evidence had not been properly assessed at trial. The jurors were presented with the conditional probability $P(\mathsf{residue}|H_d)$ of finding the firearm residue in George's coat given the defense hypothesis H_d (say, that George *did not* fire the gun). This probability was estimated to be quite low, indicating that the evidence spoke against the defense's hypothesis. But the jurors were not presented with the conditional probability $P(\mathsf{residue}|H_p)$ of finding the same evidence given the prosecutor's hypothesis H_p (say, that George *did* fire the gun). An expert witness, Mr. Keeley, was asked to provide both

_____

[15]A more detailed discussion of standards of proof can be founding in Part IV of the book.

conditional probabilities and estimated them to be $1/100$, which indicated that the firearm residue had no probative value. After new guidelines for reporting low level of firearm residue were published in 2006, the Forensic Science Service re-assessed the evidence and concluded that it was irrelevant. George appealed again in 2007, and relying on Keely's estimates, won the appeal.

At first, this case seems a good illustration of how likelihood ratios help to correctly asses the value of the evidence presented at trial. But what were the hypotheses that the expert compared? A study of the trial transcript shows that Keeley's choice of hypotheses was not transparent and the likelihood ratio based on them was therefore hard to interpret. On one occasion, Keeley took the prosecutor's hypothesis to be 'The particle found in George's pocket came from the gun that killed Dando' and the defense hypothesis to be 'The particle on George's pocket was inserted by contamination.' The problem is that the evidence is a logical consequence of either of them, so the conditional probability of the evidence given each hypothesis is one. The most charitable reading of the trial transcript suggests that the expert had in mind the hypotheses 'George was the man who shot Dando' and 'The integrity of George's coat was corrupted.' But Keeley was never clear that he compared these two hypotheses. Thus, the meaning of the likelihood ratio he provided remains elusive (see Fenton, Berger, Lagnado, Neil, & Hsu, 2014 for further details).

The confusion in the George case stems from the lack of a clear statement of the hypotheses. This lack of clarity can lead one to think that, since the likelihood ratio equals one (for a pair of not well-specified competing hypotheses), the evidence must prove nothing. To avoid this slippage, experts should perhaps only pick hypotheses that are exclusive (they cannot both be true) and exhaustive (they cannot be both false). In this way, the parties would not be able to pick *ad hoc* hypotheses and skew the assessment of the evidence in their own favor. Further, exclusive and exhaustive hypotheses would have likely forced the expert in the George case to be more precise.

But requiring that the hypotheses be always exclusive and exhaustive is not without complications either. Conceptually, it is a good idea to compare a hypothesis with its negation. Practically, this might not be viable, as often there are multiple different ways a hypothesis can turn out to be false. For consider an expert who decides to formulate the defense hypothesis by negating the prosecution hypothesis, say, 'the defendant did *not* hit the victim in the head.' This choice for the defense hypothesis can be unhelpful in assessing the evidence, because the required probabilities are hard to estimate. What is the probability that the suspect would carry such and such blood stain if he did not hit the victim in the head? This depends on whether he was present at the scene, what he was doing at the time and many other circumstances. As Evett, Jackson, & Lambert (2000) point out, in many real life cases, the choice of a particular hypothesis to be used by the expert in the evaluation of the evidence will depend on contextual factors.[16]

A third confusion stems from the fact that pairs of roughly similar hypotheses—which provide descriptions of the events at different levels of granularity—may be associated with different likelihood ratios for the same evidence (Fenton, Berger, Lagnado, Neil, & Hsu, 2014). This variability makes the likelihood ratio a seemingly arbitrary and easily manipulable measure of evidential value. An intriguing example of this phenomenon is a variation on the two-stain problem, originally formulated by Evett (1987). Suppose a crime was committed by two people, who left two stains at the crime scene: one on a pillow and another on a sheet. John Smith, who was arrested for a different reason, genetically matches the DNA on the pillow, but not the one on the sheet.[17] What likelihood ratio should we assign to the DNA match in question? Meester & Sjerps (2004) argue that there are three plausible pairs of hypotheses associated with numerically different likelihood ratios (see their paper for the derivations). The three options are listed below, where 1 in R is the random match probability of Smith's genetic profile and δ the prior probability that Smith was one of the crime scene donors.

---

[16]Another practical difficulty is that comparing exclusive and exhaustive hypotheses can be unhelpful for jurors or judges. In a paternity case, for example, the expert should not compare the hypotheses 'The accused is the father of the child' and its negation, but rather, 'The accused is the father of the child' and 'The father of the child is a man unrelated to the putative father' (Biedermann, Hicks, Taroni, Champod, & Aitken, 2014). The choice of the latter pair of competing hypotheses is preferable. Even though the relatives of the accused are potential fathers, considering such a far-fetched possibility would make the assessment of the evidence more difficult than needed.

[17]In the original formulation, two stains from two different sources were left at the crime scene, and the suspect's blood matches one of them. Let one hypothesis be that the suspect was one of the two men who committed the crime and the other hypothesis the negation of the first. Evett (1987) shows that the likelihood ratio of the match relative to these two hypotheses is $1/2q_1$ where q_1 is the estimated frequency of the characteristics of the first stain. The likelihood ratio does not depend on the frequency associated with the second stain. In general, if there are n bloodstains of different phenotypes, the likelihood ratio is $1/nq_1$. So the likelihood ratio depends on the number of stains but not on the frequency of the other characteristics.

H_p	H_d	LR
Smith was one of the crime scene donors.	Smith was *not* one of the crime scene donors.	$R/2$
Smith was the pillow stain donor.	Smith was not one of the crime scene donors.	R
Smith was the pillow stain donor.	Smith was not the pillow stain donor.	$R(2-\delta)/2(1-\delta)$

The first pair of hypotheses is the most general, encompassing any trace at the crime scene. The third pair is the most specific, focusing on just the trace on the pillow . The second pair lies in between the other two. There is significant variability across the three likelihood ratios. The first is half the second and third likelihood ratios (if δ is close to zero). In a trial, we could imagine prosecution and defense disagreeing about the right selection of hypotheses to compare. To their advantage, the prosecution might insist that the right comparison is given by the second or third pair, while the defense might insist that the right comparison is given by the first pair.

A closer scrutiny of the matter, however, reveals that the difference across likelihood ratios is not as troubling as it might seem. It would be surprising that the likelihood ratio for the same piece of evidence would remain the same when it is applied across hypotheses that make different commitments about what happened. So, once again, the likelihood ratio behaves as it should. It finely discriminates between the granularity of different hypotheses.[18]

The moral is this. The likelihood ratio is perfectly fine in its place, but its limits should be kept in mind. We should not overlook that it is relative to a pair of hypotheses and thus its assessment of the evidence is also relative. So we are dealing with a two-step process here: first, select a pair of hypotheses of interest; and second, assess whether the evidence makes a distinction between the two and to what extent. The likelihood ratio is useful for the second step, not the first. It provides no guidance in choosing the hypotheses. It is a tool to assess the value of the evidence *once* the hypotheses are chosen. The two steps must be spelled out clearly and kept separate. If they are not, confusions and misinterpretations of the evidence will occur. As noted at the start of this chapter, we are after a *local* measure of evidential value (strength, support), one that is circumscribed to a certain hypothesis or pair of hypotheses. The likelihood ratio serves this purpose well, but does nothing more than that. We elaborate on this point further as we discuss the notion of legal relevance in the next and final section.

5. Relevance and the small-town murder scenario

The Federal Rules of Evidence define relevant evidence as one that has 'any tendency to make the existence of any fact that is of consequence to the determination of the action more probable or less probable than it would be without the evidence' (rule 401). As before, we are dealing with a two-step process here: first, select a 'fact that is of consequence to the determination of the action,' sometimes called a material fact or hypothesis; and second, assess whether the evidence changes the probability of the material fact or hypothesis. More succinctly, two conditions must be met: (1) materiality and (2) probative value. An item of evidence is relevant if it has probative value for a hypothesis that is material ('of consequence to the determination of the action'). The likelihood ratio is helpful for assessing probative value, and in fact, it allows for a more fine-grained assessment because it formalizes *degrees* of probative value (or strength, support). We have illustrated how this can be done in the case of match evidence and eyewitness testimony. But the likelihood ratio does not directly address the question of materiality and therefore cannot provide a full account of relevance.

Consider now probative value more closely. The Rules define it in probabilistic language, an occurrence that has not gone unnoticed (CGG Aitken, Roberts, & Jackson, 2010; Colin Aitken & Taroni, 2004; Lempert, 1977; Lyon & Koehler, 1996; Sullivan, 2019). Specifically, the Rules invoke the idea of

---

[18]Interestingly, even though the likelihood ratios are numerically different, the posterior odds for the three pair of hypotheses are the same. To see why, note that the prior odds of the three H_p's in the table should be written in terms of δ. Following Meester & Sjerps (2004), the prior odds of the first hypothesis in the table are $\delta/1-\delta$. The prior odds of the second hypothesis are $(\delta/2)/(1-\delta)$. The prior odds of the third hypothesis are $(\delta/2)/(1-(\delta/2))$. In each case, the posterior odds — the result of multiplying the prior odds by the likelihood ratio — are the same: $R \times \delta/2(1-\delta)$. The reason for this is that the hypotheses are equivalent conditional on the evidence. Smith was one of the crime scene donors just in case he was the pillow stain donor, because he is excluded as the stain sheet donor. Smith was not one of the crime scene donors just in case he was not the pillow stain donor, because he is excluded as the sheet stain donor. Dawid (2004) cautions that the equivalence of hypotheses, conditional on the evidence, does not imply that they can all be presented in court. Meester & Sjerps (2004) recommend that the likelihood ratio should be accompanied by a tabular account of how a choice of prior odds (or prior probabilities) will impact the posterior odds, for a sensible range of priors.

a probability change, but do not mention the likelihood ratio. The two are related, however. A piece of evidence changes the probability of hypothesis H *if and only if* the likelihood ratio $P(E|H)/P(E|\neg H)$ is different from one. But note that the comparison here holds between H and its negation, not just between any two competing hypotheses H and H'. When the comparison is between hypotheses that are not one the negation of the other, likelihood ratio and probability change come apart. A piece of evidence may change the probability of a hypothesis—so that $P(H|E) \neq P(H)$—even though the likelihood ratio $prE|H/P(E|H')$ equals one. How can that be?

Suppose Fred and Bill attempted to rob a man. The victim resisted, was struck on the head and died. Say H_p stand for 'Fred struck the fatal blow' and H_d stand for 'Bill struck the fatal blow.' The hypotheses are not exhaustive. A missing hypothesis is 'The man did not die from the blow.' Suppose E is the information that the victim had a heart attack six months earlier. The likelihood ratio $P(E|H_p)/P(E|H_d)$ equals one since $P(E|H_p) = P(E|H_d)$. Yet E reduces the probability of both H_p and H_d. This might be confusing, but there is no inconsistency here. If what is material in the case is only whether either Bill or Fred did it, evidence about a heart attack would be irrelevant. The likelihood ratio correctly tracks this fact. On the other hand, if what is material is whether the victim died by natural or human causes, evidence about a heart attack would be relevant. Again, it is paramount not to confuse questions of materiality and questions of probative value.

The distinction between materiality and probative value also helps to deflect the force of a number of counterexamples against the likelihood ratio that we find in the literature. The vignette below was formulated by Ronald Allen and appears in the multi-authored discussion in (Park et al., 2010):

> **Small Town Murder.** A person accused of murder in a small town was seen driving to the small town at a time prior to the murder. The prosecution's theory is that he was driving there to commit the murder. The defense theory is an alibi: he was driving to the town because his mother lives there to visit her. The probability of this evidence if he is guilty equals that if he is innocent, and thus the likelihood ratio is 1 ..., and under what is suggested as the "Bayesian" analysis, it is therefore irrelevant. Yet, every judge in every trial courtroom ... would say it is relevant. And so we have a puzzle.

Counterexamples of this sort abound.[19] The puzzle is that the evidence is intuitively relevant and yet the likelihood ratio would seem to tell otherwise. There are a few things to say in response. First, as noted already, the likelihood ratio provides an account of probative value (strength, support), not relevance (which includes both materiality and probative value). Further, probative value is relative to a specific pair of hypotheses. The likelihood ratio may change depending on the selection of hypotheses.

Rule 401 makes clear that relevant evidence should have 'any tendency to make the existence of *any fact that is of consequence* [emphasis ours] to the determination of the action more probable or less probable.' Just because the likelihood ratio equals one for a specific selection of H and H', it does not follow that it equals one for *any* selection of H and H' which are of consequence to the determination of what happened. For example, the likelihood ratio of the evidence 'suspect was seen driving to town' relative to the pair of hypotheses 'suspect was in town' *versus* 'suspect was not in town' would be different from one. Whether the suspect was in town at all is presumably of consequence for determining what happened. So the fact that he was seen driving is helpful information for establishing whether or not he was in town, and the likelihood ratio can model this fact.

This discussion makes clear that relevance is a *global, holistic* notion. Legal cases often have a very complex structure. They consist of several pieces of evidence as well as several factual hypotheses to be assessed in light of the evidence. Some may be all-encompassing, such as the ultimate prosecutor's hypothesis 'The defendant committed insider trading.' Other hypotheses may be more circumscribed, such as 'The defendant visited the victim on Sunday.' An item of evidence is relevant so long as it can affect the probability of *any* material hypothesis, in one way or another.[20] Evidence is relevant so long as

---

[19]Here is another scenario: Suppose a prisoner and two guards had an altercation because the prisoner refused to return a food tray. The prisoner had not received a package sent to him by his family and kept the tray in protest. According to the defense, the prisoner was attacked by the guards, but according to the prosecution, he attacked the guards. The information about the package sent to the prisoner and the withholding of the tray fails to favor either version of the facts, yet it is relevant evidence (Pardo, 2013).

[20]A complication is that the choice of hypotheses needed to determine relevance might depend on other items of evidence, and so it might be difficult to determine relevance until one has heard all the evidence. This fact—Ronald Allen and Samuel Gross argue in (Park et al., 2010)—makes the probabilistic account of relevance impractical. But, in response, David Kaye points out that deciding whether a reasonable juror would find evidence E helpful requires only looking at what hypotheses or stories the juror would reasonably consider. Since the juror will rely on several clues about which stories are reasonable, this task is computationally easier than going over all possible combinations of hypotheses (Park et al., 2010).

it has a probabilistic impact on a sub-hypothesis involved in the case, even without having a recognizable probabilistic impact on the prosecutor's or defense's ultimate hypotheses. A local assessment of the value of the evidence using the likelihood ratio can help to make decisions about relevance. But in order to model the relationships between hypotheses formulated at different levels of complexity, we need to move past the likelihood ratio. As we shall see in later chapters, Bayesian networks can make better sense of this holistic perspective (de Zoete, Fenton, Noguchi, & Lagnado, 2019).

6. Conclusion

We discussed the likelihood ratio as a measure of evidential value (strength, support). We argued that it is a better fit in the legal settings than the Bayes factor. It is not heavily dependent on prior probabilities, it is flexible and is not liable to the problem of irrelevant conjuncts. We illustrated how the likelihood ratio can be deployed for evaluating match evidence and eyewitness testimony. We emphasized how it provides a conceptual framework that can be filled in with numerical data about the risks of error, false positive and false negative probabilities. This approach promises to strengthen the role of cross-examination as a tool for weeding out bad evidence.

But despite its promise, the likelihood ratio has limits. It is relative to a pair of (often rather circumscribed) hypotheses and its meaning should not be severed from them. Strictly speaking, there is no single likelihood ratio that corresponds to a single piece of evidence. Multiple hypotheses can be at issue in a legal case and the same piece of evidence can speak to them to different degrees. There is nothing surprising about that. Circumscribed sub-hypotheses are part of larger, more complex hypotheses, and finally, they all converge into the ultimate hypothesis. To model these complex relationships, however, a more elaborate theoretical framework is needed, Bayesian networks, our topic in several chapters to follow.

A. Appendix:Confirmation measures

Our terminology in this chapter included expressions such as 'strength of evidence,' 'value of evidence' or 'evidential support.' An expression often used in the philosophical literature is 'confirmation.' How is confirmation related to the discussion in this chapter? Bayesian confirmation theory is a sub-field of philosophy of science and epistemology. It aims to provide a probabilistic account of what it means for evidence to confirm a scientific theory. This naturally bears similarities with the question of what it means for evidence to support a hypothesis put forward in a trial. We will summarize some of the theories of Bayesian confirmation and then compare confirmation to evidential strength.

Confirmation can be understood as a function of an agent's degrees of belief—a measure of an agent's *firmness* of belief. On this account, the degree of confirmation that a piece of evidence E provides in favor of scientific theory T is defined as the posterior probability $P(T|E)$. This account can be defended by invoking three plausible requirements that any measure of confirmation should satisfy.

The first requirement is that the degree of confirmation is a continuous function of $P(T)$ and $P(E|T)$ which is non-decreasing in the first argument and non-increasing in the second argument. That is, increasing the prior, should not lower the confirmation level, and increasing the likelihood should not increase the confirmation level (Sprenger & Hartmann, 2019). Call this condition the *prior-posterior dependence*.[21] A corollary of the prior-posterior-dependence is that confirmation of T by E, $c(T,E)$, should track the posterior order-wise, that is, $c(T,E) > c(T,E')$ just in case $P(T|E) > P(T|E')$. A second requirement is that there should be a neutral point n such that E confirms (or disconfirms) T just in case $c(T,E) > n$ (or $c(T,E) < n$) and is neutral exactly at n. Call this the *qualitative-quantitative bridge*. Finally, a third requirement is *local equivalence*. Theories that are logically equivalent given the evidence should receive equal confirmation from this evidence.

Interestingly, all confirmation measures which satisfy prior-posterior dependence, qualitative-quantitative bridge, and local equivalence are strictly increasing functions of $P(H|E)$. Such measures are said to explicate confirmation as *firmness of belief*. Moreover, all functions satisfying these three

---

[21] Some formulations (Crupi, 2015) are a bit more general and include background knowledge K. In that setting, the corresponding requirement is called *Formality* and takes the confirmation to be a function of $P(H \wedge E|K), P(H|K)$ and $P(E|K)$. For the sake of simplicity, we will suppress the reference to K, unless required by the context.

conditions are ordinally equivalent.[22]

But the concept of confirmation cannot be fully captured by posterior probabilities. Even if the posterior T is low, one might still think that a given experiment speaks strongly in favor of T. Another problematic feature of confirmation understood as firmness of belief is this. If E confirms T, then for any T' that is excluded by T, E disconfirms T'. But now think of the small town murder scenario we already discussed: the fact that the suspect was seen in town seems to support both the prosecution's hypothesis that he committed the murder, and the defense's hypothesis, that he was in town to visit his mother. Confirmation as firmness cannot capture these intuitions.

Following the second edition of (Carnap, 1962), it is customary to distinguish another notion in the vicinity: confirmation as *increase in firmness* of belief. If we replace local equivalence with tautological equivalence $c(T, \top) = c(T', \top)$, where $\top$ is a logical tautology—the idea being that hypotheses are equally supported by empty evidence—we end up with another class of confirmation measures, those meant to capture *probabilistic relevance*. On this approach, E confirms (or disconfirms) T just in case $P(H|E) > P(H)$ (or $P(H|E) < P(H)$).

Here is a list of the most prominent confirmation measures on offer (Sprenger & Hartmann, 2019), normalized so that they all have neutral points at 0:

$$D(T, E) = P(T|E) - P(T) \qquad \text{(Difference)}$$

$$Lr(T, E) = log\left(\frac{P(T|E)}{P(T)}\right) \qquad \text{(Log-ratio)}$$

$$LL(T, E) = log\left(\frac{P(E|T)}{P(E|\neg T)}\right) \qquad \text{(Log-likelihood)}$$

$$K(T, E) = \frac{P(E|T) - P(E|\neg T)}{P(E|T) + P(E|\neg T)} \qquad \text{(Kemeny-Oppenheim)}$$

$$Z(T, E) = \begin{cases} \frac{P(T|E) - P(T)}{1 - P(T)} & \text{if } P(T|E) \geq P(T) \\ \frac{P(T|E) - P(T)}{P(T)} & \text{if } P(T|E) < P(T) \end{cases} \qquad \text{(Generalized entailment)}$$

$$S(T, E) = P(T|E) - P(T|\neg E) \qquad \text{(Christensen-Joyce)}$$

$$C(T, E) = P(E)(P(T|E) - P(T)) \qquad \text{(Carnap)}$$

$$R(T|E) = 1 - \frac{P(\neg T|E)}{P(-T)} \qquad \text{(Rips)}$$

The Log-likelihood, the Kemeny–Oppenheim measure and the simple likelihood ratio (not in the list above) are ordinally equivalent (and no other pair on the list is). Log-likelihood makes calculations of joint support additive, and the Kemeny–Oppenheim measure has the nice feature of ranging from -1 to 1 and having 0 as a neutral point. If one accepts the likelihood ratio, these two can be used on some occasions when these additional features are needed.

So we face a seemingly radical plurality of confirmation measures. These measures, however, can be unified by a normalizing procedure described in (Crupi, Tentori, & Gonzalez, 2007). Leaving out the details, this normalization process yields another family of confirmation measures:

$$Z_\alpha = \begin{cases} Z(h, e)^\alpha & \text{if } P(T|E) \geq P(T) \\ -|Z(h, e)|^\alpha & \text{otherwise.} \end{cases}$$

Z_α is an S-shaped function where the parameter α describes the curvature. This function defines a family of measures, call them Z-measures. A good reason to favor Z-measures is that they generalize logical entailment. Take any $k > 0$ and say $v(E, T) = k$ iff $E \models T$ and $v(E, T) = -k$ iff $E \models \neg T$, and $v(E, T) = 0$ otherwise. Call *logical closure* the requirement that if $v(E, T) > v(E', T')$, then $c(E, T) > c(E', T')$. Logical closure is satisfied only by the generalized entailment measure and the likelihood ratio (along with the ordinally equivalent measures, the Log-likelihood and Kemeny-Oppenheim measure). The other measures in the table all fail to satisfy this requirement.

Besides failure of logical closure, there are further reasons to reject the other measures. First, Carnap's and Christensen-Joyce's measure fail to meet the requirement of prior-posterior dependence. Second,

---

[22]Measure c is ordinally equivalent to measure c' just in case always $c(E, T) \gtreqless c(E', T')$ iff $c'(E, T) \gtreqless c'(E', T')$.

Measure	Reason not to use
(Z)	dependence on priors, contradiction between confirmational independence and conflicting evidence
(Difference)	dependence on priors, logical closure failure
(Log-ratio) and (Bayes factor)	satisfies law of likelihood, symmetry, dependence on priors, failure to satisfy logical closure
(Generalized entailment)	dependence on priors, independent conflicting evidence
(Christensen-Joyce)	excluded by final probability incrementality with prior-posterior dependence
(Carnap)	excluded by final probability incrementality with prior-posterior dependence, symmetry, logical closure failure
(Rips)	dependence on priors, failure of logical closure
(Christensen-Joyce)	excluded by final probability incrementality with prior-posterior dependence
(Kemeny-Oppenheim)	none of the above, but unnecessarily complex
(Log likelihood)	none of the above, but logarithms are hard for humans
(Likelihood ratio)	none of the above

Table 1: Reasons not to use various confirmation measures in legal fact-finding applications.

Carnap's measure and Log-ratio have the counterintuitive consequence that $C(T, E) = C(E, T)$ (call this *symmetry*). Third, the difference measure, generalized entailment, the Log-ratio, Carnap's and Rips's measures depend on the prior of T.

Dependence on prior probabilities is particularly problematic in the legal setting. For we would like the expert's assessment not to depend on the expert's prior convictions about the hypothesis. Further, the expert's statement should carry the same meaning for various agents involved in the fact-finding process, even if they assign different priors to the hypothesis. For this reason, dependence on priors in legal evidence evaluation is a particularly undesirable feature. This might be less pressing in the scientific setting. Say a scientific community agrees on the status of a given theory prior to an experiment. Then, after the experiment, it is a legitimate question what impact the experiment has on the status of that theory, and perhaps it makes sense that the prior status of that theory plays a role in the theory's confirmation.

So the only measures that withstand scrunity so far are the likelihood ratio (and ordinarilly equivalent measures) and Z-measures. But the latter suffer from another problem. Say E and E' are confirmationally independent regarding H just in case both $c(T, E|E') = c(T, E)$ and $c(T, E'|E) = c(T, E')$. Say E and E' are conflicting evidence regarding T iff $P(T|E) > P(T)$ while $P(T|E') < P(T)$. Branded Fitelson (2021) has proven that any measure ordinally equivalent with Z excludes the fairly intuitive possibility of the existence of confirmationally independent and yet conflicting evidence.

Table 1 summarizes the reasons not to adopt a certain measure of confirmation. The likelihood ratio emerges as the clear winner. The Log likelihood and Kemeny-Oppenheim's measures are ordinally equivalent to the likelihood ratio. The reason not to use them in the legal setting is a matter of convenience. Kemeny-Oppenheim's measure is conceptually more complex than likelihood ratio, and thinking in terms of logarithms is unnatural for ordinary reasoners.

Finally, we should mention a general argument for preferring the likelihood ratio. This argument was developed by Heckerman (1988). Suppose you have background information b, and consider a hypothesis H when you obtain a piece of evidence E. A belief update, $U(H, E, b)$, together with your prior stance about H, should determine your posterior belief in H. If we denote conditional belief in H given E, without assuming it is probabilistic yet, as $H|E$, this means that there should be a function f such that:

$$H|E, b = f(U(H, E, b), H|b)$$

where, on this approach, $H|E, b$ and $U(H, E, b)$ and $H|b$ are all real numbers, and f is required to be continuous in both arguments and monotonically increasing when the other is held constant. Call this the *update requirement*. Bayesian updating is a particular case: the selection of H and b, with a joint distribution in the background, determines a function from $P(H|b)$ to $P(H|E, b)$.

Further, assume the *consistency property*: if the arguments are logically equivalent, then belief update yields the same value:

$$[H_1 \Leftrightarrow H_2, E_1 \Leftrightarrow E_2] \Rightarrow U(H_1, E_1, b_1) = U(H_2, E_2, b_2)$$

The next assumption has it that when you update on two items of evidence, the order of updating shouldn't matter and the result should be the same as updating on the joint evidence. The *combination property* requires that there is a function g that is continuous in both arguments and monotonically increasing in each argument when the other is held constant such that:

$$U(H, E_1 E_2, b) = g(U(H, E_1, b), U(H, E_2, E_1, b))$$

A general result in group theory by Aczel is that any continuous monotonic function of two arguments that satisfies the associativity relation must be additive in some transformed space. In this particular case, the result entails that any update that satisfies the update condition, and the consistency and combination properties is the arithmetic difference of a posterior and prior belief up to an arbitrary monotonic transformation, that is, that there are monotonic functions h and i such that:

$$h(U(H, E, b)) = i(H|E, b) - i(H|b)$$

In the probabilistic context, where $H|E$ is $P(H|E)$, the likelihood ratio satisfies the update requirement and has the combination and consistency properties. Accordingly, the transformation that makes it additive is the logarithmic function.

Now, the *independence correspondence property* has it that if E_1 and E_2 are conditionally independent given H and given $\neg H$, then $U(H, E_2, E_1, b) = U(H, E_2, b)$. The key result is that any probabilistic update satisfying the independence correspondence property must be a monotonic transformation of the likelihood ratio. This means that a list of fairly intuitive general conditions on what an update function should be like entails that this update function is just a variant of the likelihood ratio.

References

Aitken, CGG, Roberts, P., & Jackson, G. (2010). Fundamentals of probability and statistical evidence in criminal proceedings (Practitioner Guide No. 1), Guidance for judges, lawyers, forensic scientists and expert witnesses. *Royal Statistical Society's Working Group on Statistics and the Law*.

Aitken, Colin, & Taroni, F. (2004). *Statistics and the evaluation of evidence for forensic scientists* (Vol. 16). Wiley Online Library. https://doi.org/10.1002/0470011238

Aitken, Colin, & Taroni, F. (2008). Fundamentals of statistical evidence - a primer for legal professionals. *The International Journal of Evidence and Proof*, *12*(3), 181–207. https://doi.org/10.1350/ijep.2008.12.3.296

Aitken, Colin, Taroni, F., & Thompson, W. (2003). How the probability of a false positive affects the value of DNA evidence. *Journal of Forensic Science*, *48*(1), 1–8. https://doi.org/10.1520/jfs2001171

Behrman, B. W., & Davey, S. L. (2001). Eyewitness identification in actual criminal cases: An archival analysis. *Law and Human Behavior*, *25*(5), 475–491. https://doi.org/10.1023/A:1012840831846

Bickel, D. R. (2012). The strength of statistical evidence for composite hypotheses: Inference to the best explanation. *Statistica Sinica*, 1147–1198.

Biedermann, A., Hicks, T., Taroni, F., Champod, C., & Aitken, C. (2014). On the use of the likelihood ratio for forensic evaluation: Response to Fenton et al. *Science & Justice*, *54*(4), 316–318. https://doi.org/10.1016/j.scijus.2014.04.001

Buckleton, J. S., Bright, J.-A., & Taylor, D. (2018). *Forensic DNA evidence interpretation*. CRC press.

Carnap, R. (1962). *Logical foundations of probability*.

Crupi, V. (2015). Confirmation. In E. N. Zalta (Ed.), *Stanford encyclopedia of philosophy*.

Crupi, V., & Tentori, K. (2010). *Irrelevant conjunction: Statement and solution of a new paradox. 77*(1), 1–13. https://doi.org/10.1086/650205

Crupi, V., Tentori, K., & Gonzalez, M. (2007). On Bayesian Measures of Evidential Support: Theoretical and Empirical Issues. *Philosophy of Science*, *74*(2), 229–252. https://doi.org/10.1086/520779

Dawid, A. P. (2004). Which likelihood ratio? (Comment on 'Why the effect of prior odds should accompany the likelihood ratio when reporting DNA evidence,' by Ronald Meester and Marjan Sjerps). *Law, Probability and Risk*, *3*(1), 65–71.

de Zoete, J. C., Fenton, N., Noguchi, T., & Lagnado, D. (2019). Resolving the so-called "probabilistic paradoxes in legal reasoning" with Bayesian networks. *Science & Justice*, *59*(4), 367–379. https://doi.org/10.1016/j.scijus.2019.03.003

Di Bello, M. (2021). *A probabilistic analysis of cross-examination using bayesian networks*. https://doi.org/10.1111/phis.12209

Di Bello, M., & Verheij, B. (2018). Evidential reasoning. In *Handbook of legal reasoning and argumentation* (pp. 447–493). Springer.

Dror, I. E., & Hampikian, G. (2011). Subjectivity and bias in forensic DNA mixture interpretation. *Science & Justice*, *51*(4), 204–208. https://doi.org/10.1016/j.scijus.2011.08.004

ENFSI. (2015). *Guidelines for evaluative reporting in forensic sciences*.

Evett, I. W. (1987). On meaningful questions: A two-trace transfer problem. *Journal of the Forensic Science Society*, *27*(6), 375–381. https://doi.org/10.1016/s0015-7368(87)72785-6

Evett, I. W., Jackson, G., & Lambert, J. A. (2000). More on the hierarchy of propositions: Exploring the distinction between explanations and propositions. *Science & Justice*, *40*(1), 3–10. https://doi.org/10.1016/S1355-0306(00)71926-5

Fenton, N., Berger, D., Lagnado, D., Neil, M., & Hsu, A. (2014). When 'neutral' evidence still has probative value (with implications from the Barry George Case). *Science & Justice*, *54*(4), 274–287. https://doi.org/10.1016/j.scijus.2013.07.002

Fitelson, Branden. (1999). The plurality of bayesian measures of confirmation and the problem of measure sensitivity. *Philosophy of Science*, *66*, S362–S378. https://doi.org/10.1086/392738

Fitelson, Branden. (2002). Putting the irrelevance back into the problem of irrelevant conjunction. *Philosophy of Science*, *69*(4), 611–622. https://doi.org/10.1086/344624

Fitelson, Branded. (2021). *A problem for confirmation measure Z*. [online manuscript].

Gillies, D. (1986). In defense of the popper-miller argument. *Philosophy of Science*, *53*(1), 110–113. https://doi.org/10.1086/289295

Gross, S. R., O'Brien, B., Hu, C., & Kennedy, E. H. (2014). Rate of false conviction of criminal defendants who are sentenced to death. *Proceedings of the National Academy of Sciences*, *111*(20), 7230–7235. https://doi.org/10.1073/pnas.1306417111

Hawthorne, J., & Fitelson, B. (2004). *Discussion: Re-solving irrelevant conjunction with probabilistic independence*. *71*(4), 505–514. https://doi.org/10.1086/423626

Heckerman, D. (1988). Ax axiomatic framework for belief updates. *Machine Intelligence and Pattern Recognition*, *5*, 11–22.

Klobuchar, A., Steblay, N. K. M., & Caligiuri, H. L. (2006). Improving eyewitness identifications: Hennepin county's blind sequential lineup pilot project. *Cardozo Pub. L. Pol'y & Ethics J.*, *4*, 381–413.

Lempert, R. O. (1977). Modeling relevance. *Michigan Law Review*, *75*, 1021–1057.

Lindsay, R. C. L., Wells, G. L., & Rumpel, C. M. (1981). Can people detect eyewitness-identification accuracy within and across situations? *Journal of Applied Psychology*, *66*(1), 79–89.

Lyon, T. D., & Koehler, J. J. (1996). Relevance ratio: Evaluating the probative value of expert testimony in child sexual abuse cases. *Cornell Law Review*, *82*, 43.

Mayo, D. (2018). *Statistical inference as severe testing*. Cambridge University Press.

Meester, R., & Sjerps, M. (2004). Why the effect of prior odds should accompany the likelihood ratio when reporting DNA evidence. *Law, Probability and Risk*, *3*(1), 51–62. https://doi.org/10.1093/lpr/3.1.51

Pardo, M. S. (2013). The Nature and Purpose of Evidence Theory. *Vanderbilt Law Review*, *66*, 547–613.

Park, R. C., Tillers, P., Moss, F. C., Risinger, D. M., Kaye, D. H., Allen, R. J., . . . Kirgis, P. F. (2010). Bayes Wars Redivivus – An Exchange. *International Commentary on Evidence*, *8*(1). https://doi.org/10.2202/1554-4567.1115

Royall, R. M. (1997). *Statistical evidence: A likelihood paradigm*. Chapman; Hall/CRC.

Shaer, M. (2016). The false promise of DNA testing. *The Atlantic*. Retrieved from https://www.theatlantic.com/magazine/archive/2016/06/a-reasonable-doubt/480747/

Sprenger, J., & Hartmann, S. (2019). *Bayesian philosophy of science*. Oxford University Press.

Sullivan, S. P. (2019). A Likelihood Story: The Theory of Legal Fact-Finding. *University of Colorado Law Review*, *90*(1), 1–66.

Thompson, S. G. (2007). Beyond a reasonable doubt-reconsidering uncorroborated eyewitness identification testimony. *UC Davis L. Rev.*, *41*, 1487–1545.

Thompson, W. C. (2013). Forensic DNA evidence: The myth of infallibility. In S. Krimsky & J. Gruber (Eds.), *Genetic explanations: Sense and nonsense* (pp. 227–347). Harvard University Press.

Triggs, C. M., & Buckleton, J. S. (2004). Comment on: Why the effect of prior odds should accompany the likelihood ratio when reporting DNA evidence. *Law, Probability and Risk*, *3*, 73–82.

Wells, G. L., & Olson, E. A. (2003). Eyewitness Testimony. *Annual Review of Psychology*, *54*(1), 277–295. https://doi.org/10.1146/annurev.psych.54.101601.145028

Wixted, J. T., & Wells, G. L. (2017). The Relationship Between Eyewitness Confidence and Identification Accuracy: A New Synthesis. *Psychological Science in the Public Interest*, *18*(1), 10–65. https://doi.org/10.1177/1529100616686966

Wright, D. B., & McDaid, A. T. (1996). Comparing system and estimator variables using data from real line-ups. *Applied Cognitive Psychology*, *10*, 75–84.